

# Task Parallelism

Prof. Guido Araujo

[www.ic.unicamp.br/~guido](http://www.ic.unicamp.br/~guido)

# Agenda

- Task parallelism: programming model
- Where are the limitations?

# Agenda

- Task Parallelism: programming model
- Where are the limitations?

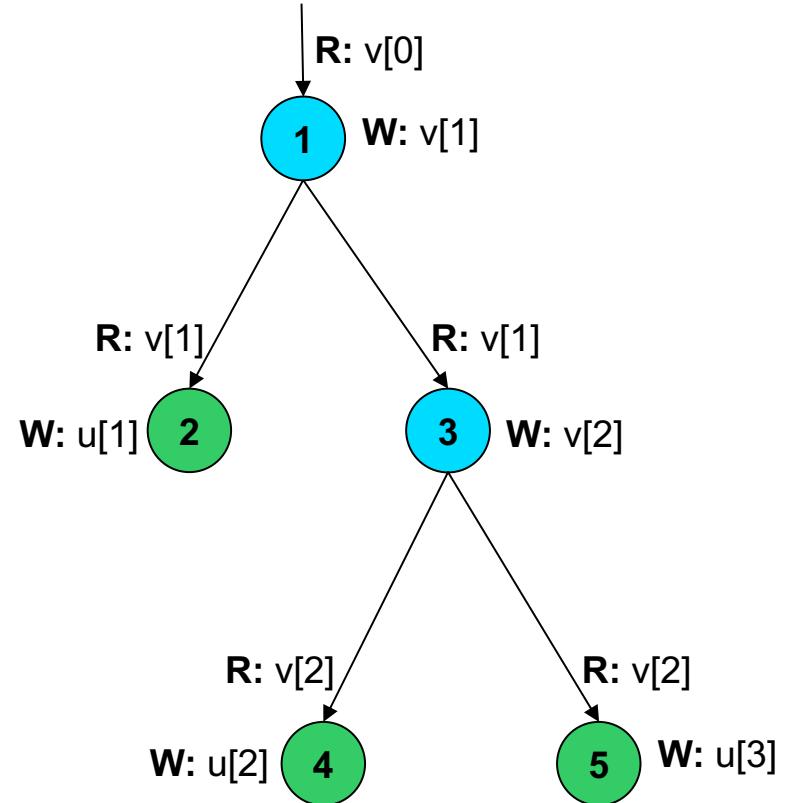
# Task Parallelism

## The OpenMP 4.0 task programming model

```
for (int i=1, j=1; i<N; i++) {  
    #pragma omp task in(v[i-1]) out(v[i])  
    fun1(&v[i-1], &v[i]); —————→ Create a task!  
  
    for (int k=0; k<i; k++, j++) {  
        #pragma omp task in(v[i]) out(u[j])  
        fun2(&v[i], &u[j]); —————→ Create a task!  
    }  
  
    fun3( 3 * i ); —————→ Create NO tasks!  
}
```

# Task Parallelism

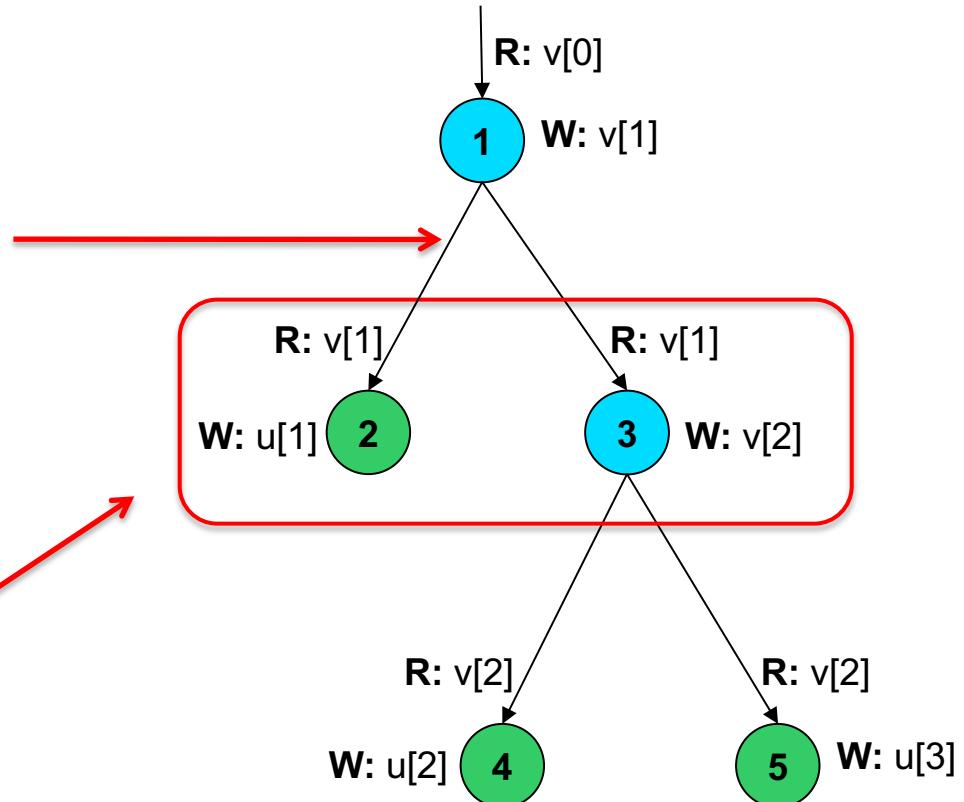
```
for (int i=1, j=1; i<N; i++) {  
    #pragma omp task in(v) out(v)  
    fun1(&v[i-1], &v[i]);  
  
    for (int k=0; k<i; k++, j++) {  
        #pragma omp task in(v) out(u)  
        fun2(&v[i], &u[j]);  
    }  
    fun3( 3 * i );  
}
```



# Task Parallelism Advantages

No need to communicate data. Just detect dependence

Parallelism between iterations of different loops

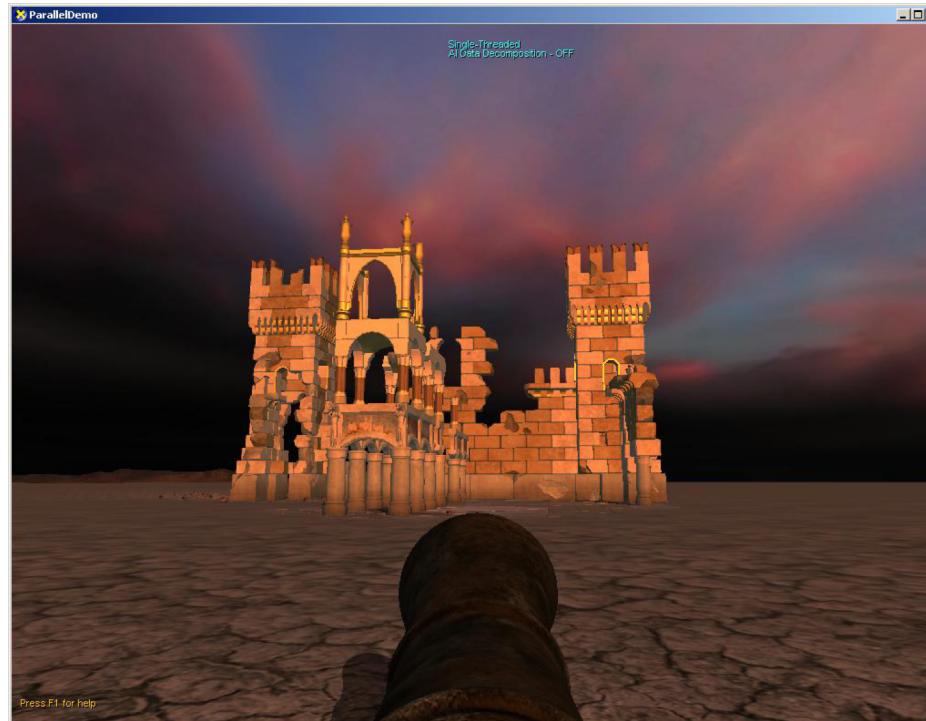


# But graphs are not always that simple!



# Task parallelism in action

- Destroy the Castle
  - Developed by Intel using OpenMP 3.X
  - Game based on DirectX 8.1



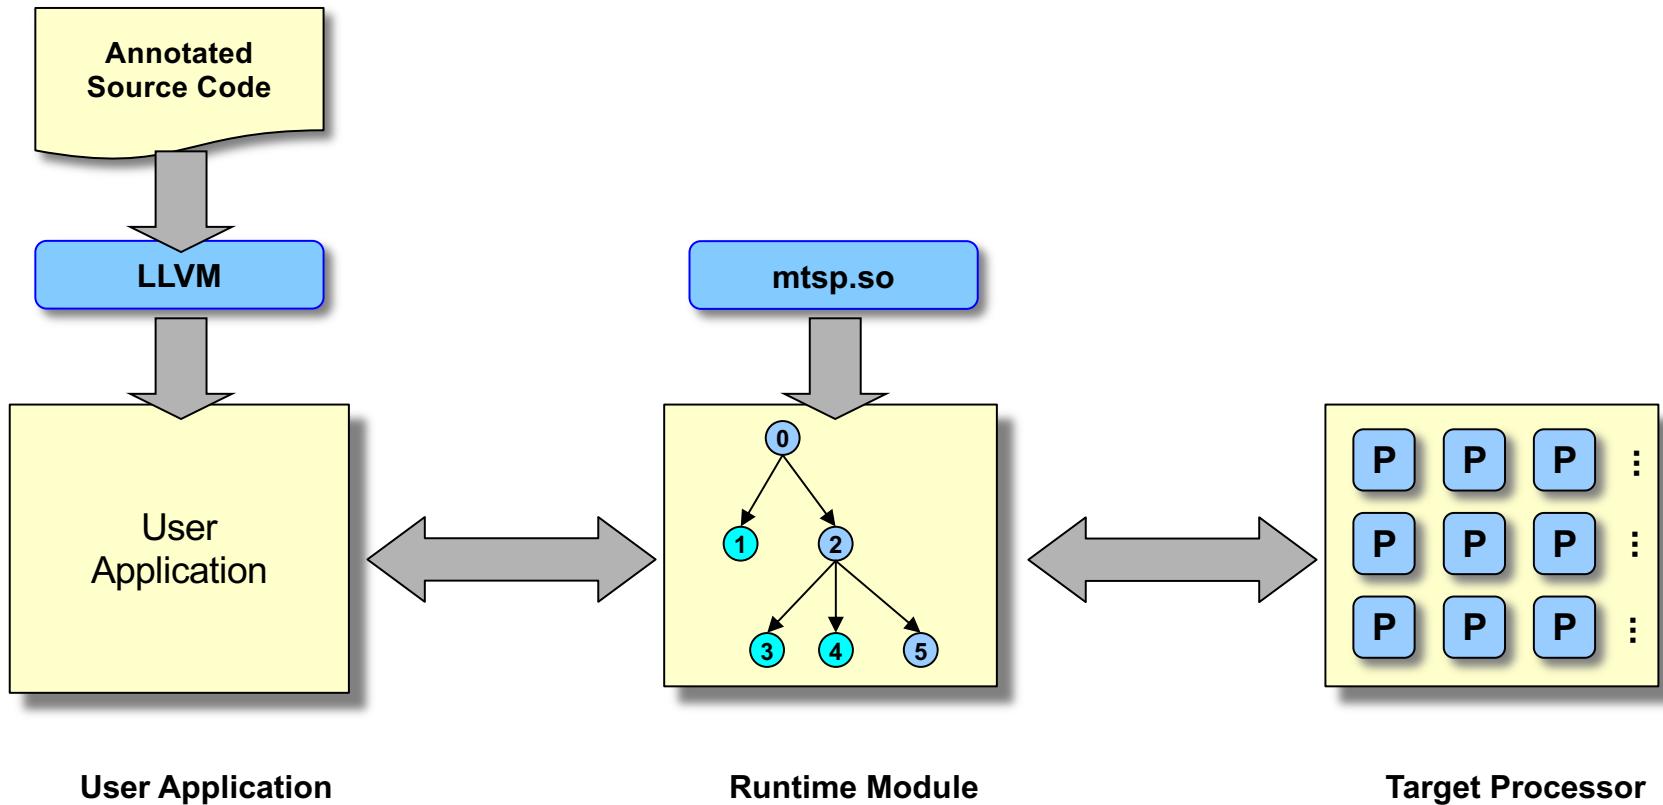
# Task parallelism in action



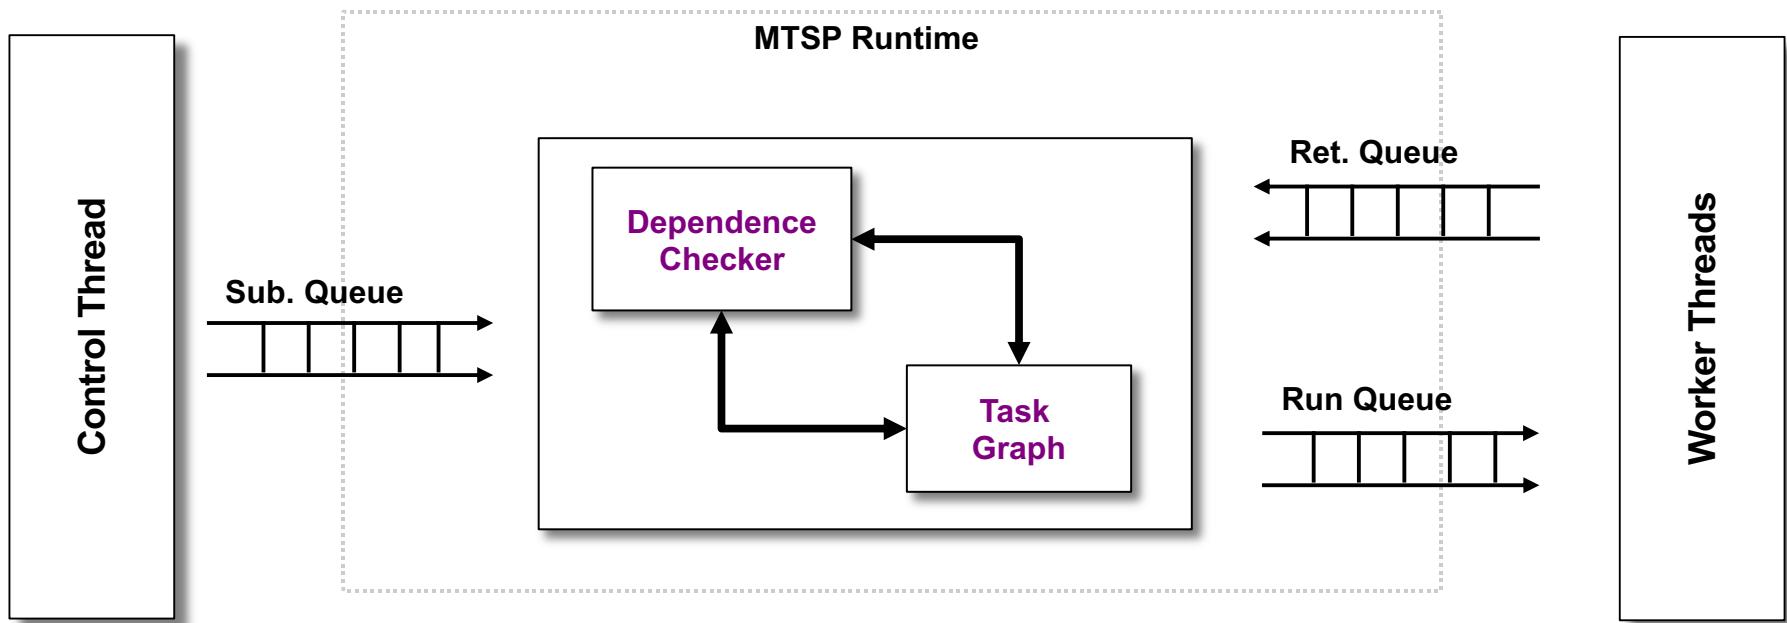
# Agenda

- Task parallelism: programming model
- Where are the limitations?

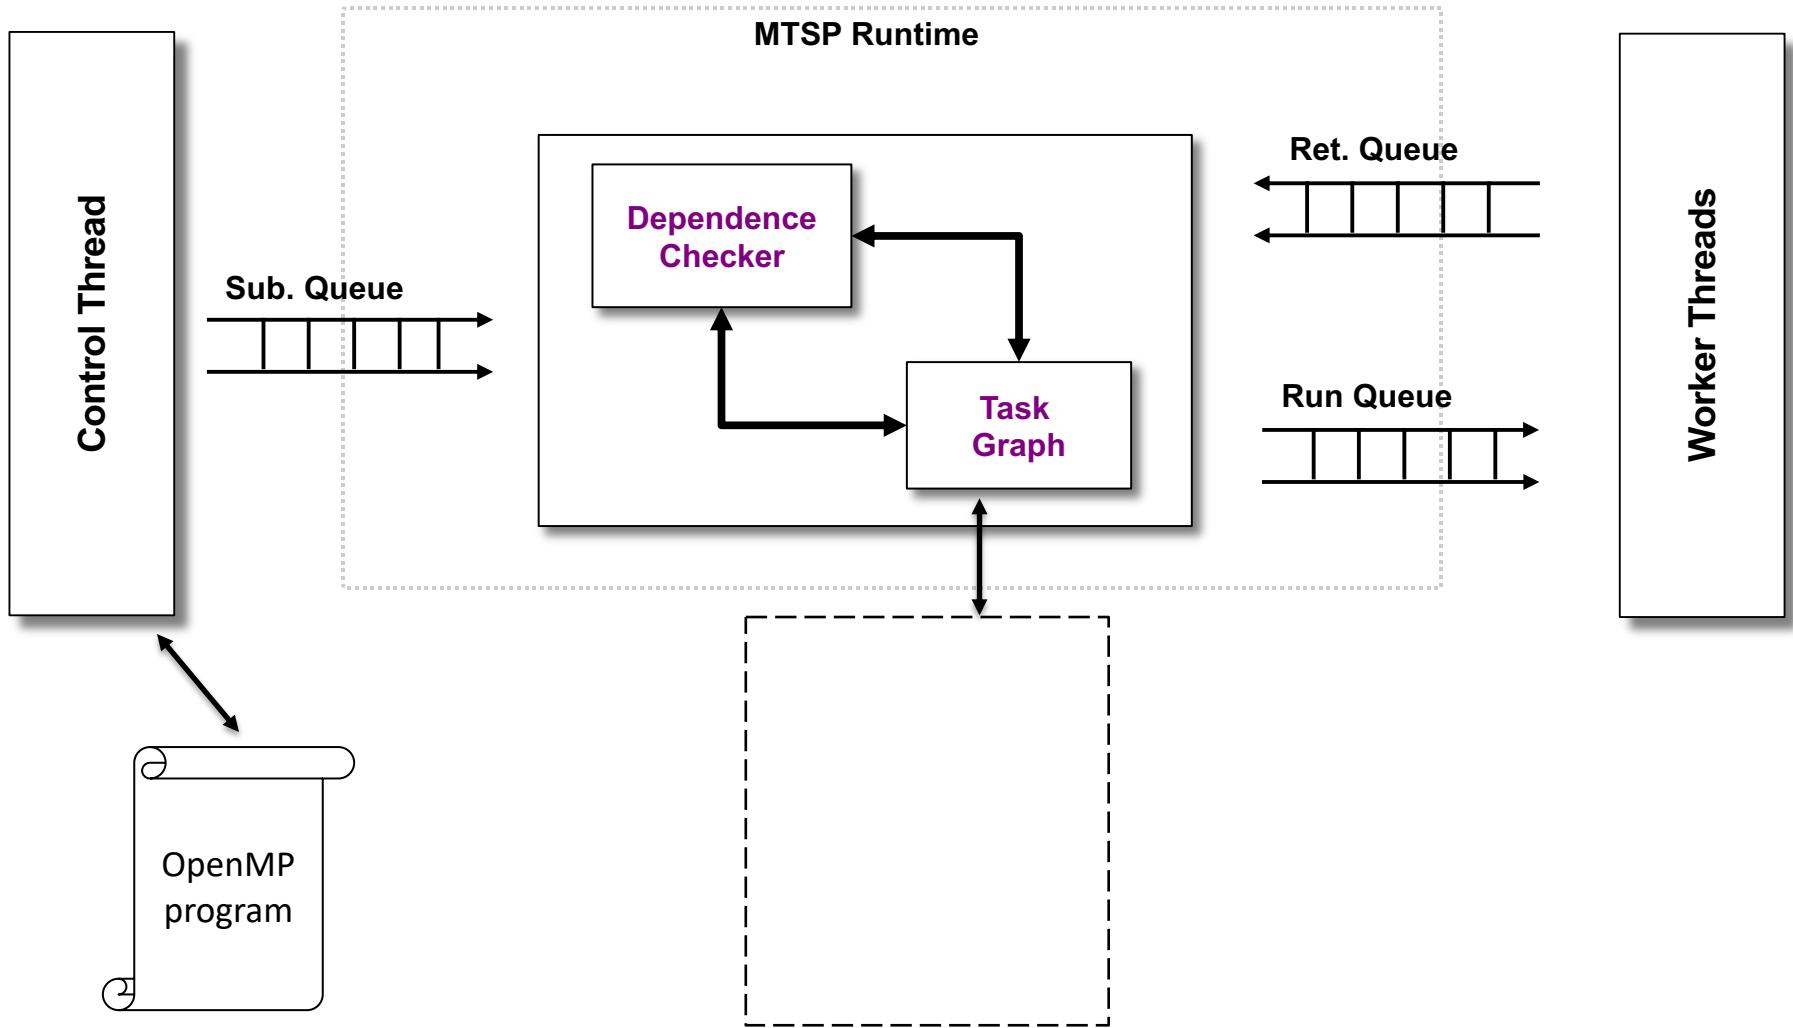
# Basic elements of a task scheduling system



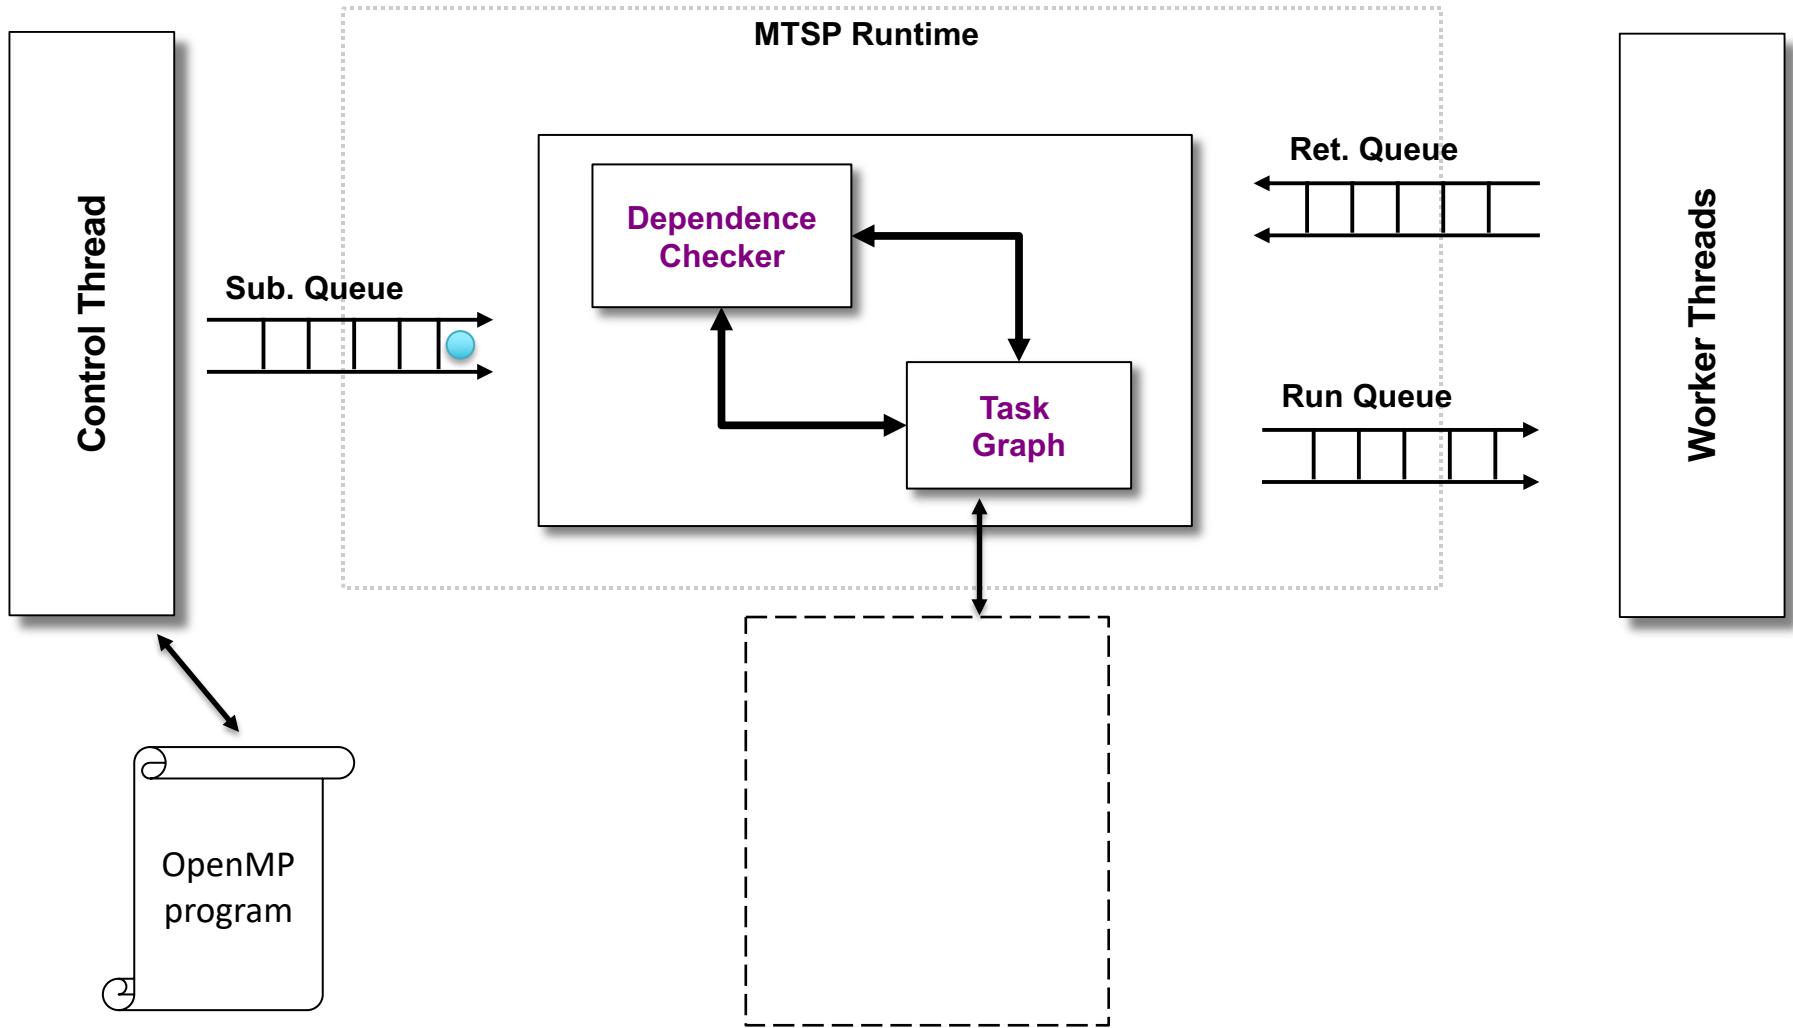
# Basic elements of a task scheduling system



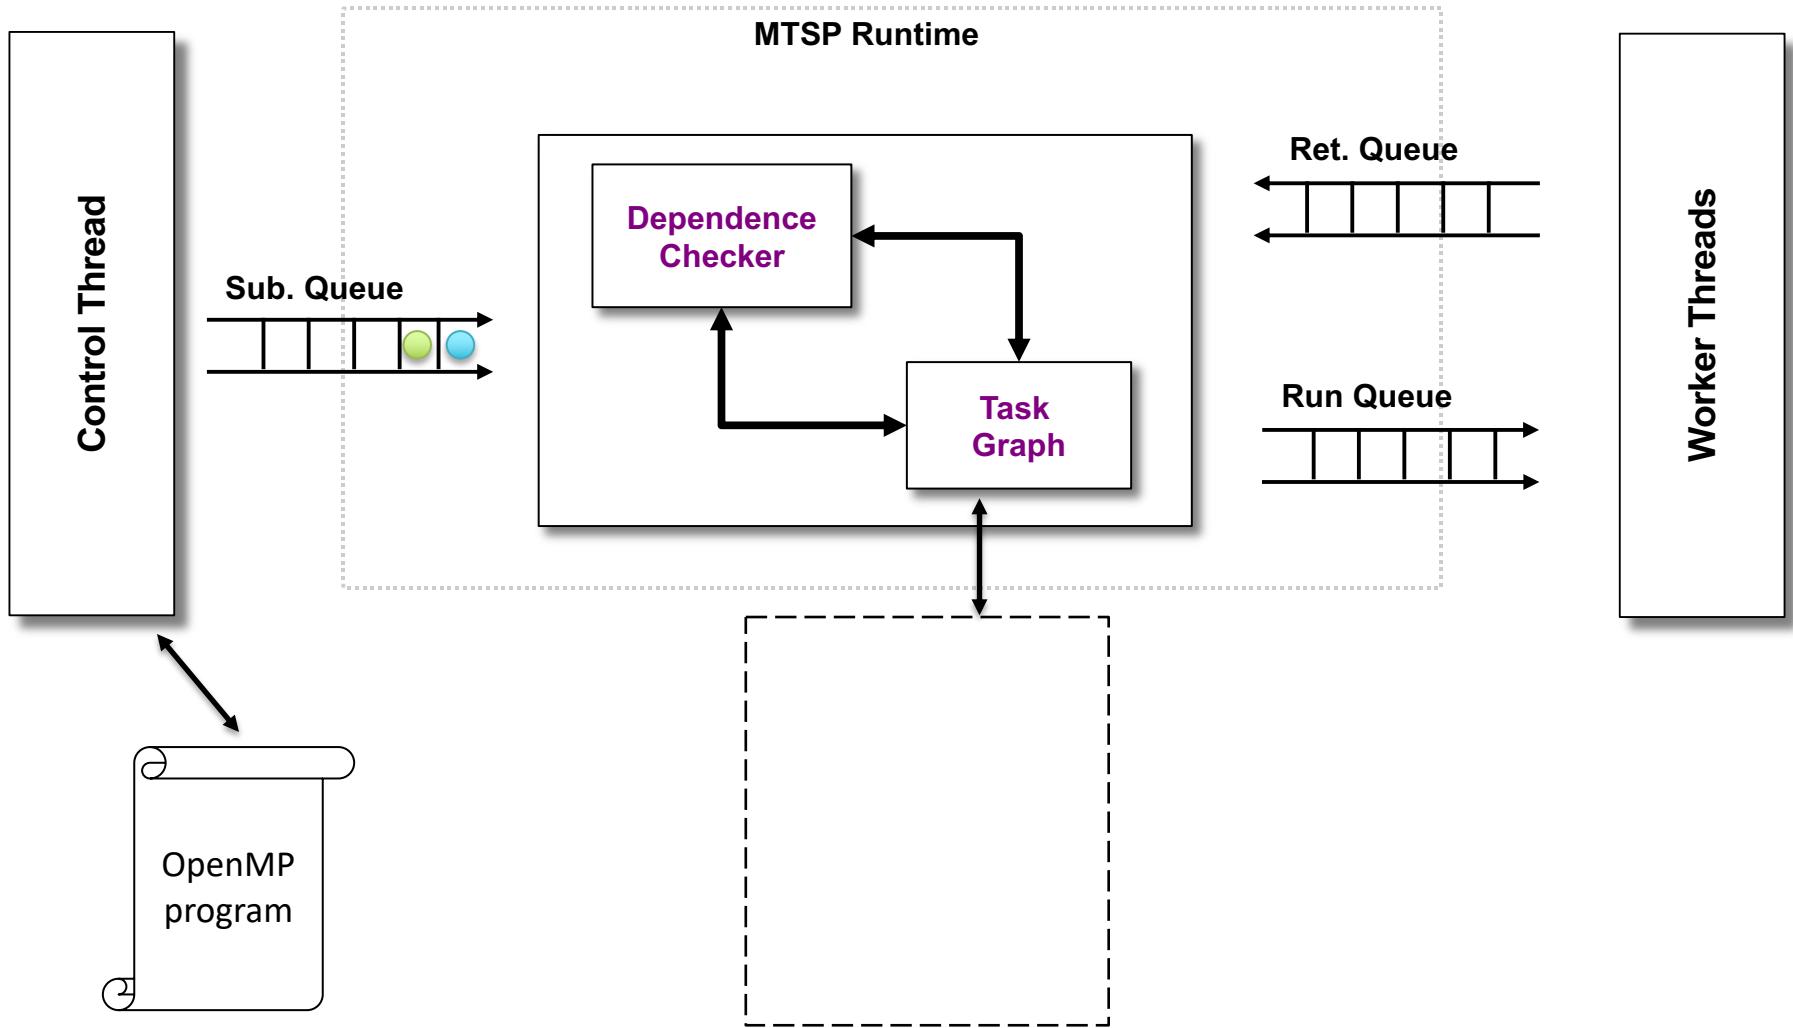
# Basic elements of a task scheduling system



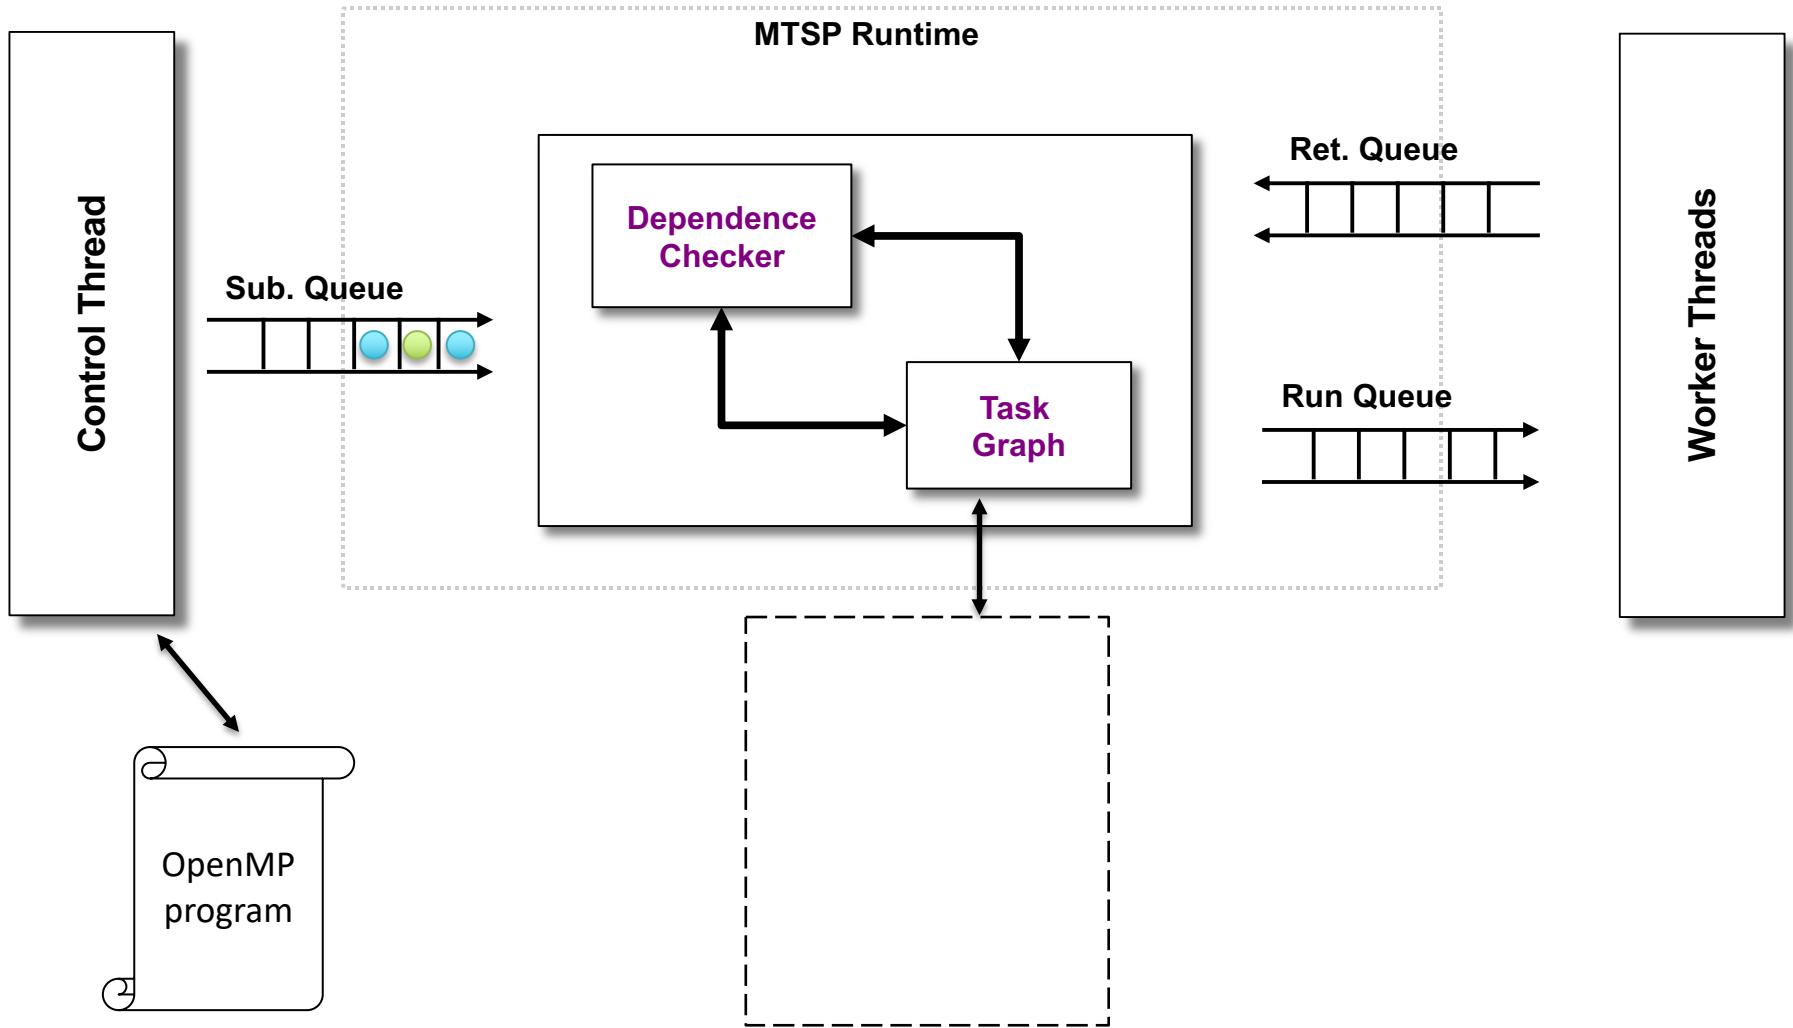
# Basic elements of a task scheduling system



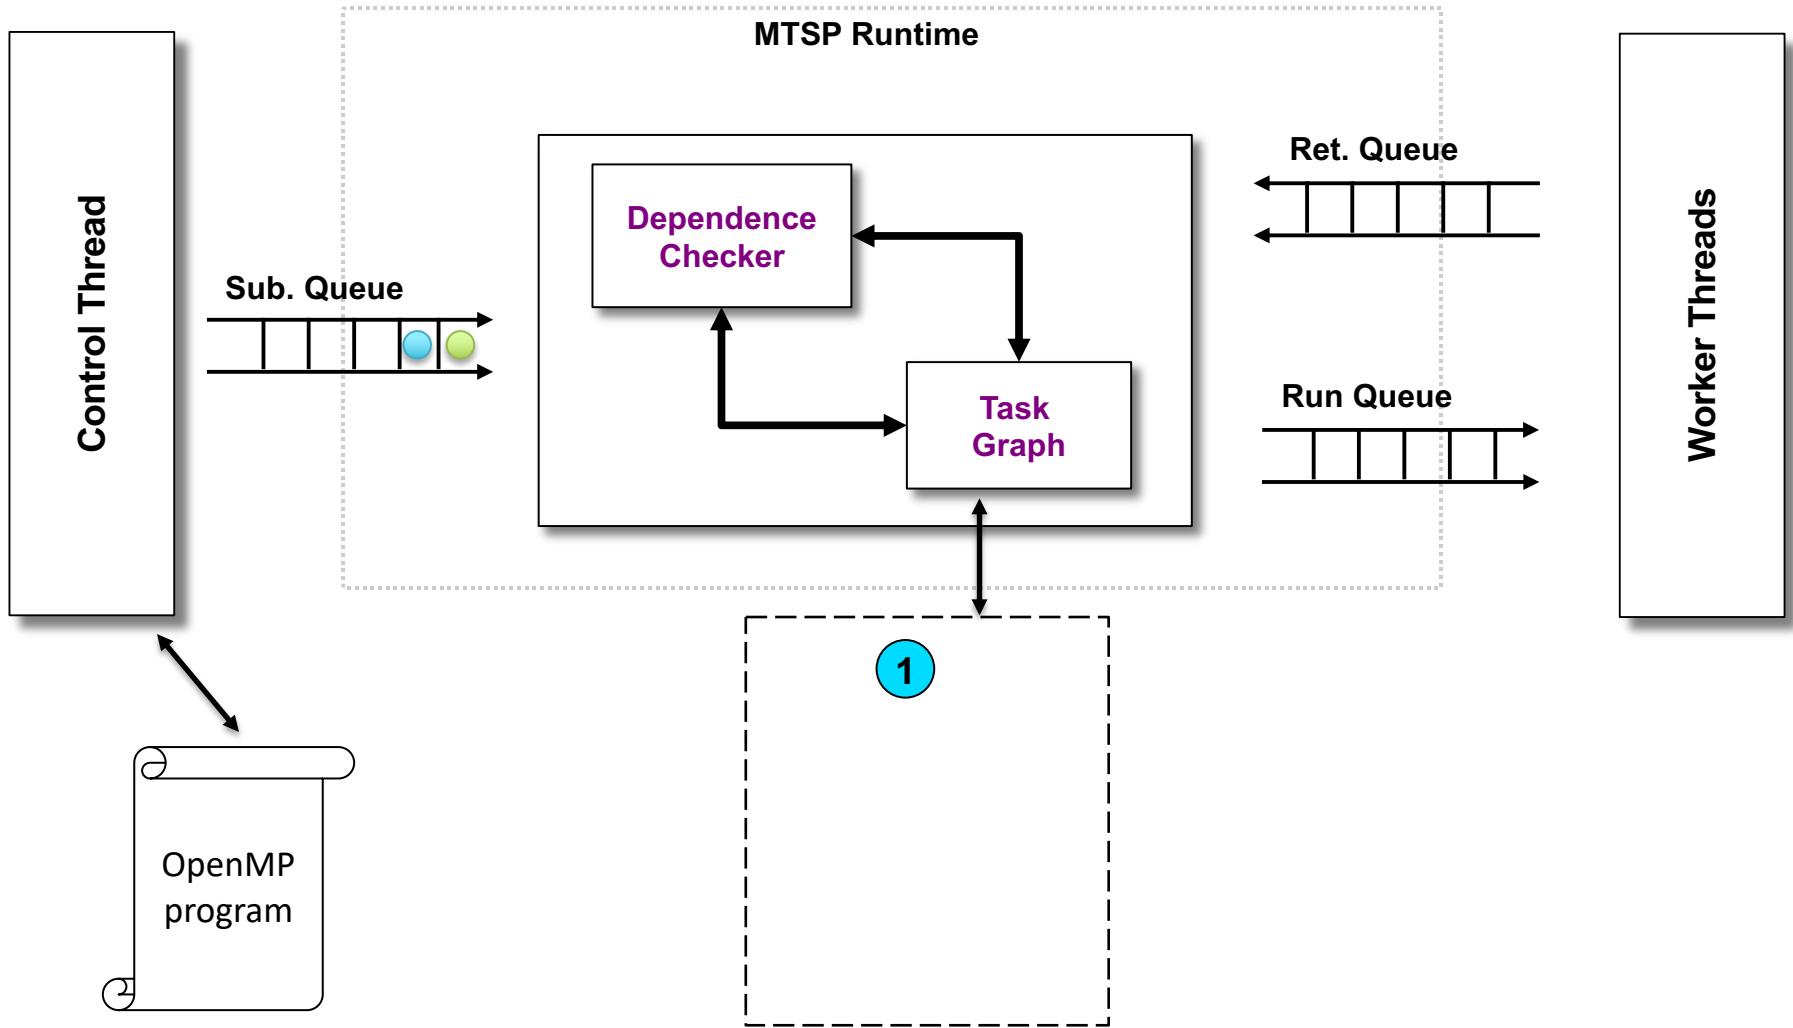
# Basic elements of a task scheduling system



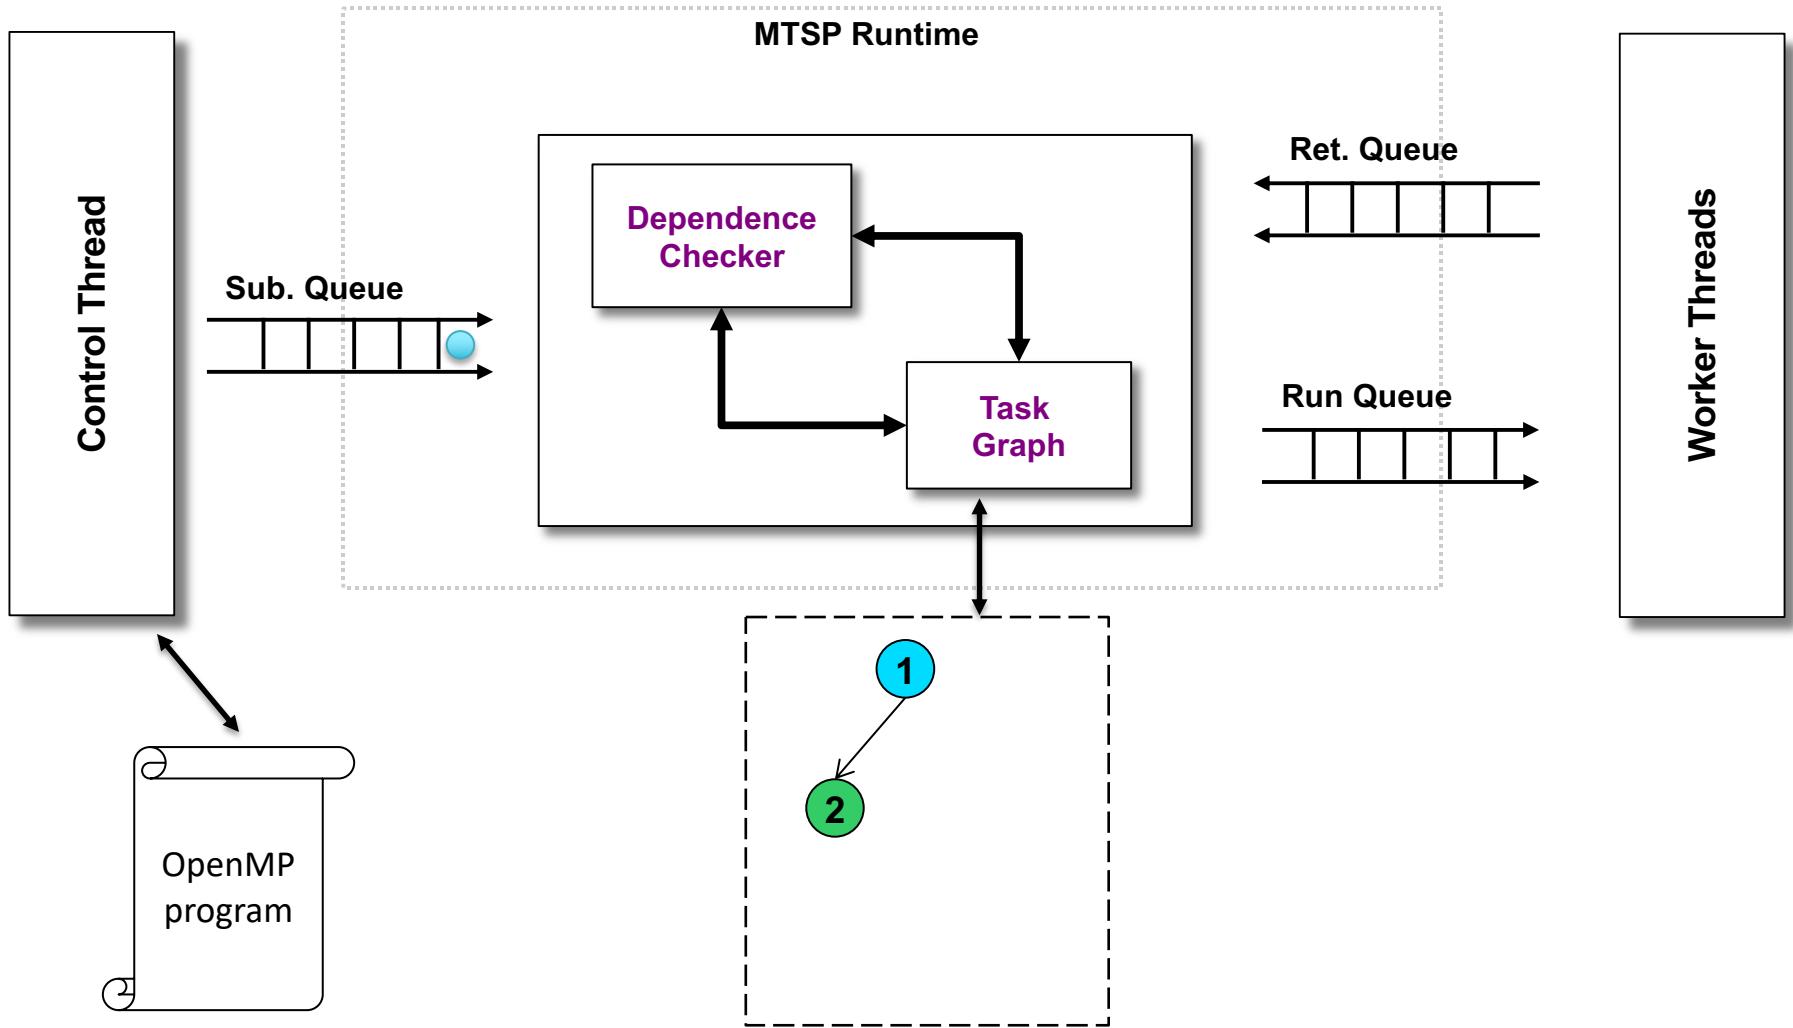
# Basic elements of a task scheduling system



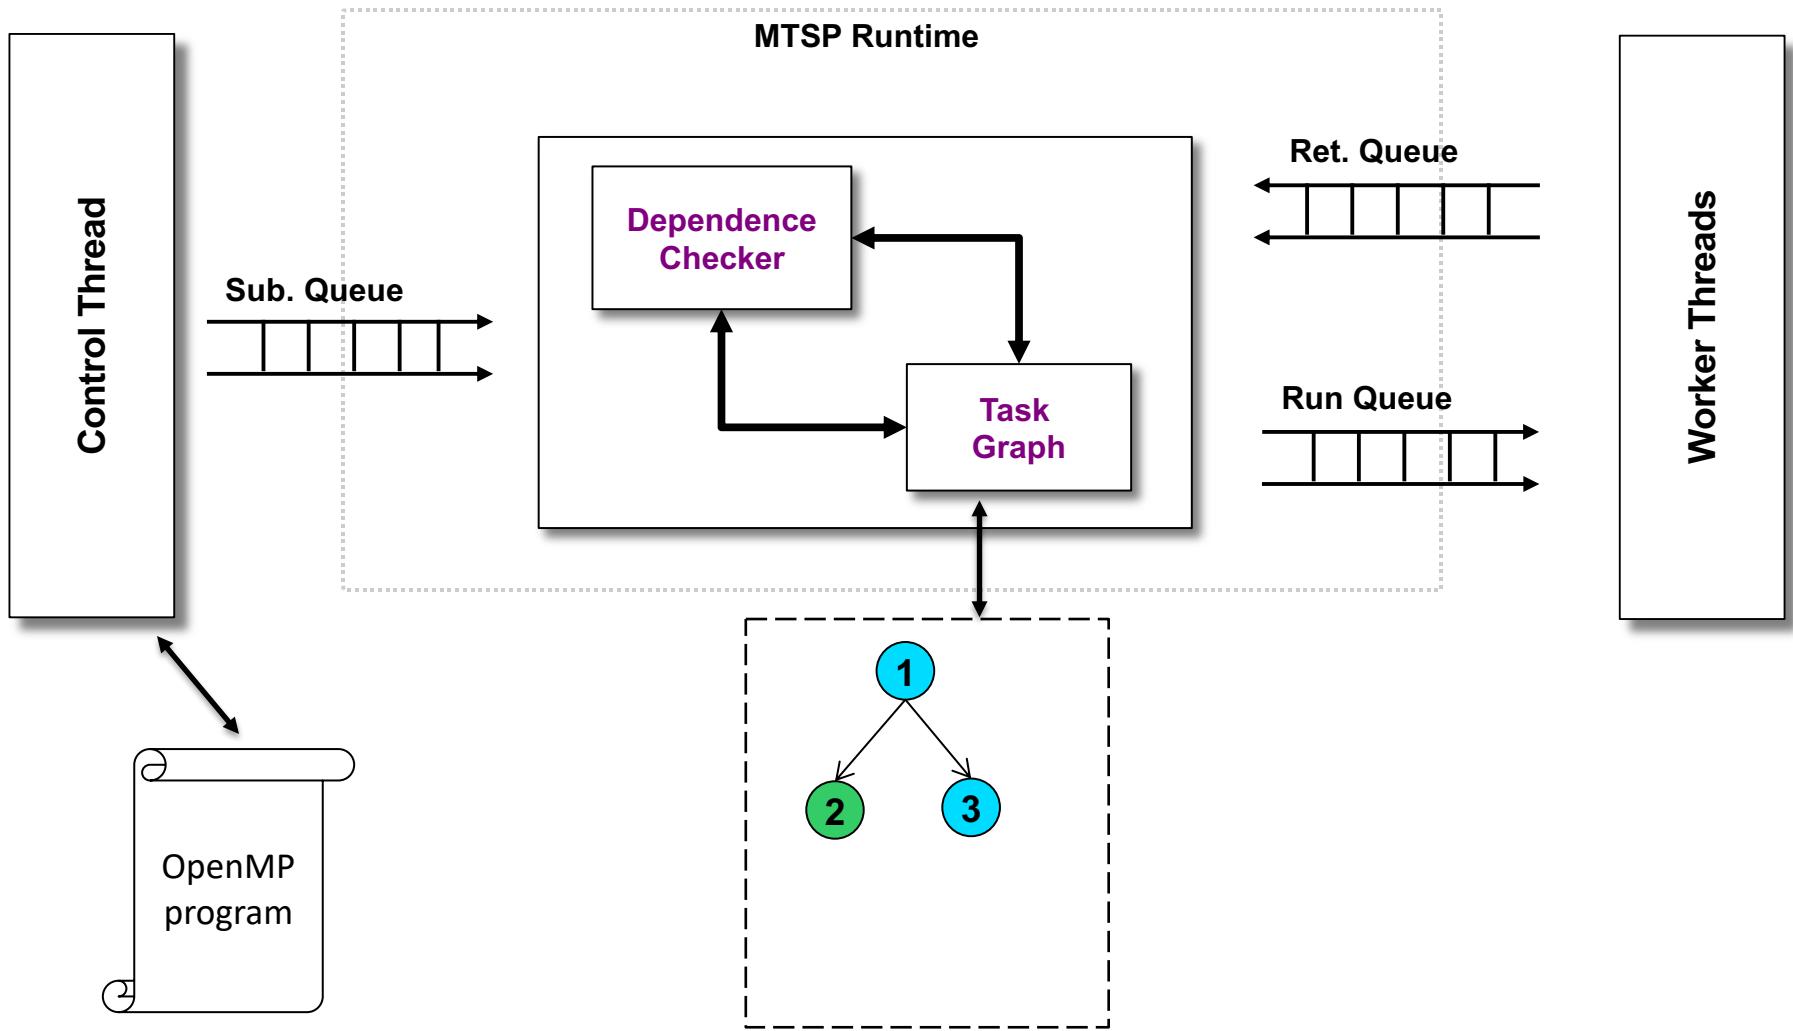
# Basic elements of a task scheduling system



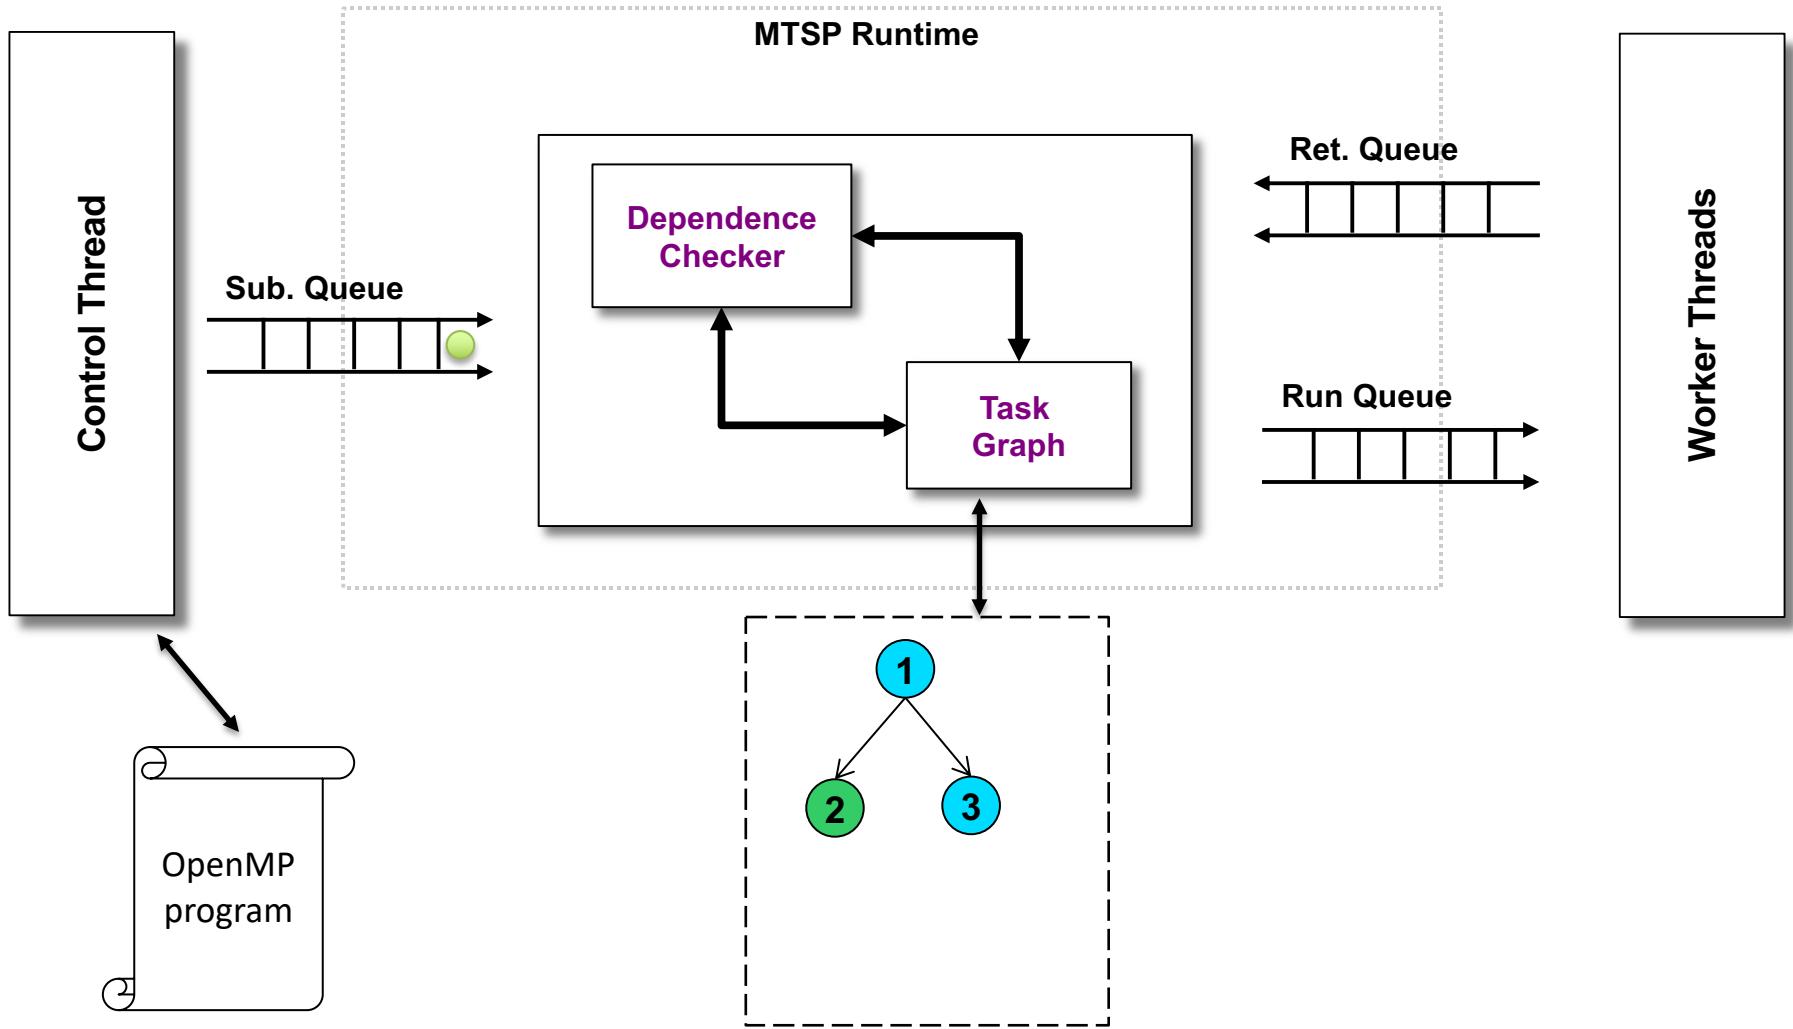
# Basic elements of a task scheduling system



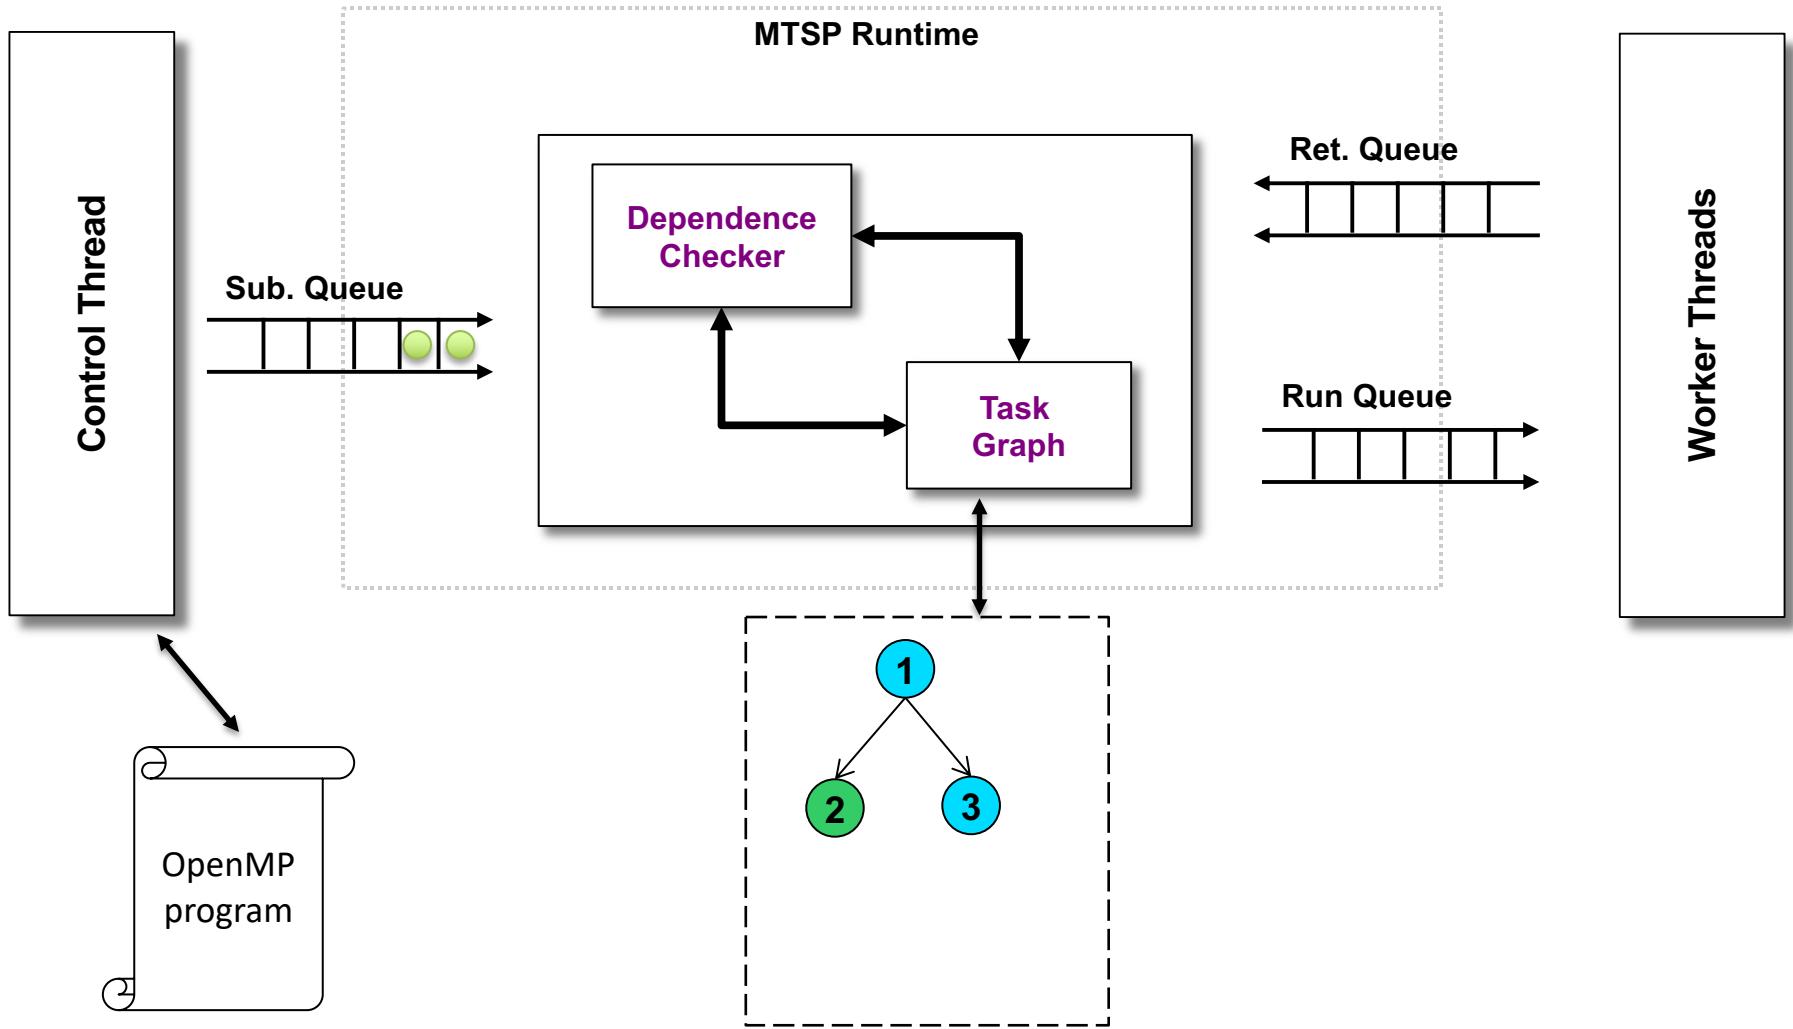
# Basic elements of a task scheduling system



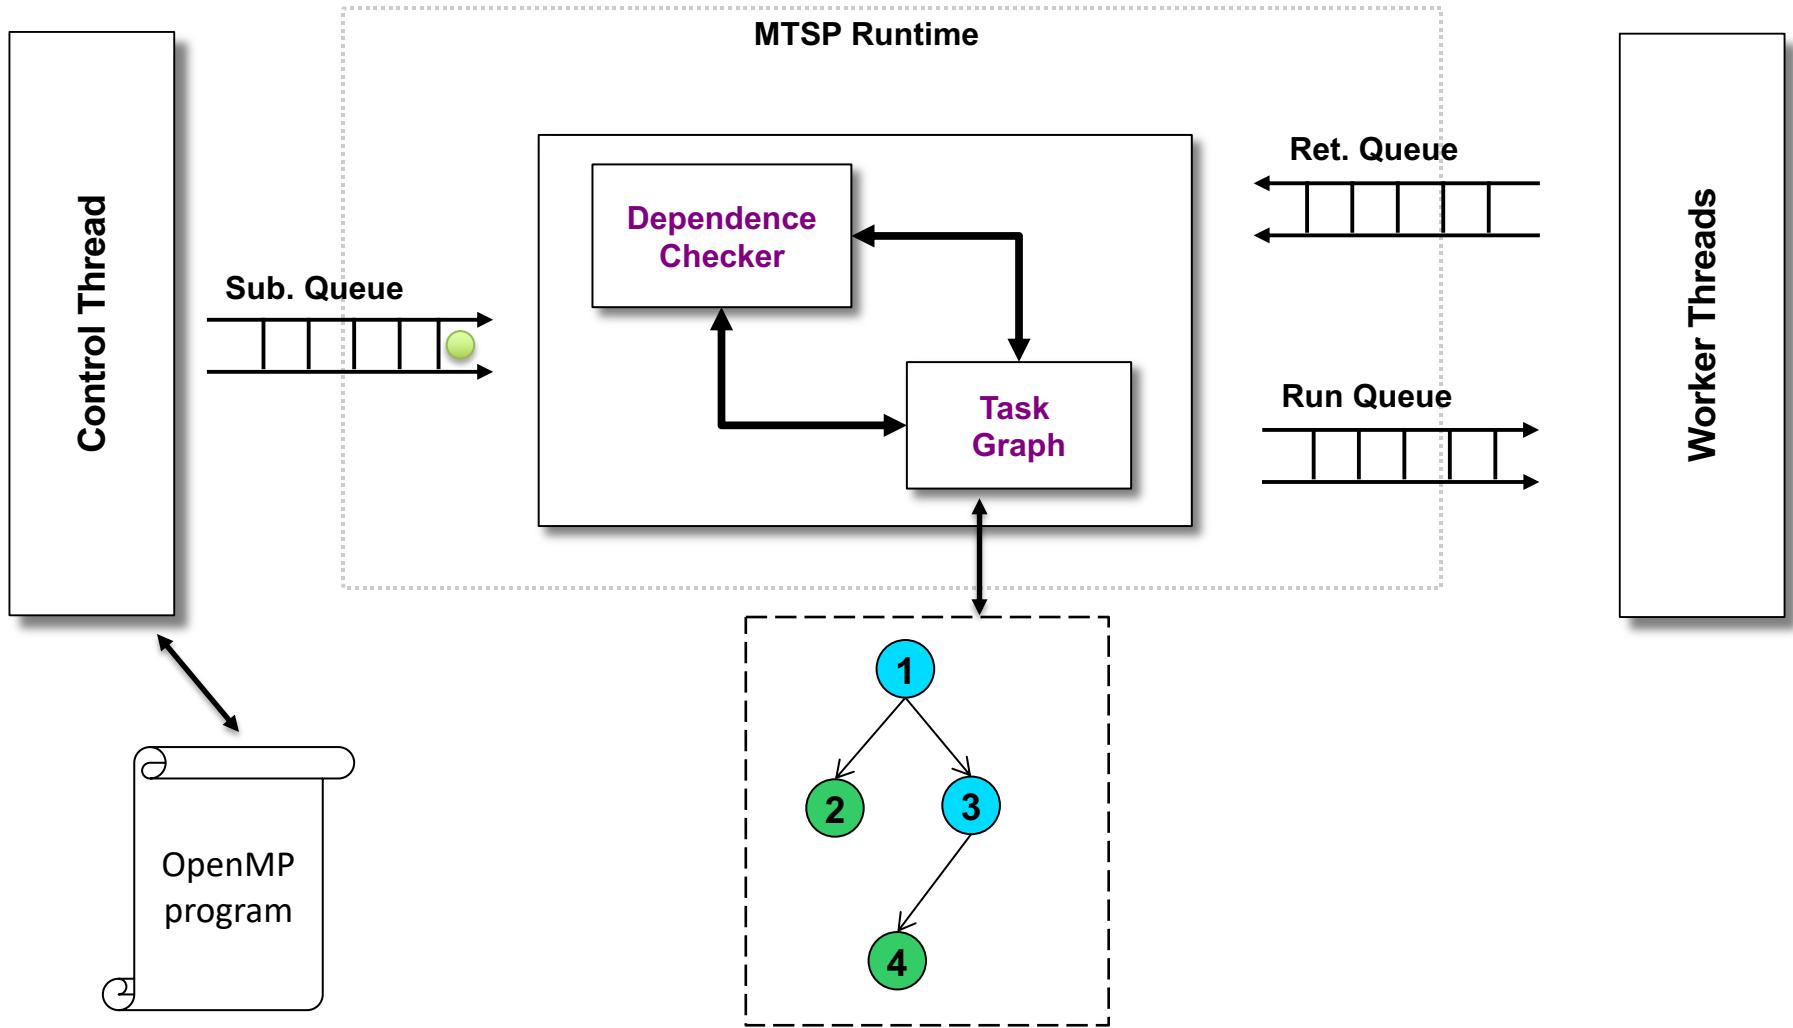
# Basic elements of a task scheduling system



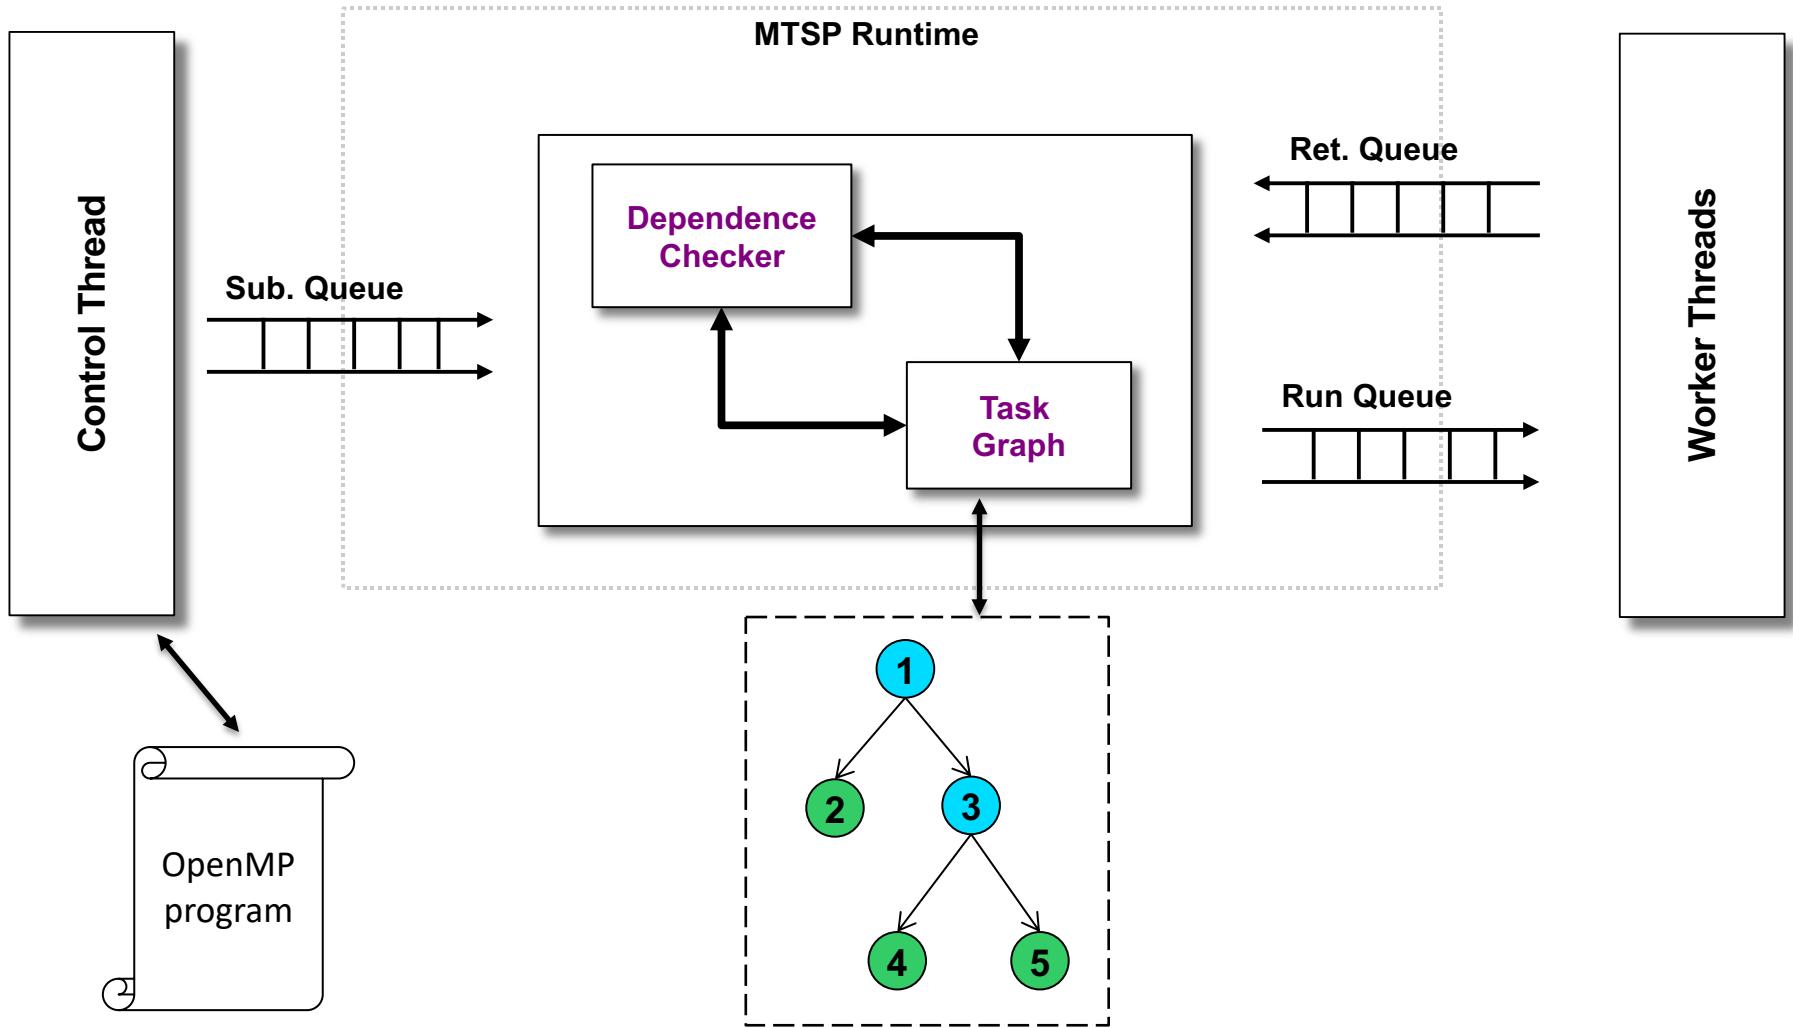
# Basic elements of a task scheduling system



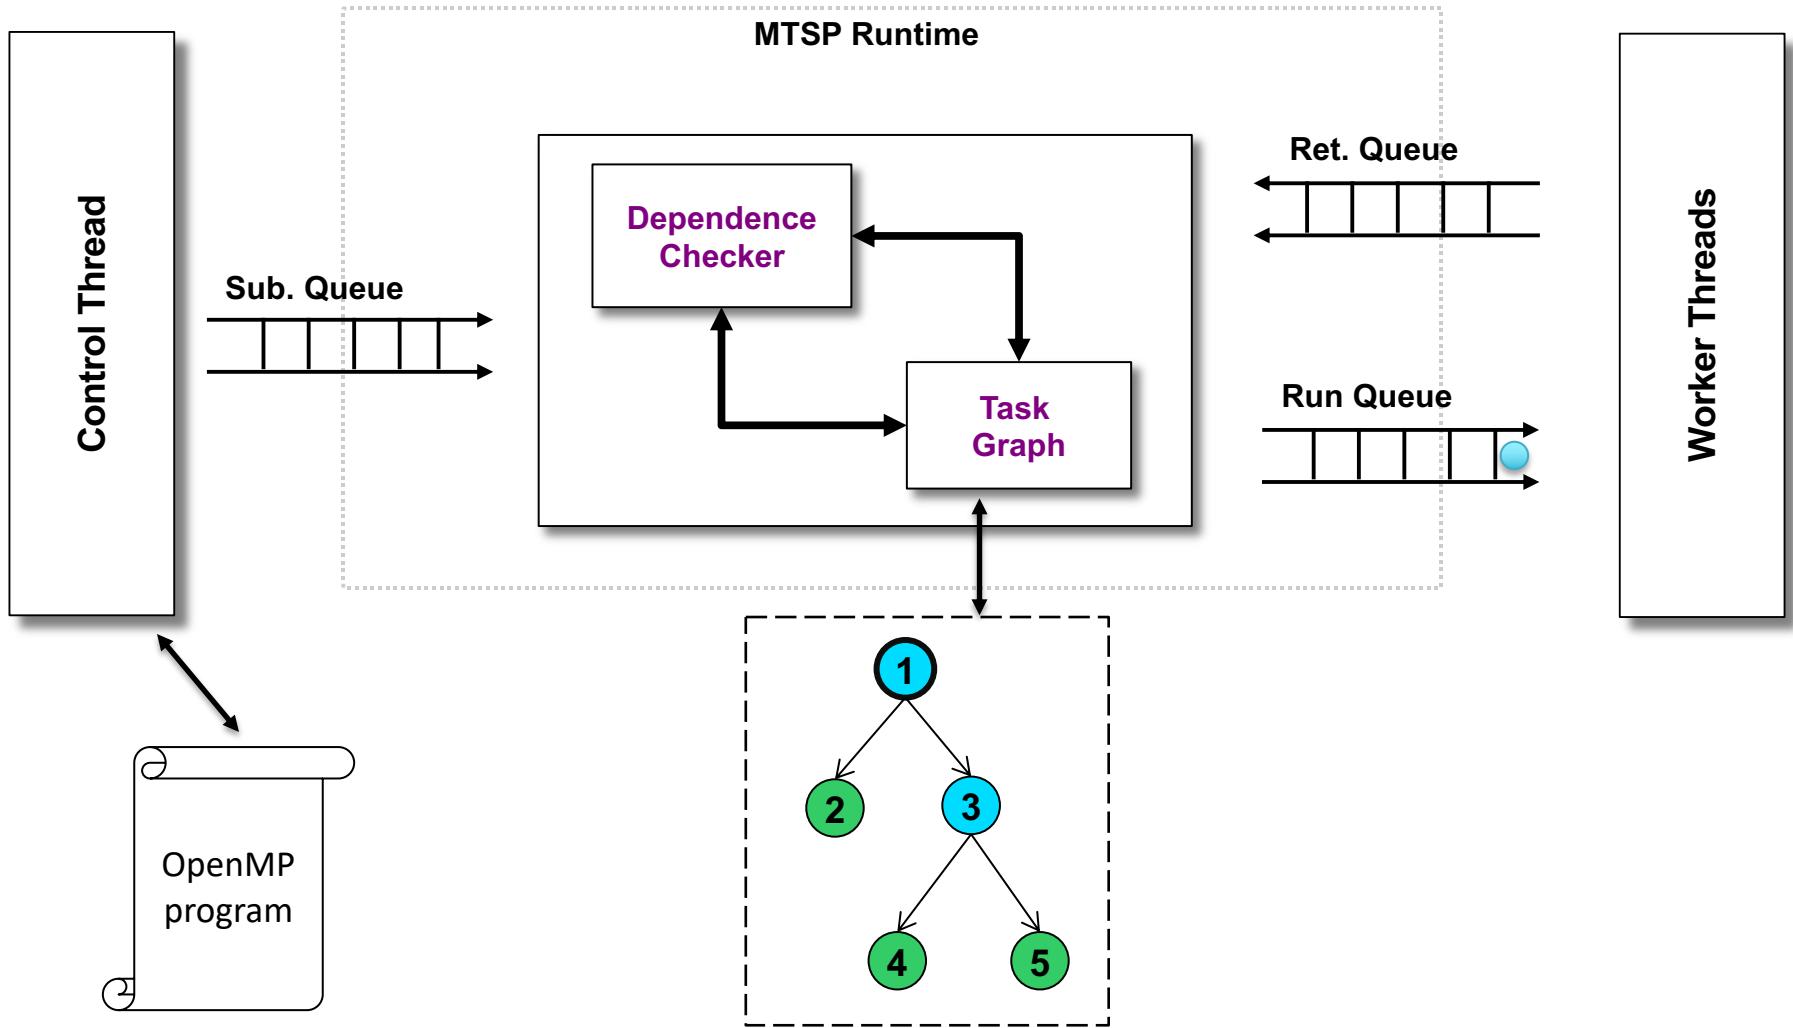
# Basic elements of a task scheduling system



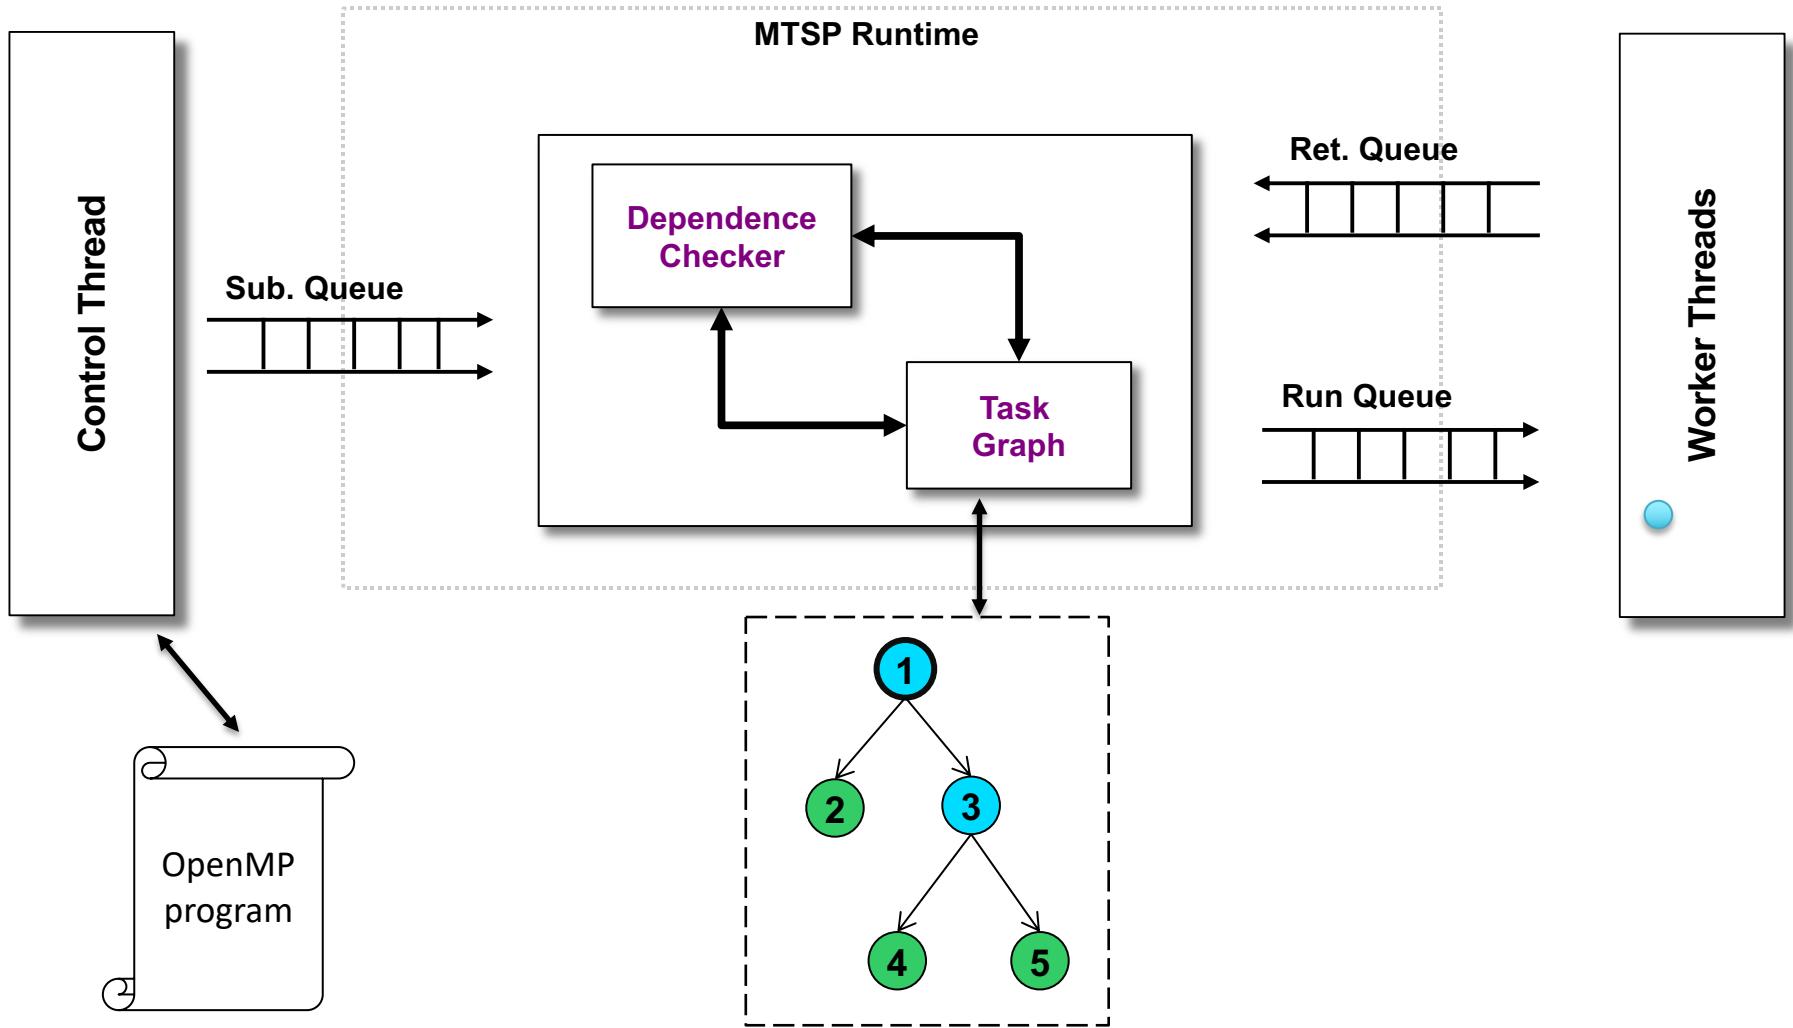
# Basic elements of a task scheduling system



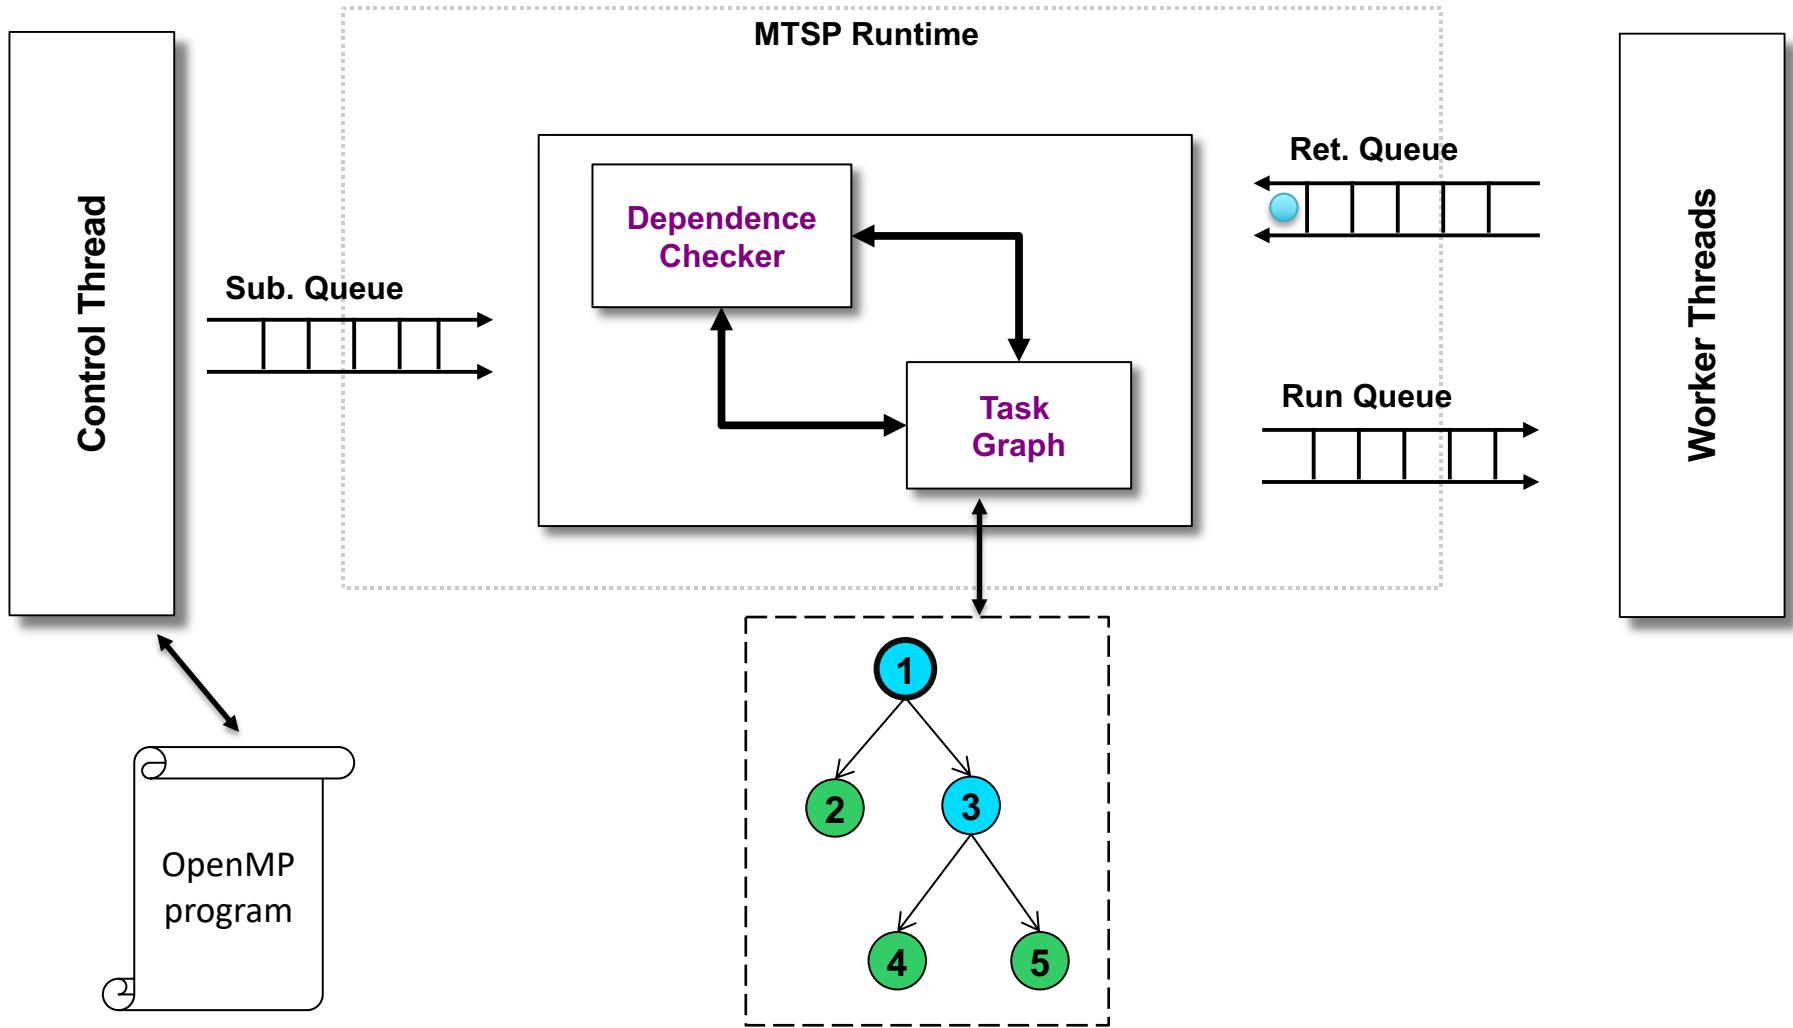
# Basic elements of a task scheduling system



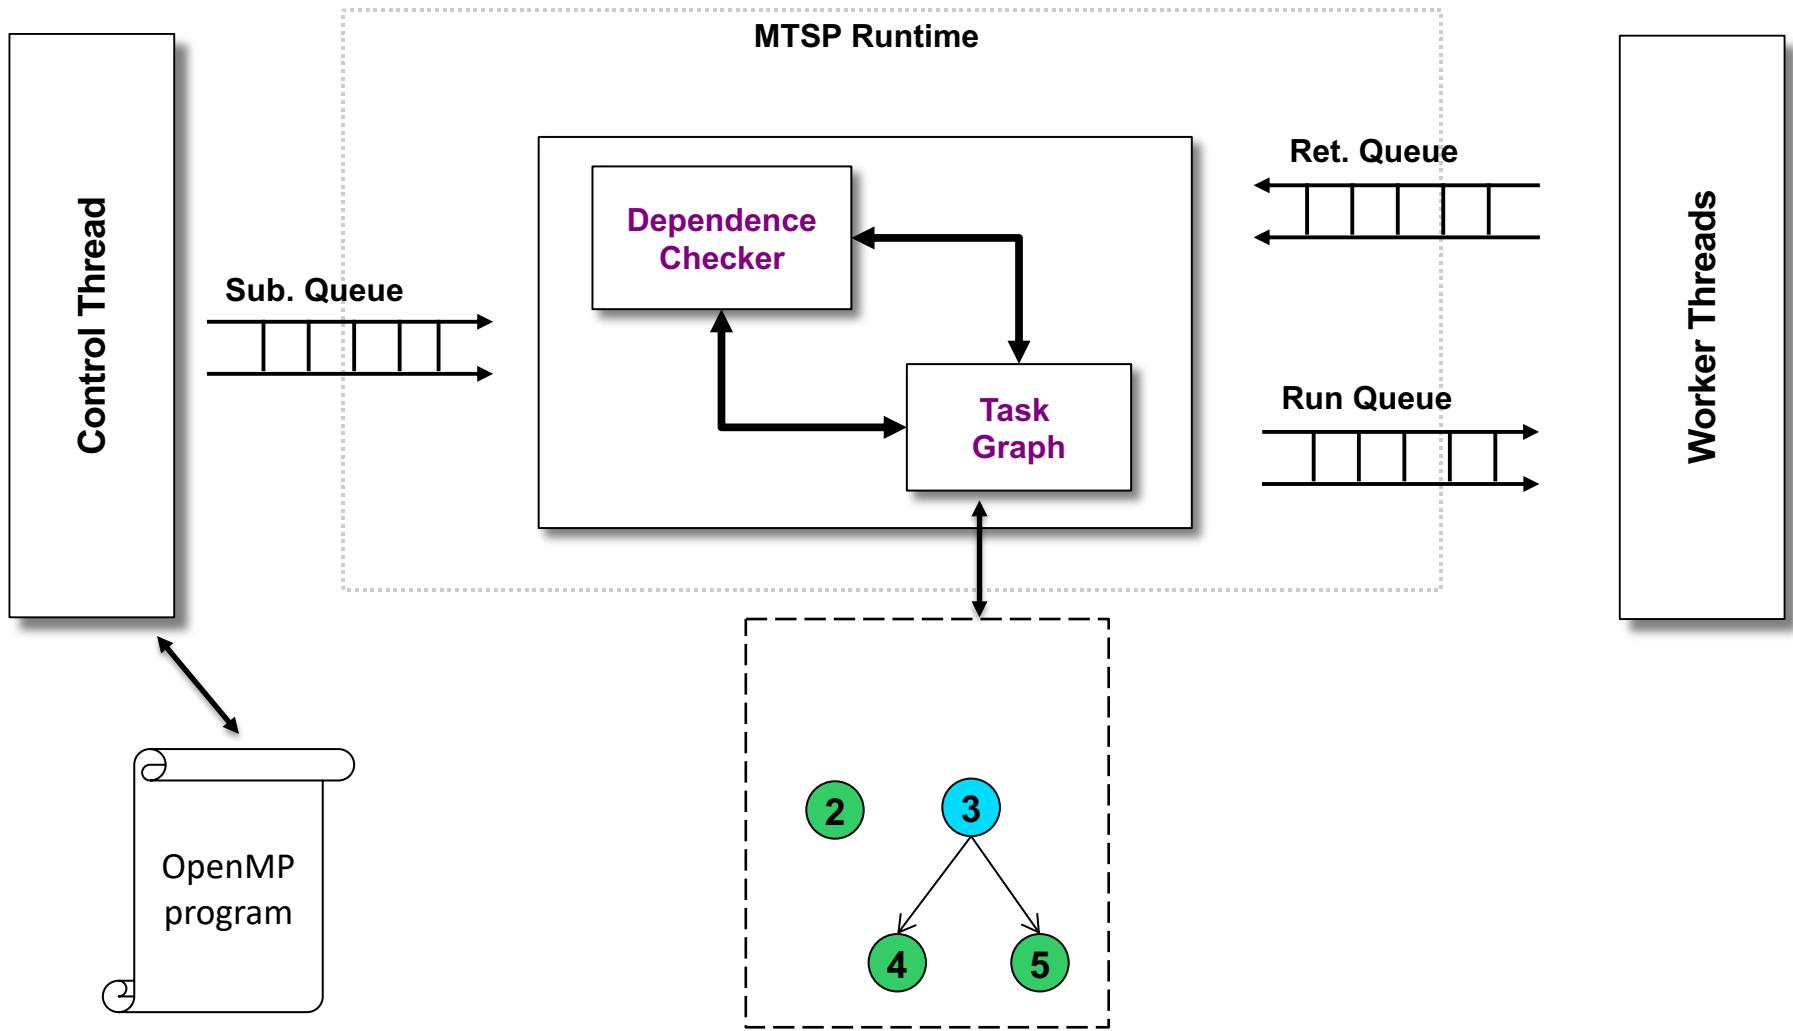
# Basic elements of a task scheduling system



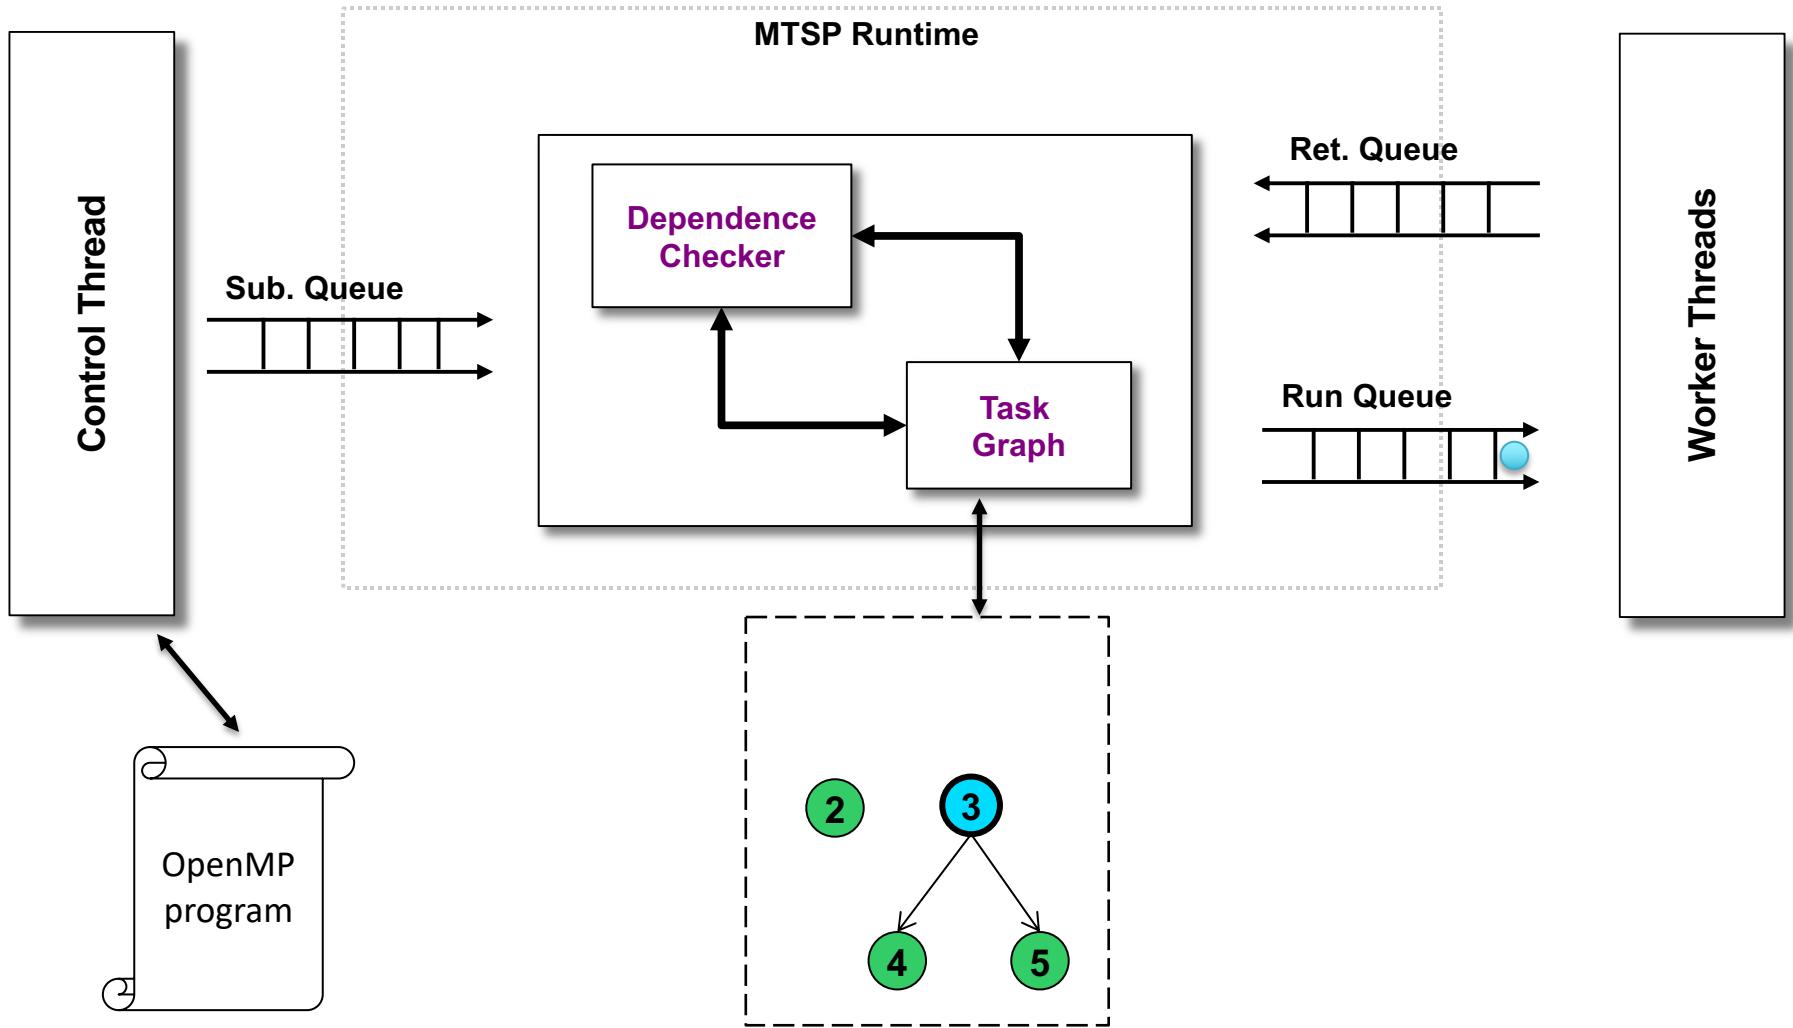
# Basic elements of a task scheduling system



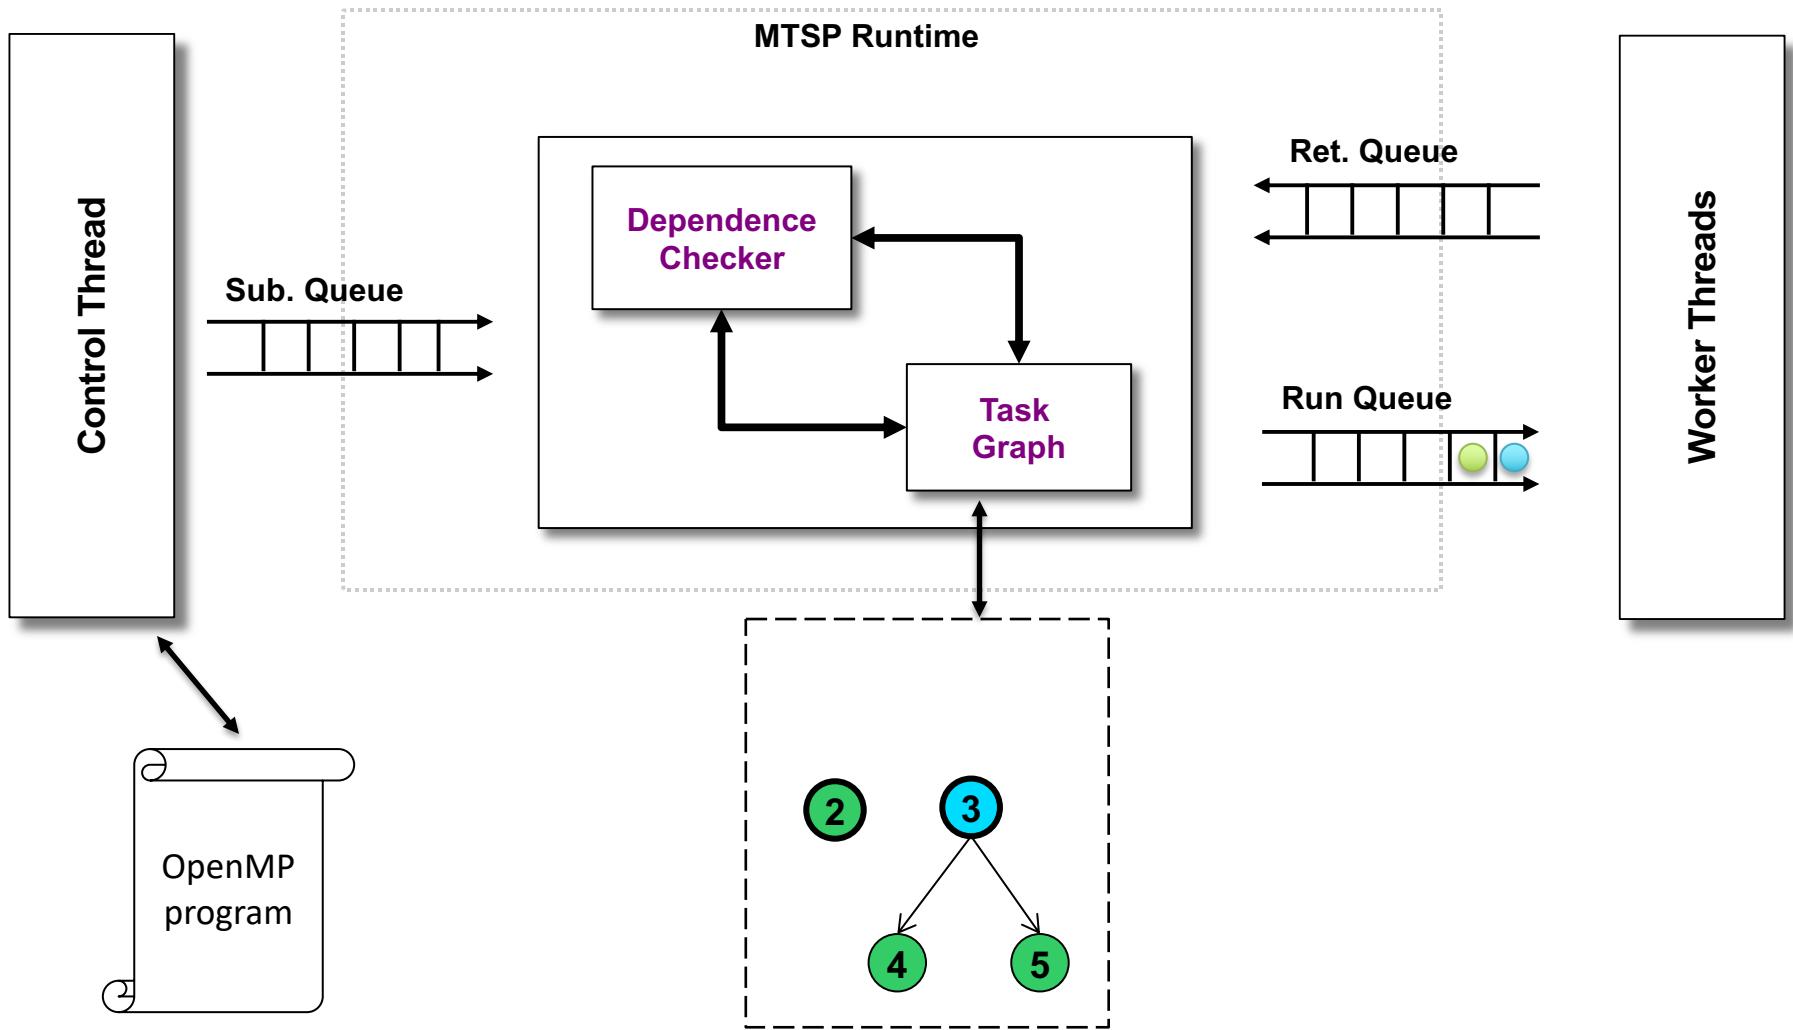
# Basic elements of a task scheduling system



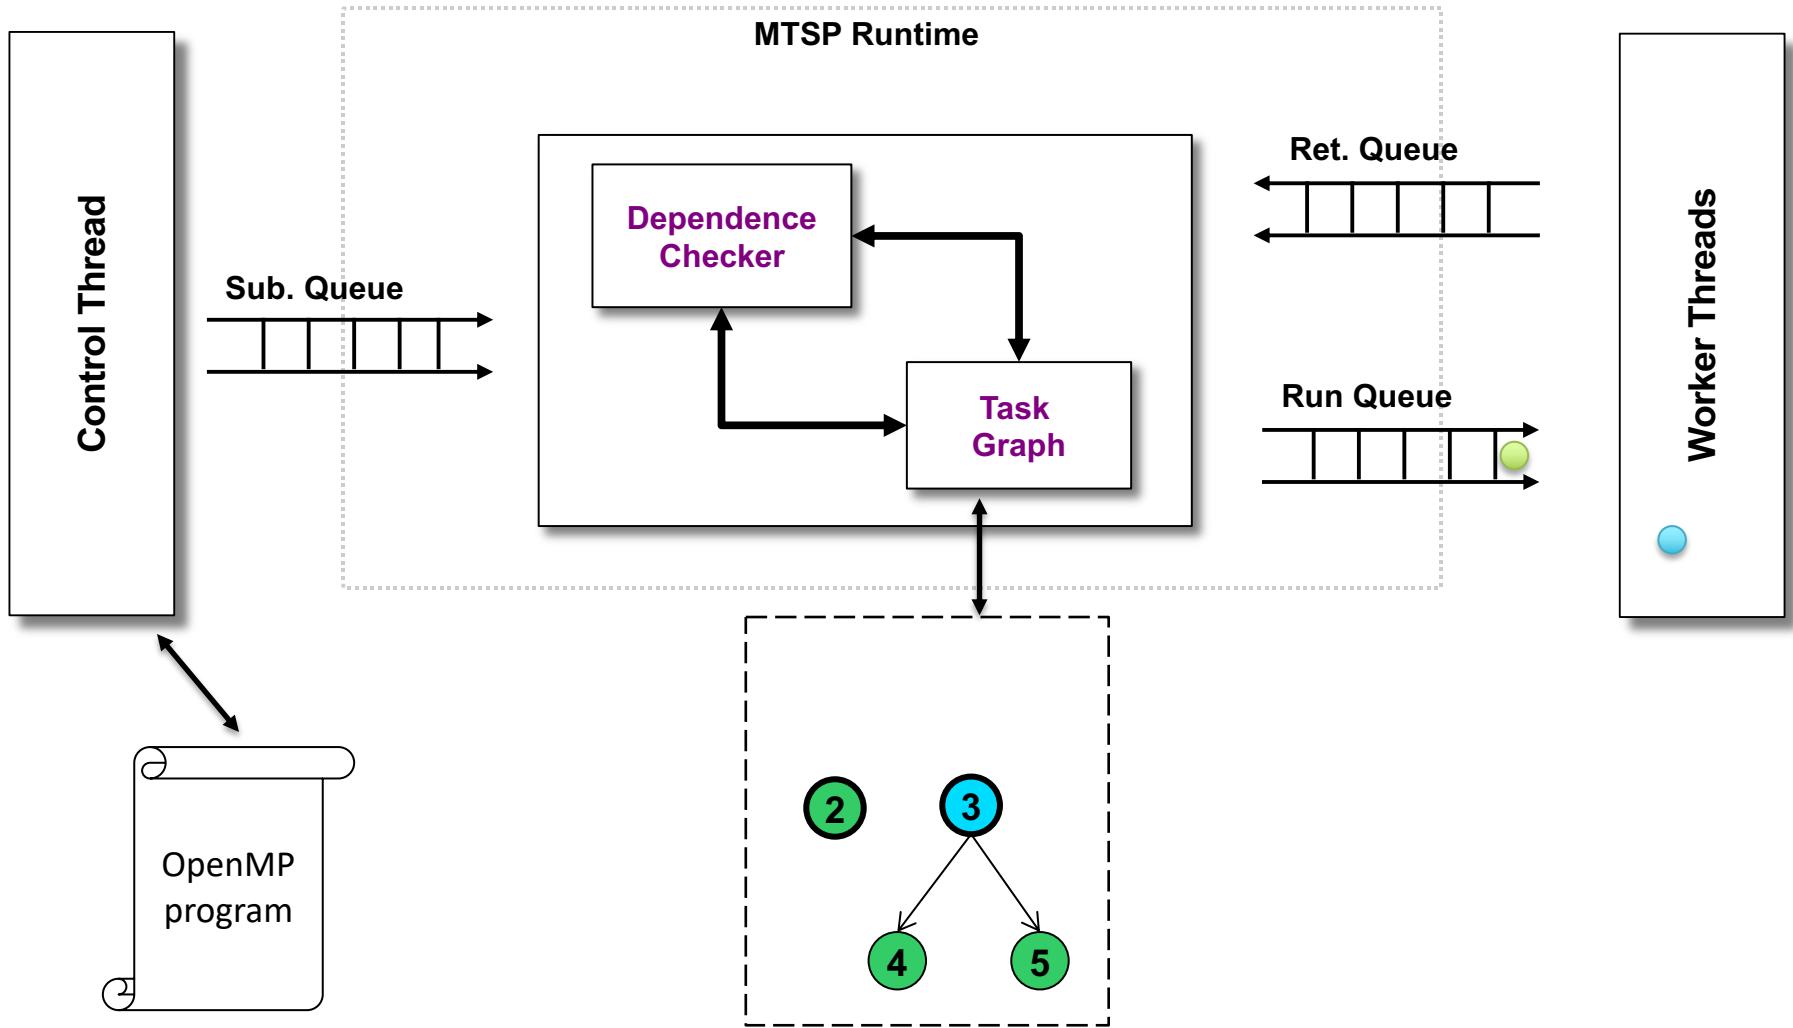
# Basic elements of a task scheduling system



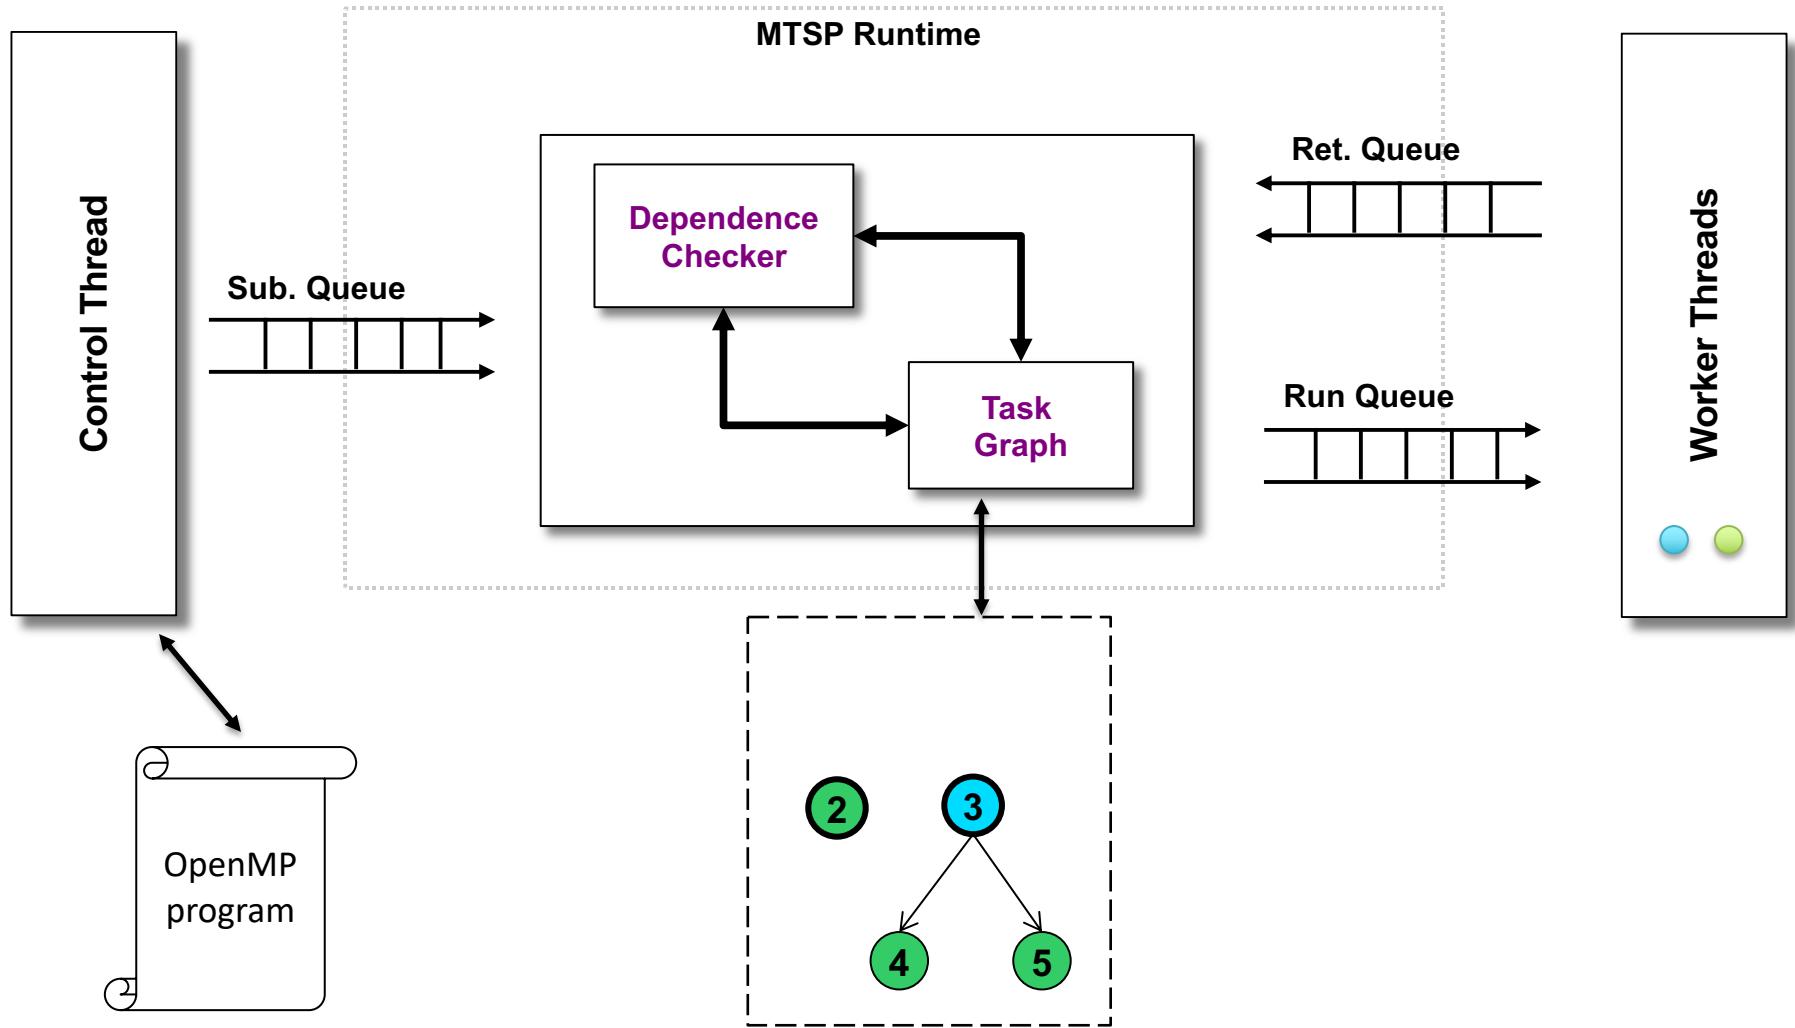
# Basic elements of a task scheduling system



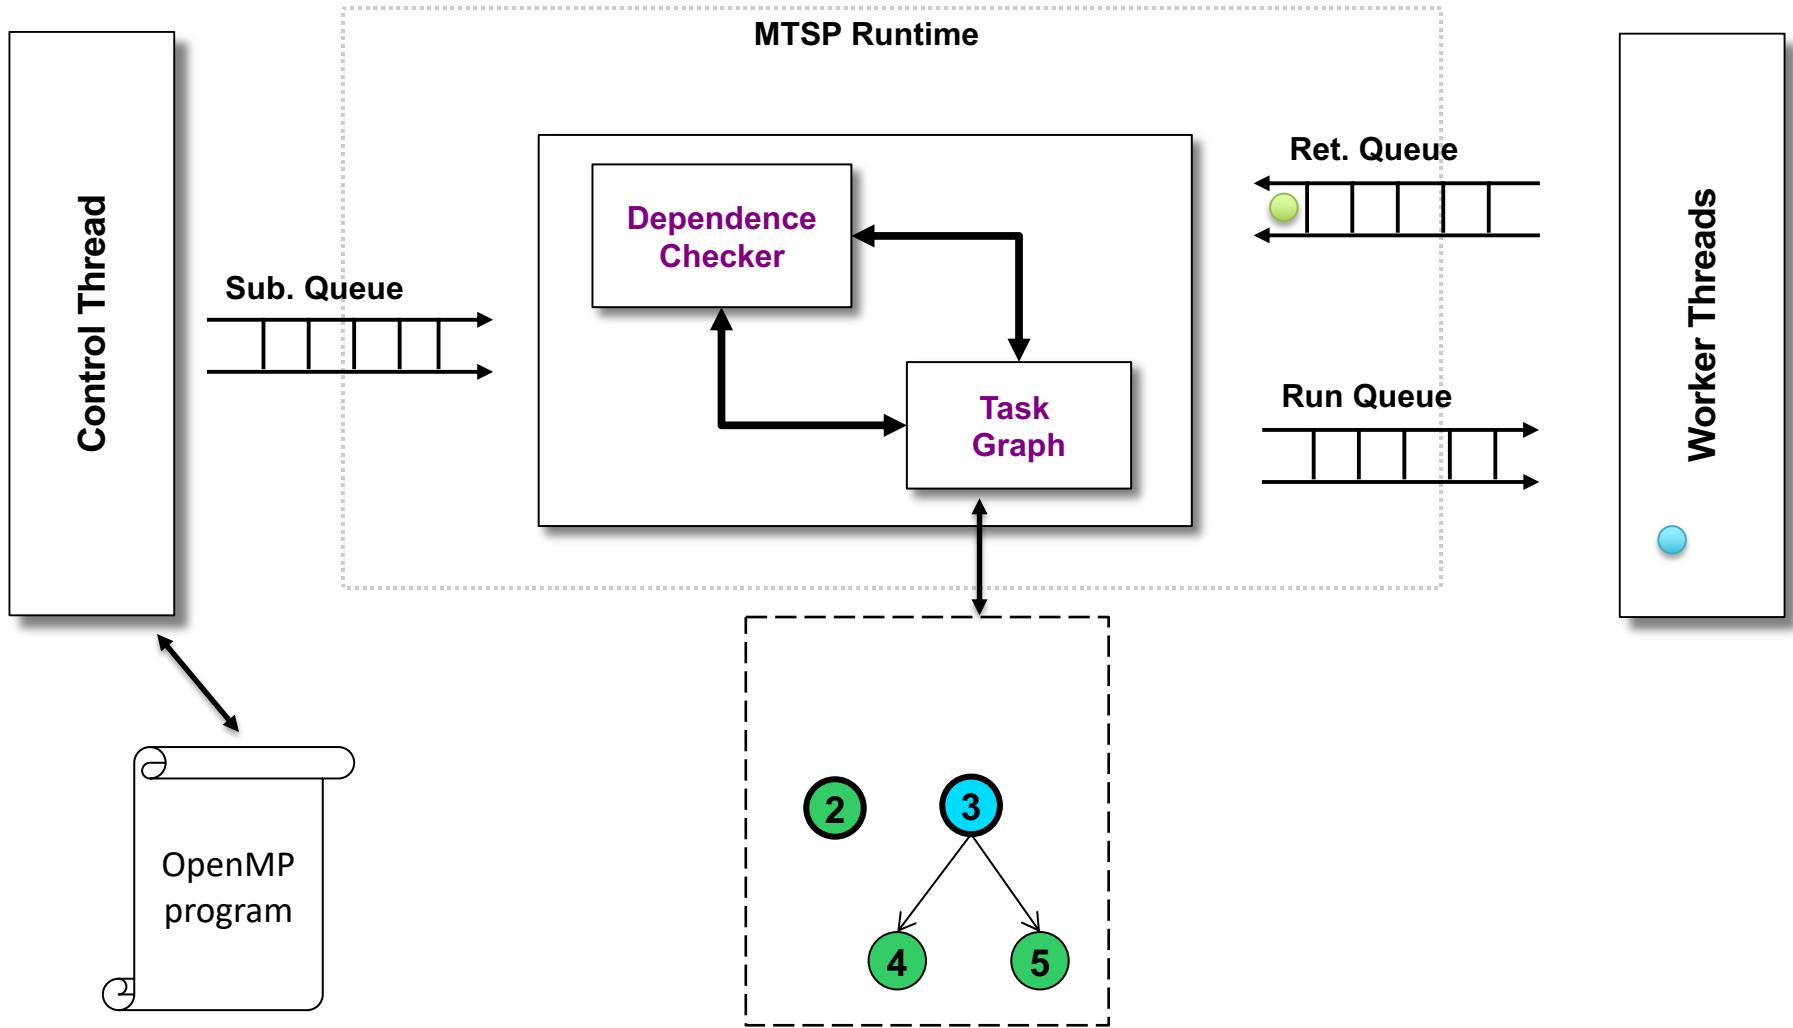
# Basic elements of a task scheduling system



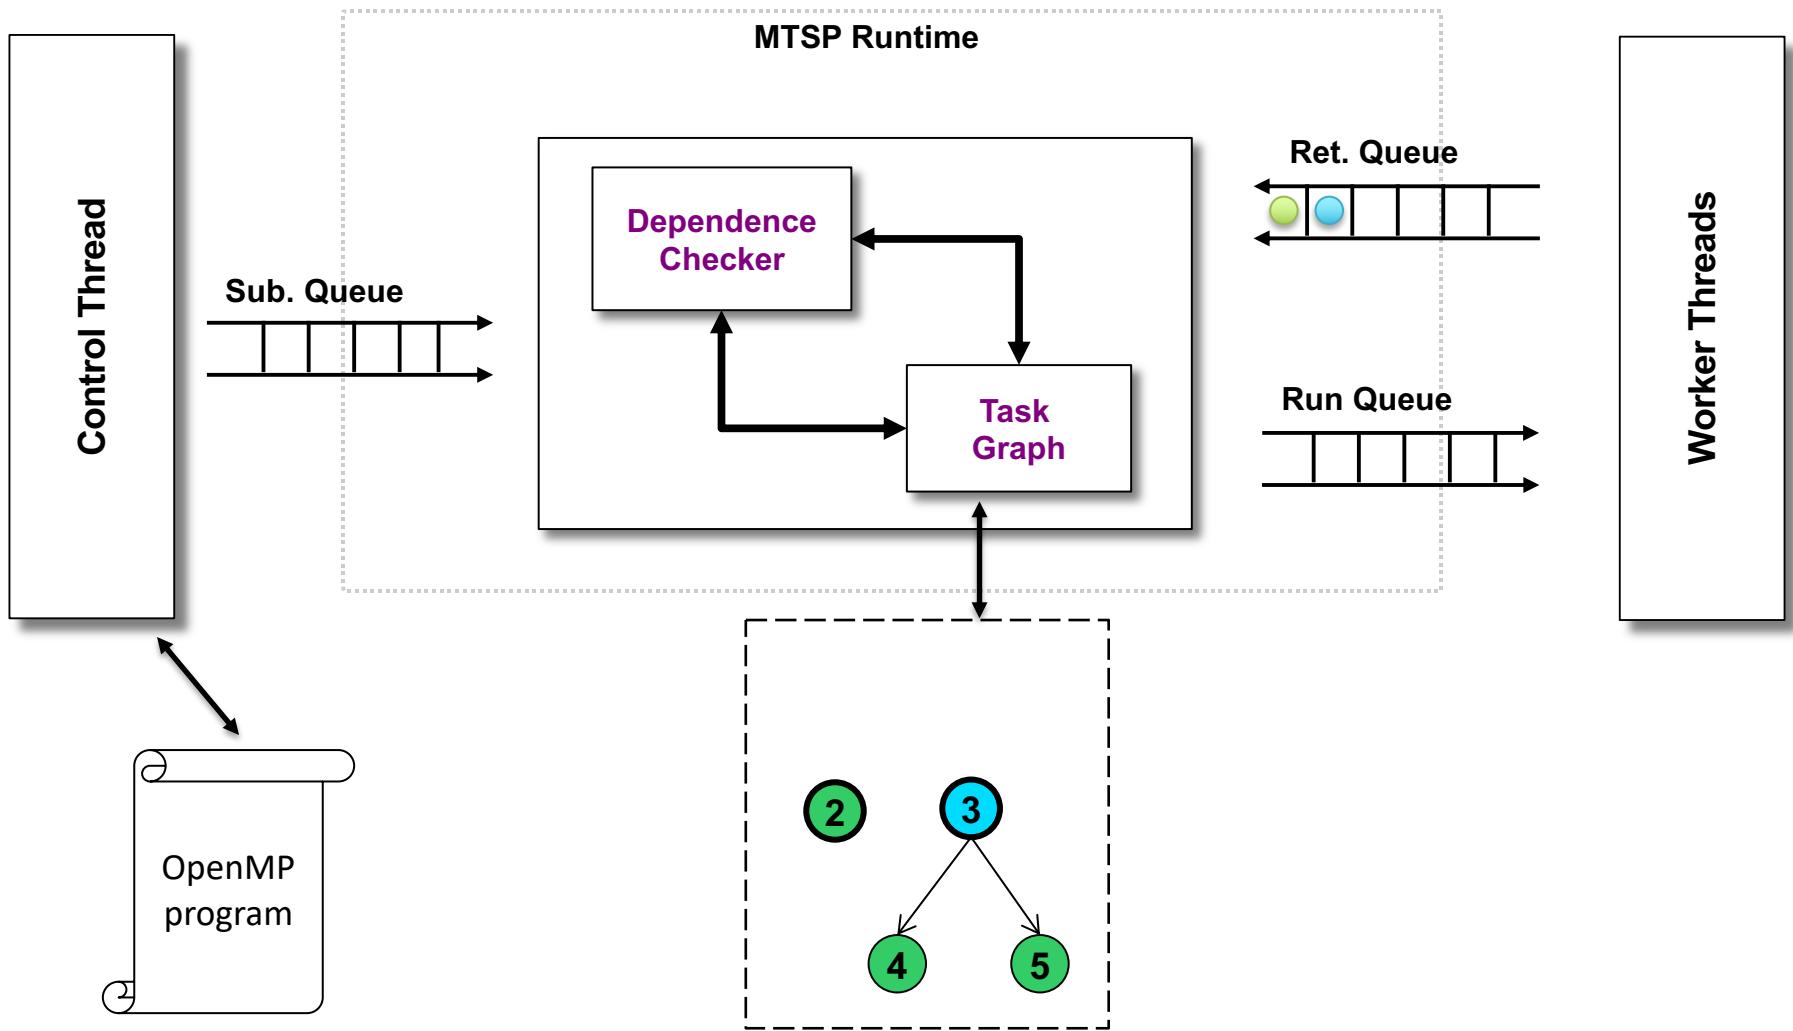
# Basic elements of a task scheduling system



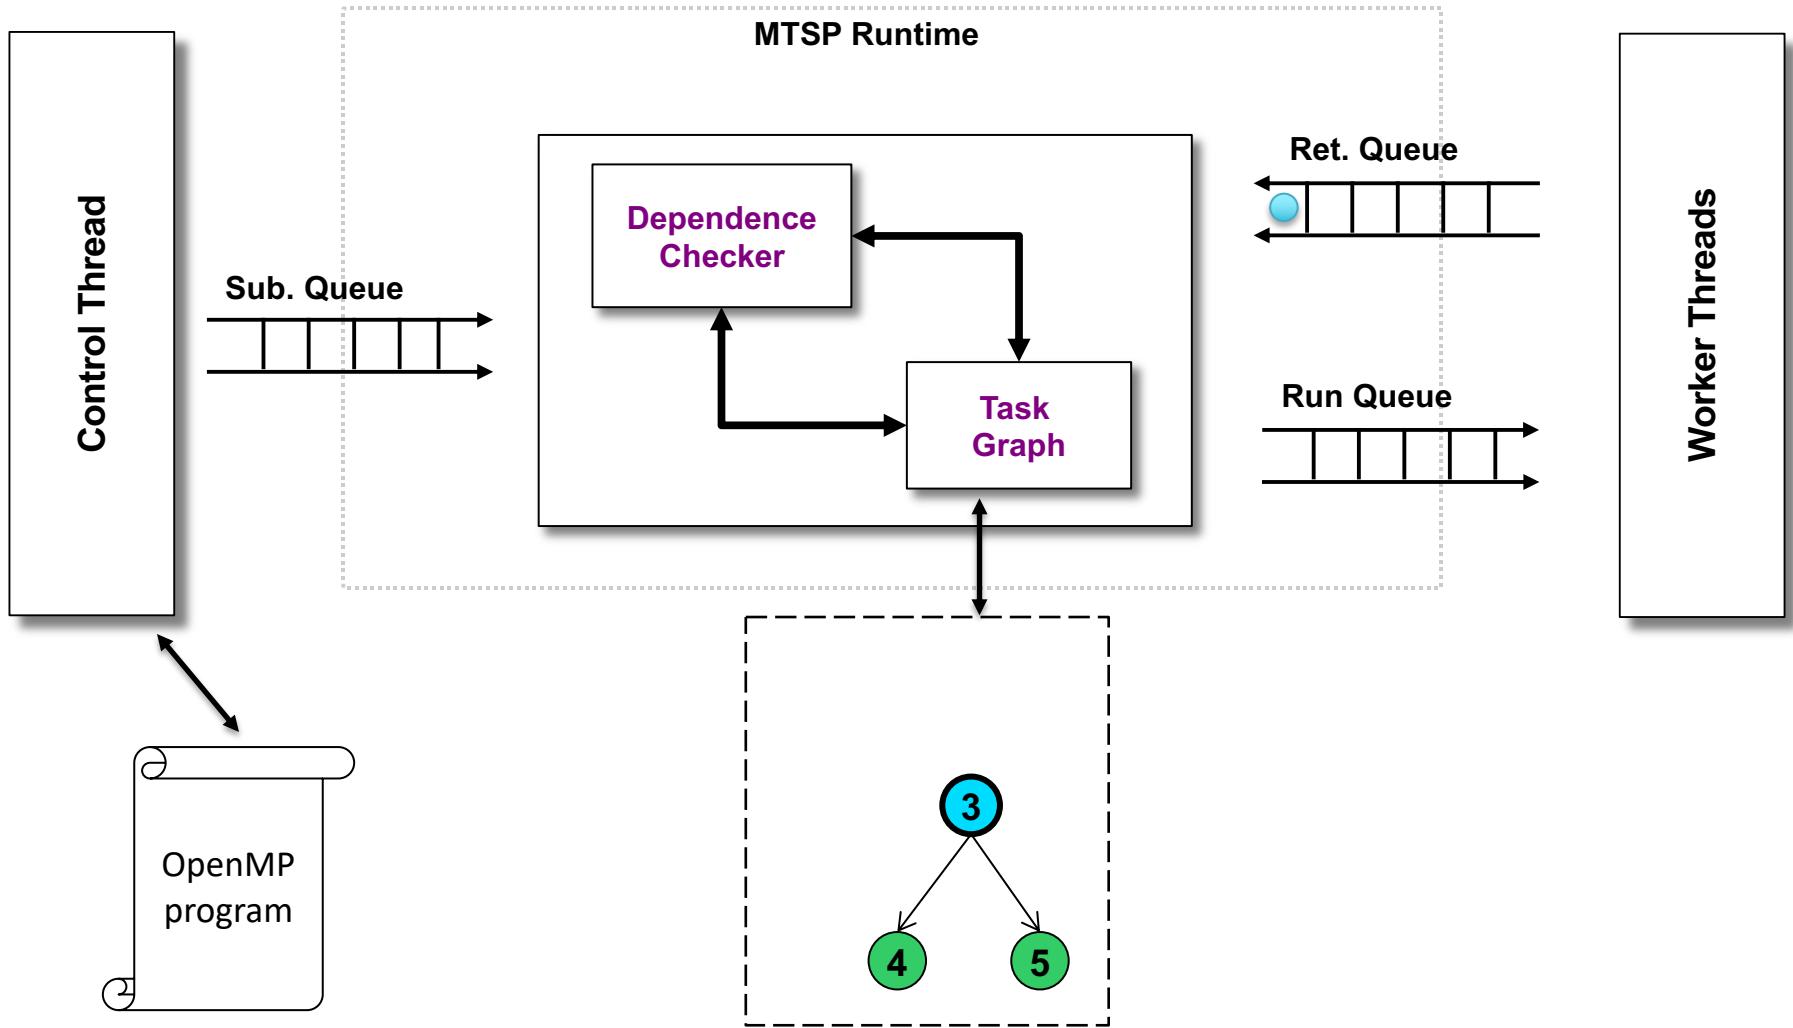
# Basic elements of a task scheduling system



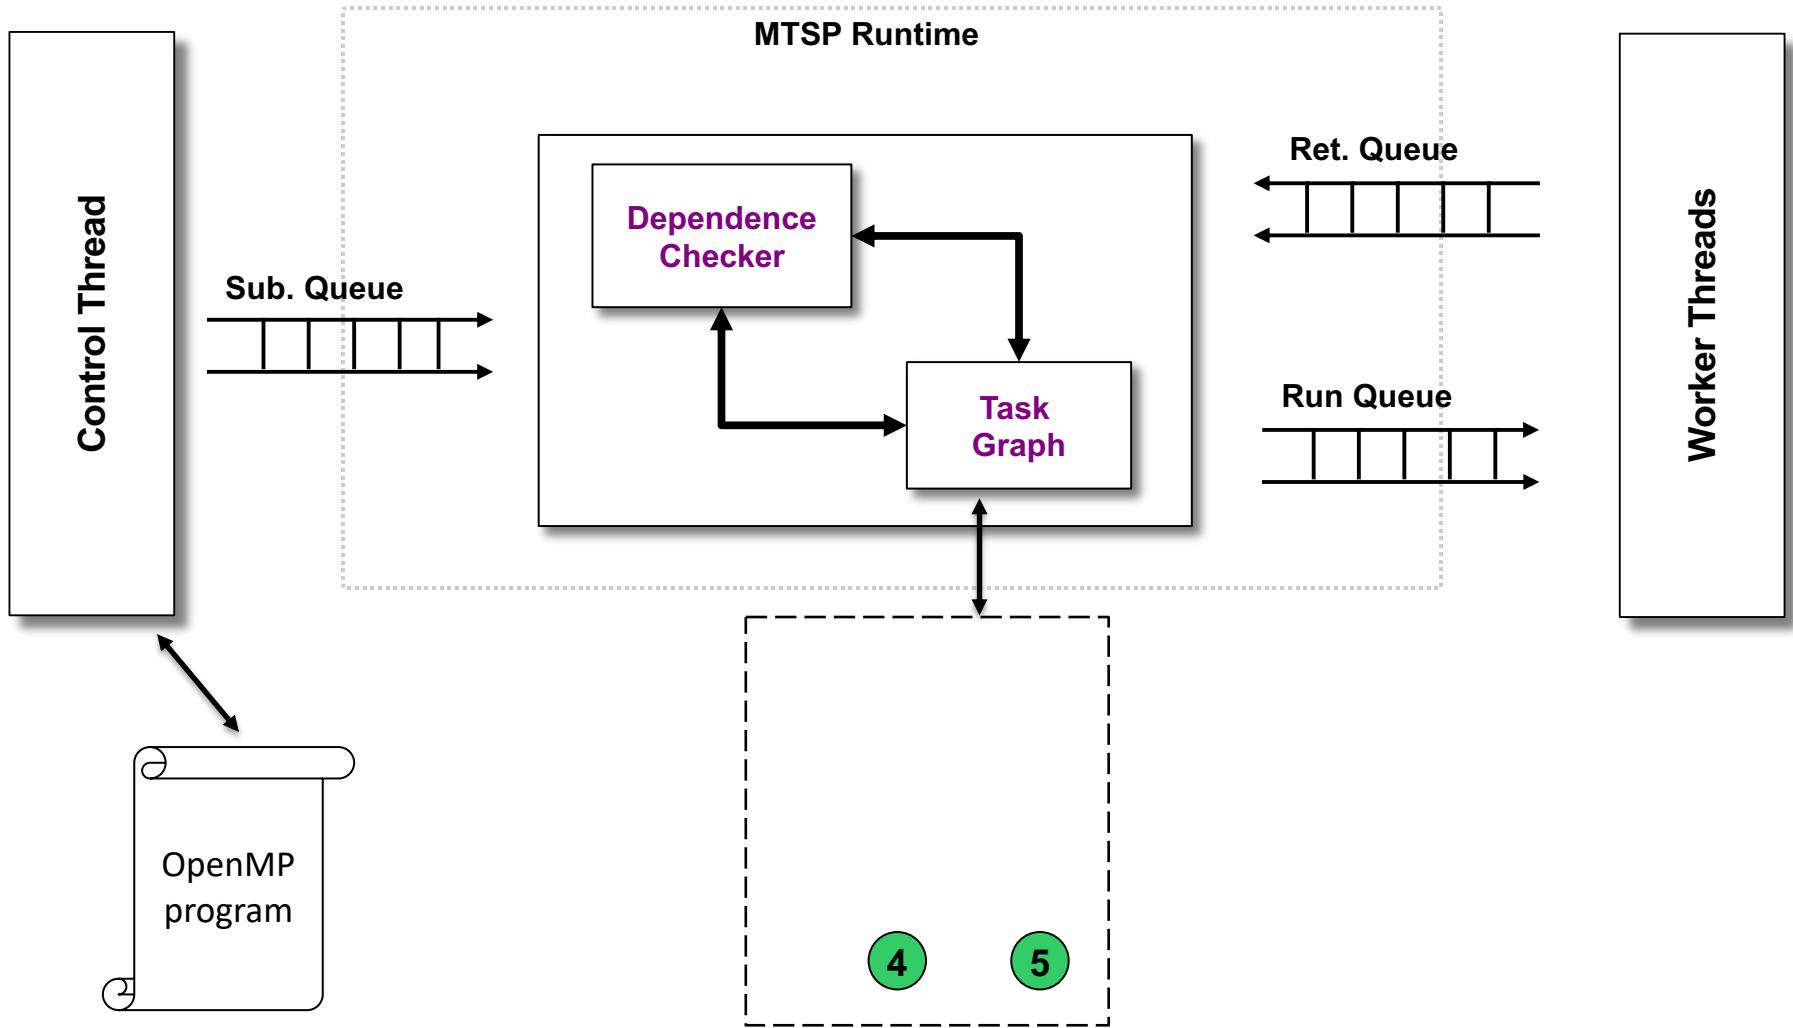
# Basic elements of a task scheduling system



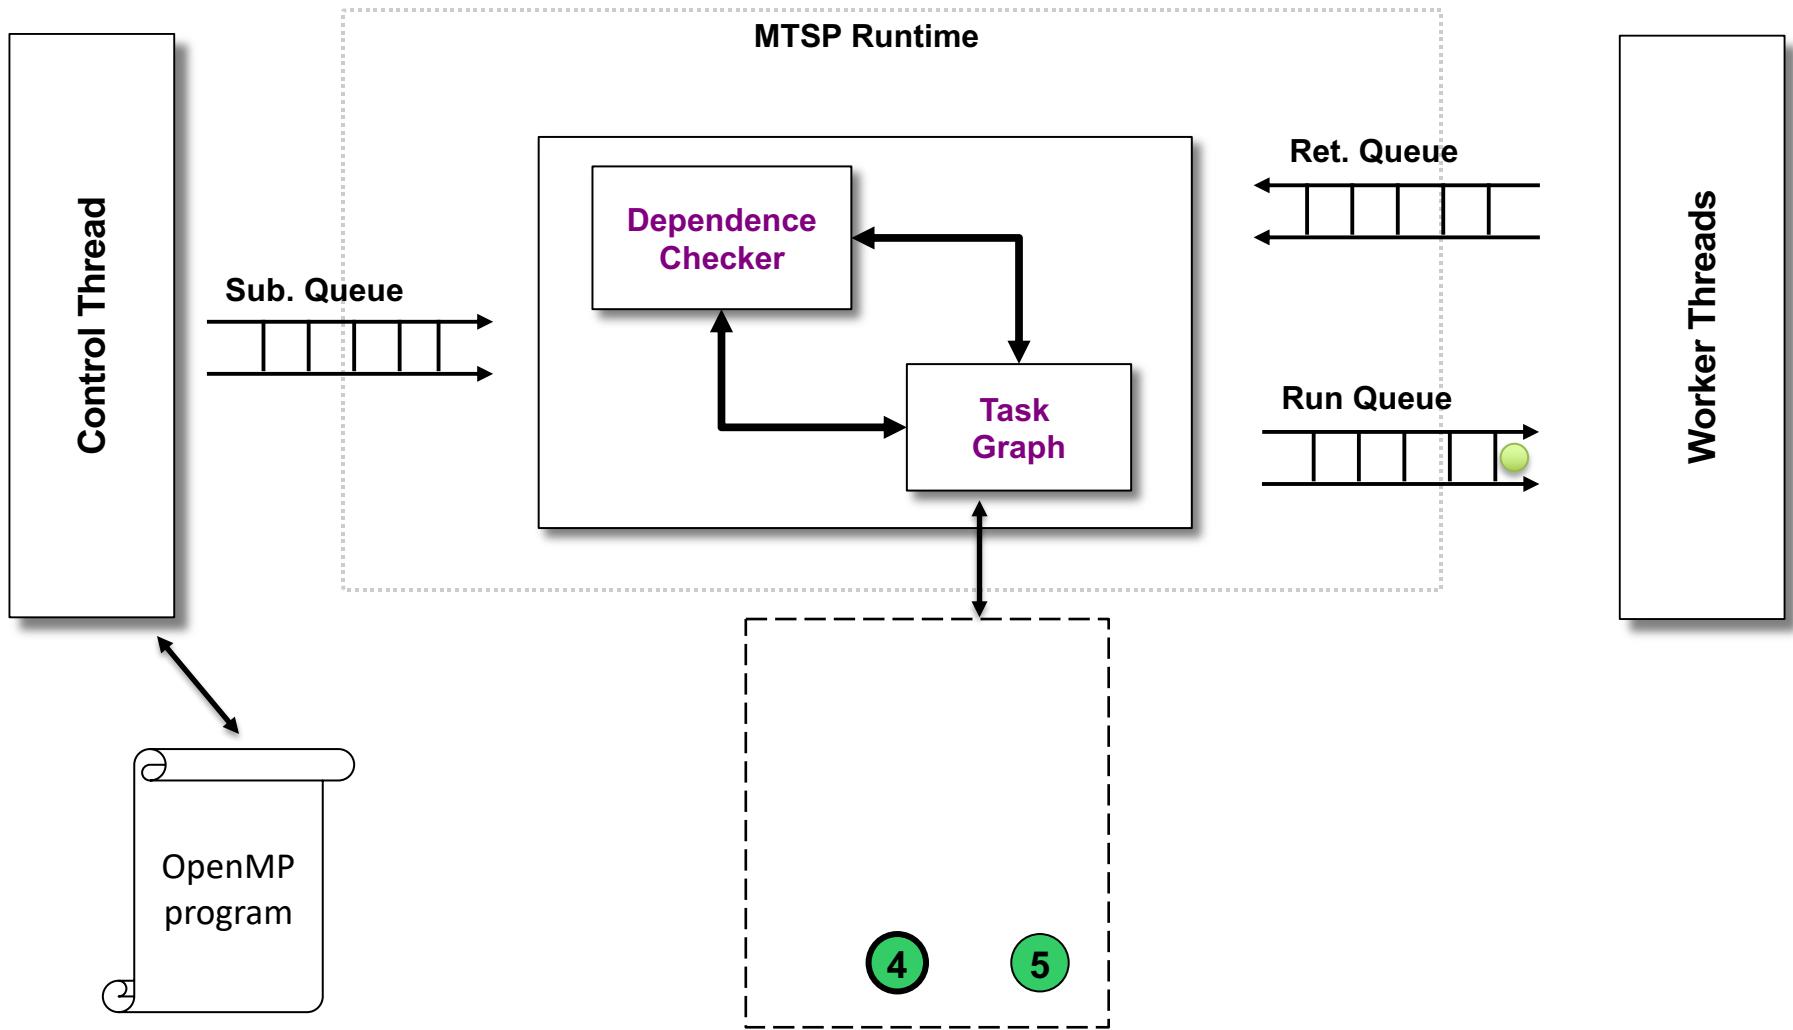
# Basic elements of a task scheduling system



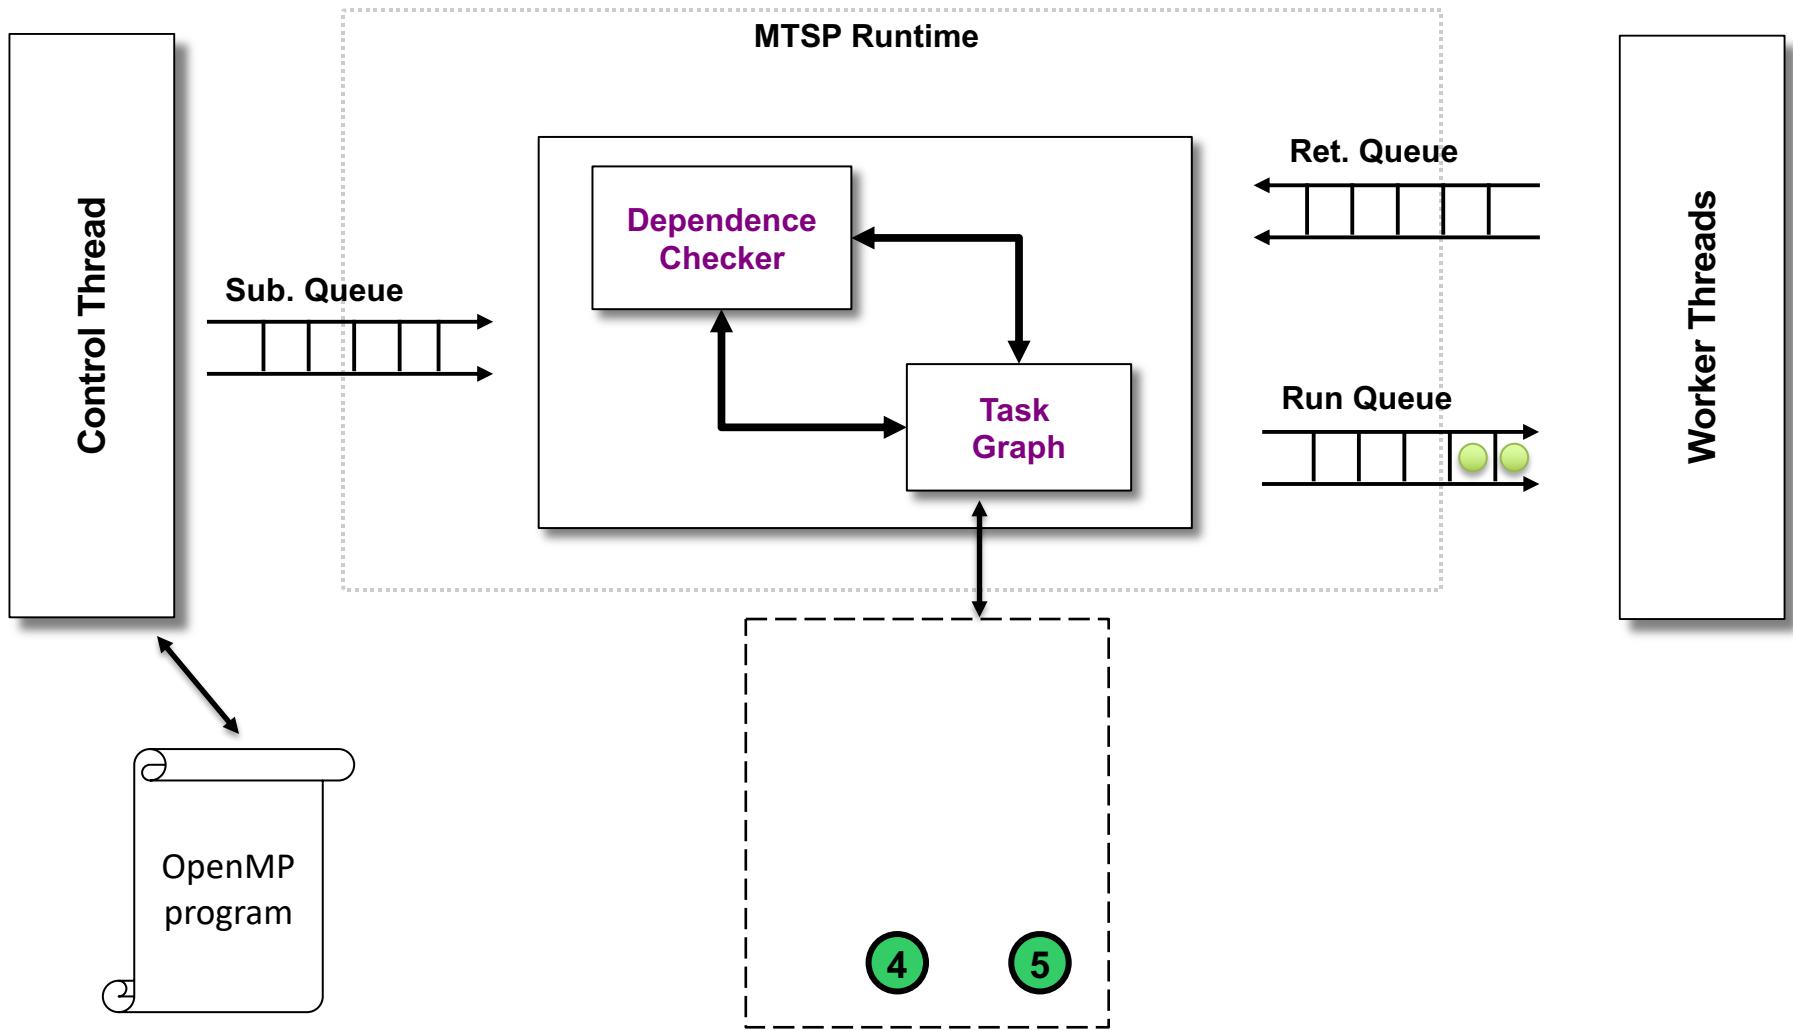
# Basic elements of a task scheduling system



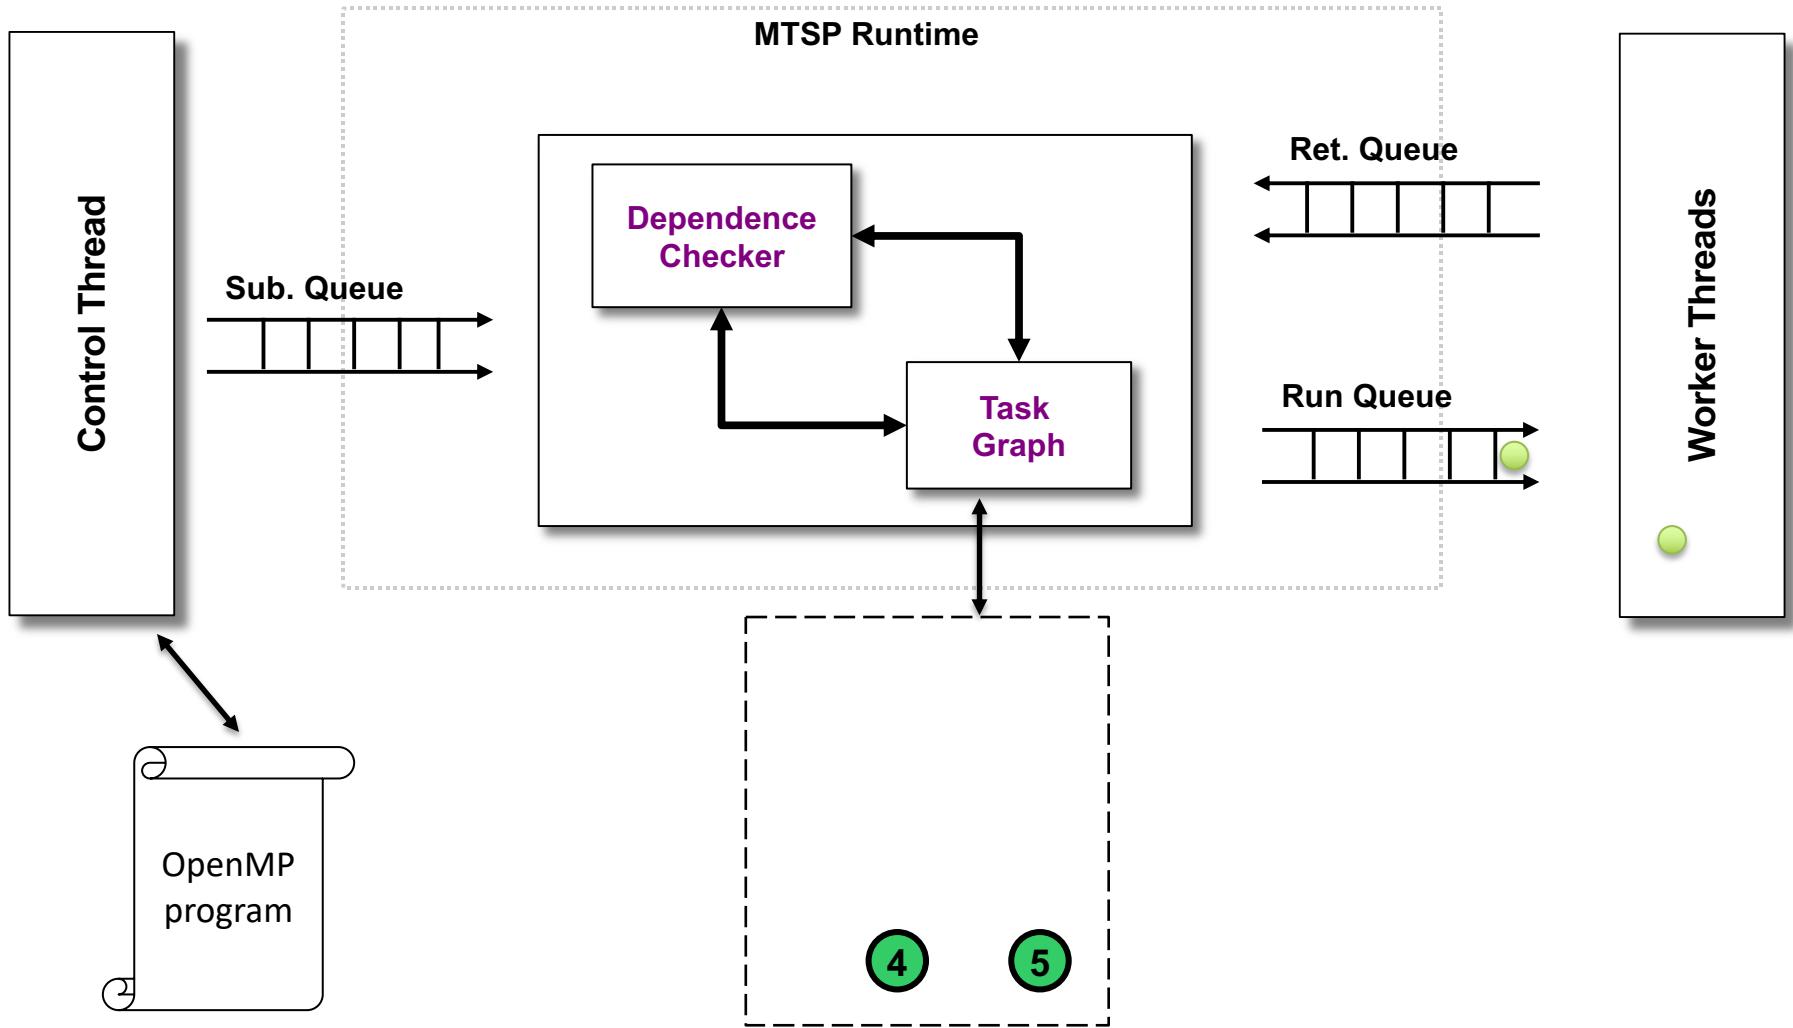
# Basic elements of a task scheduling system



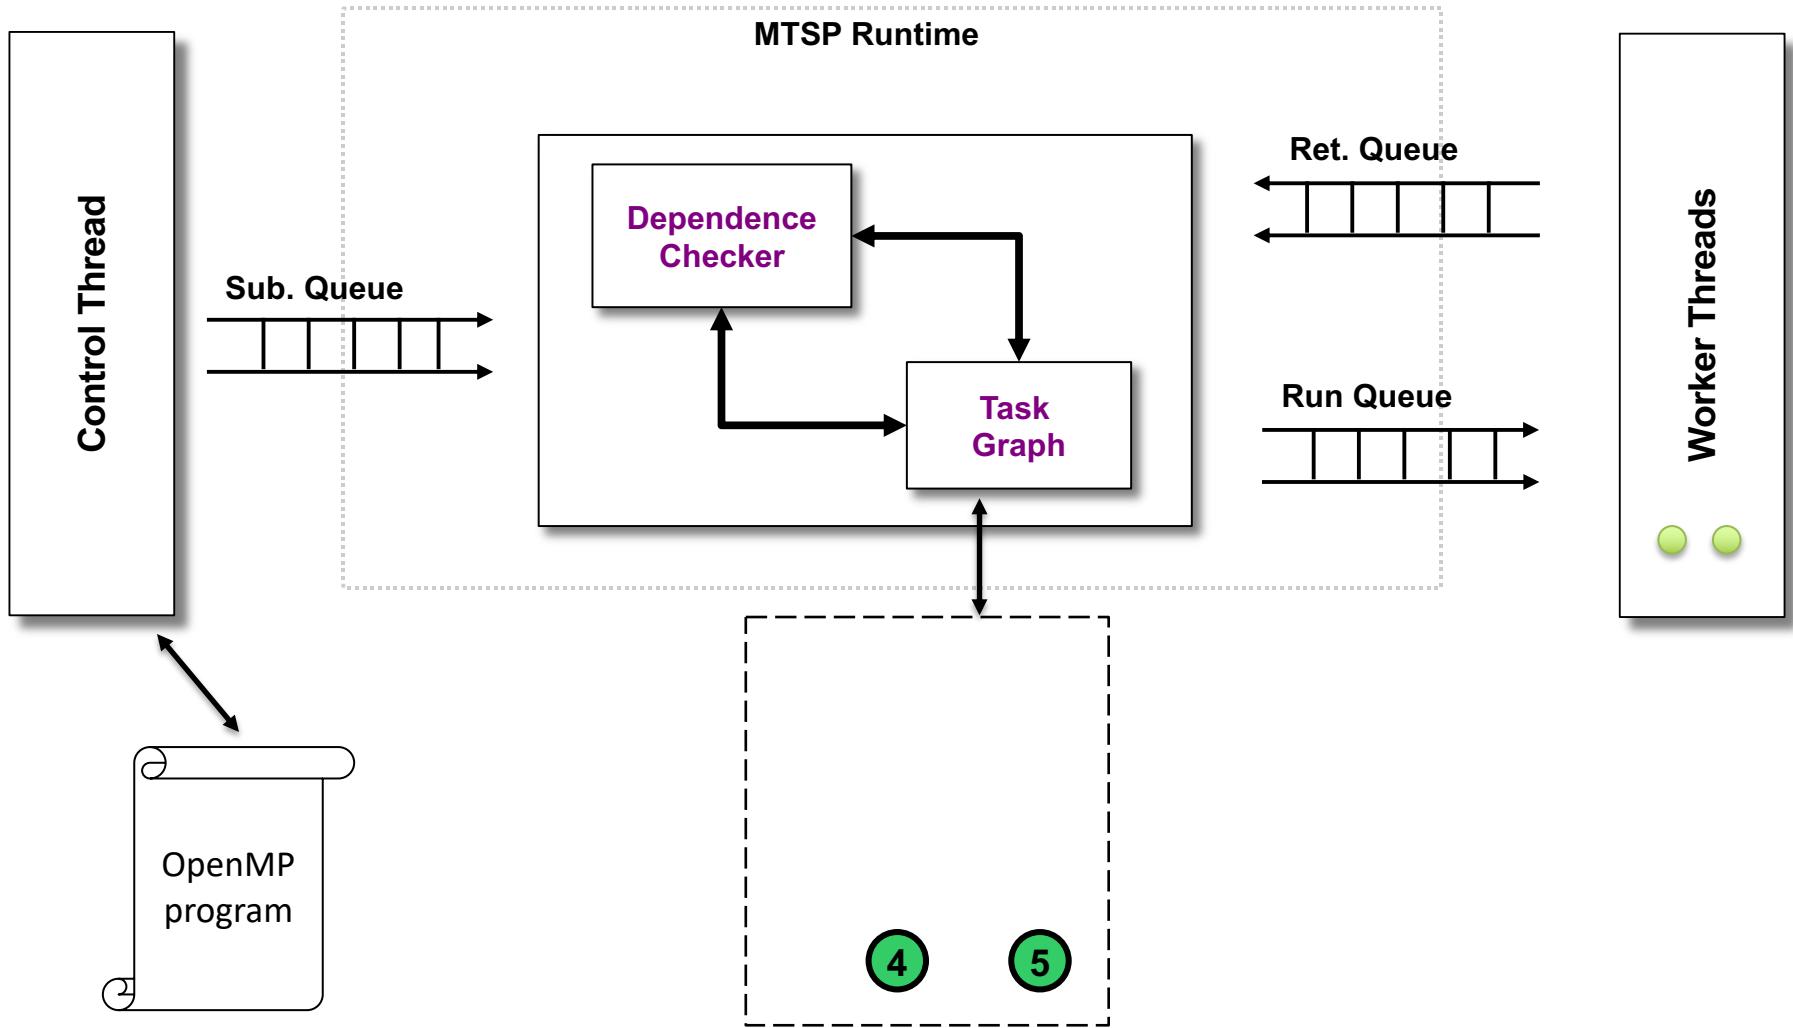
# Basic elements of a task scheduling system



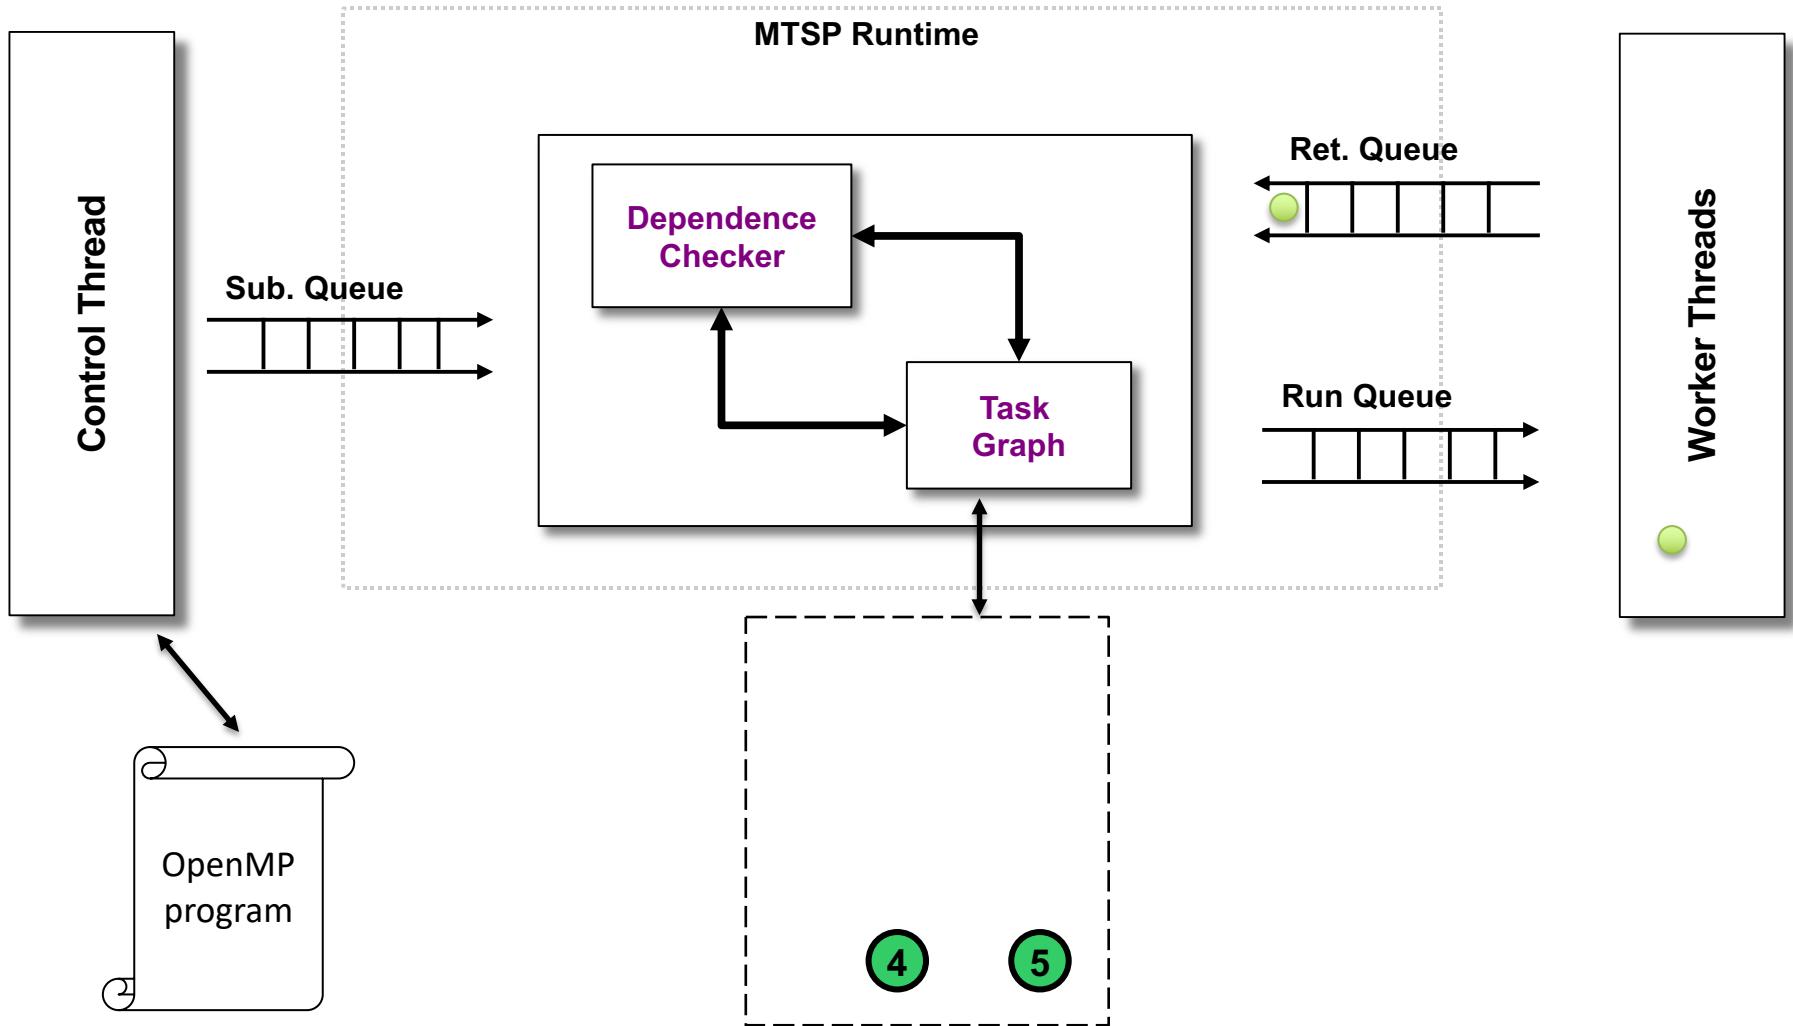
# Basic elements of a task scheduling system



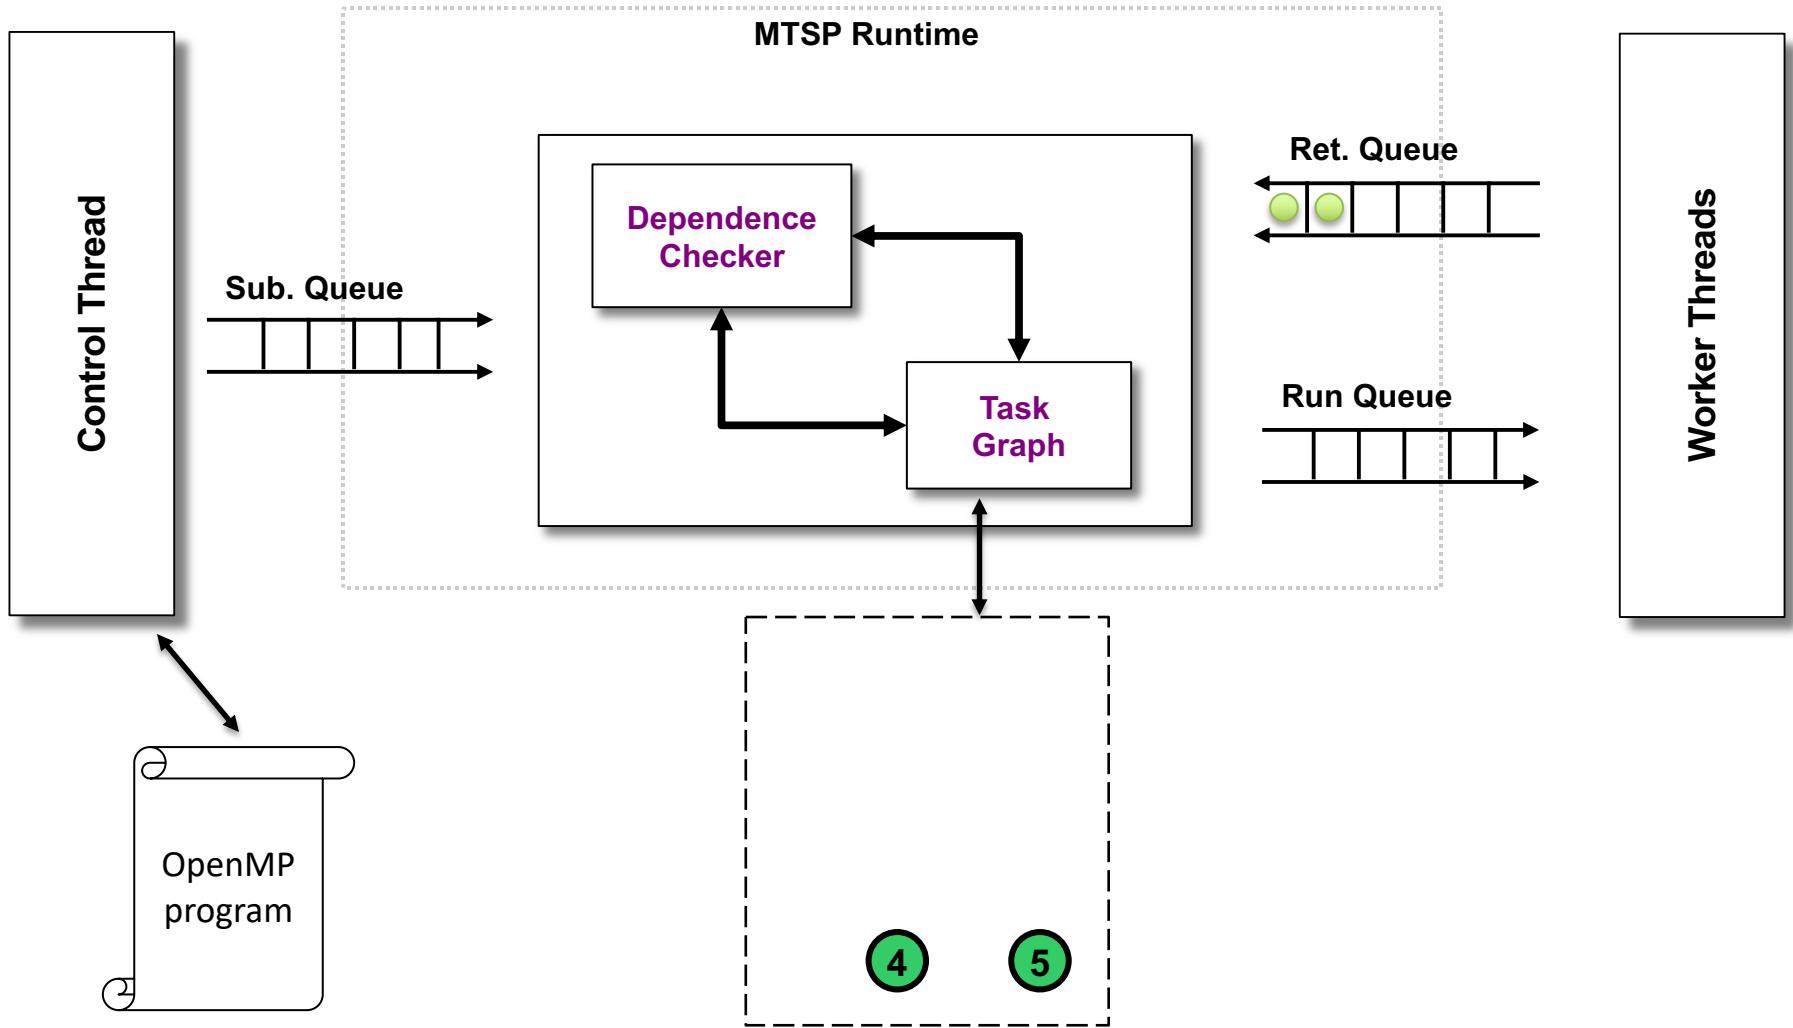
# Basic elements of a task scheduling system



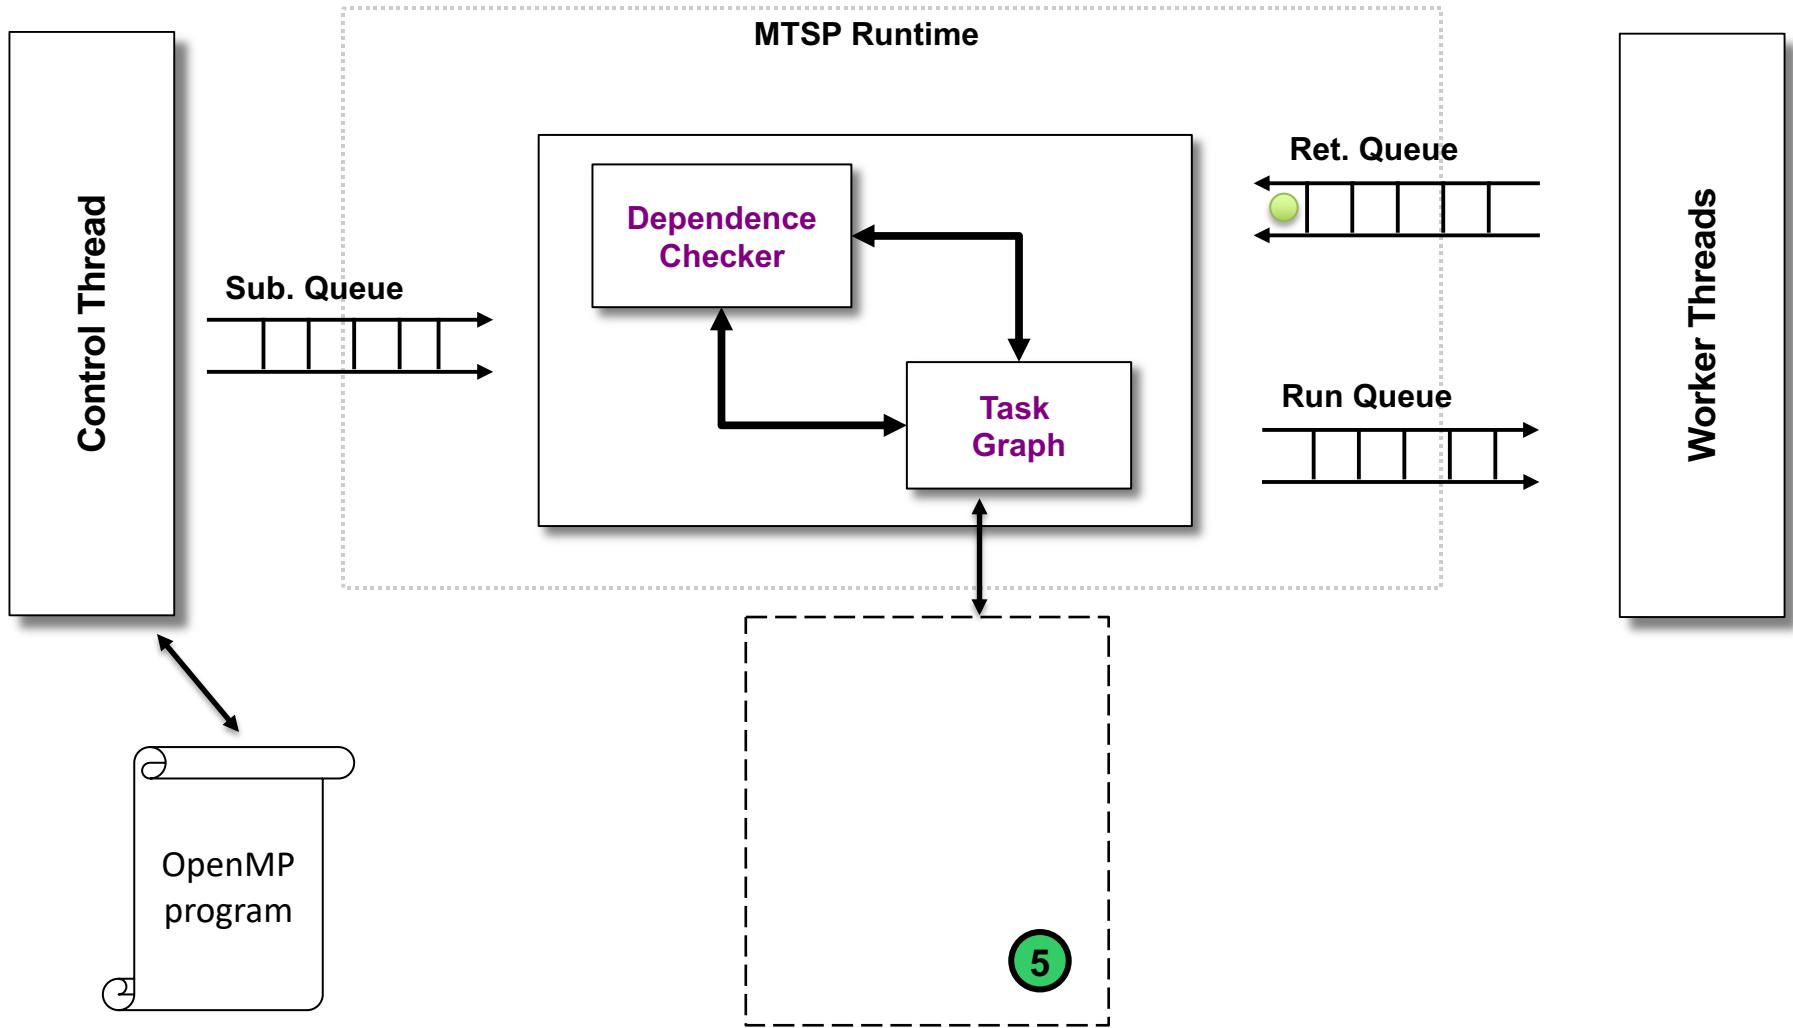
# Basic elements of a task scheduling system



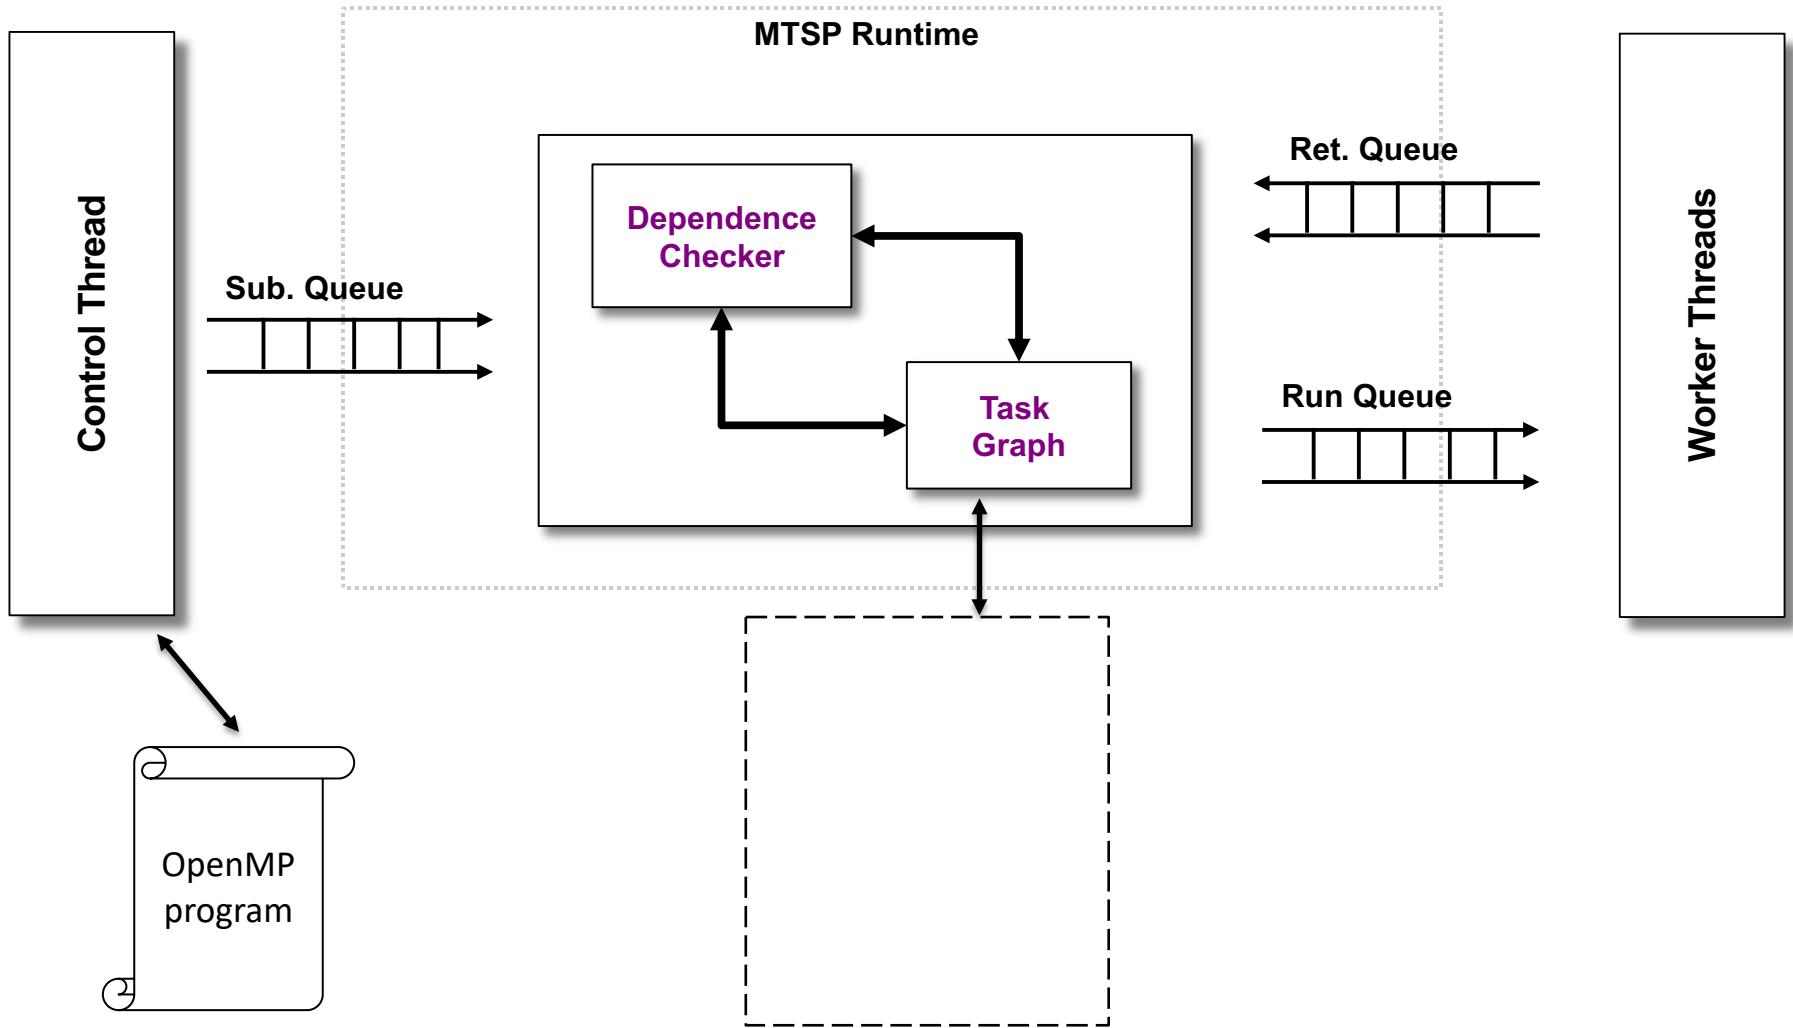
# Basic elements of a task scheduling system



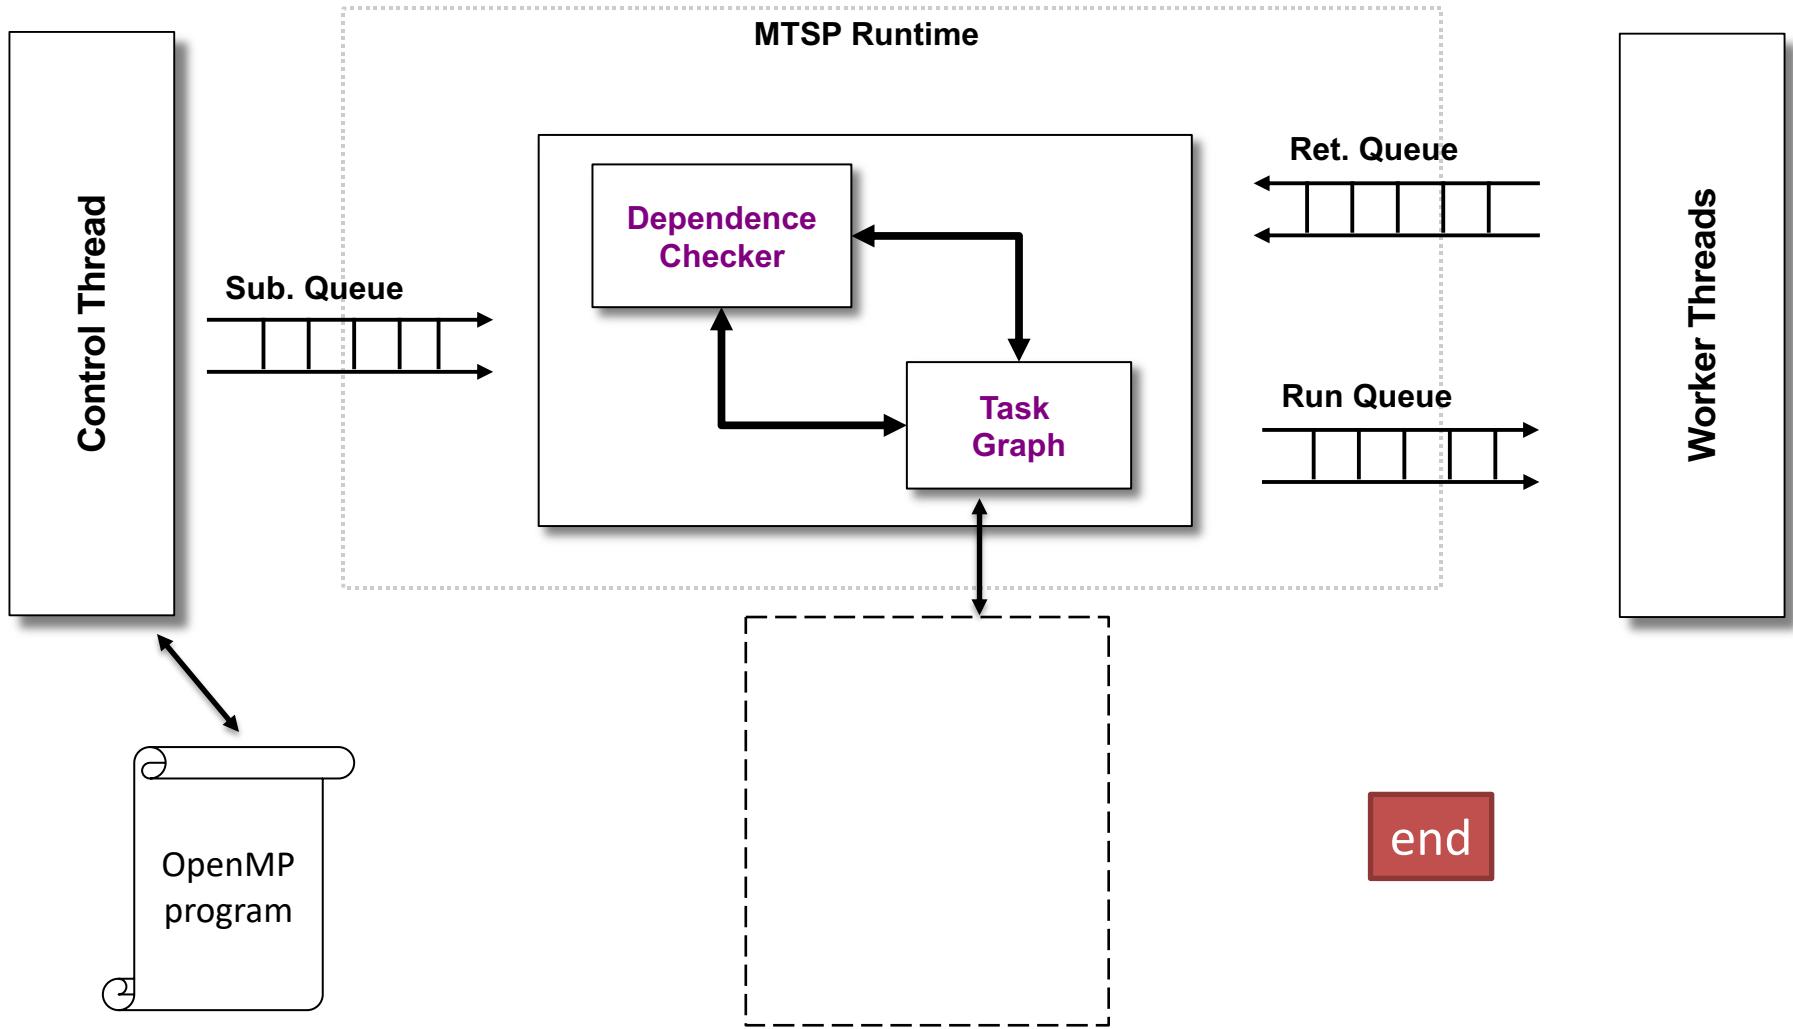
# Basic elements of a task scheduling system



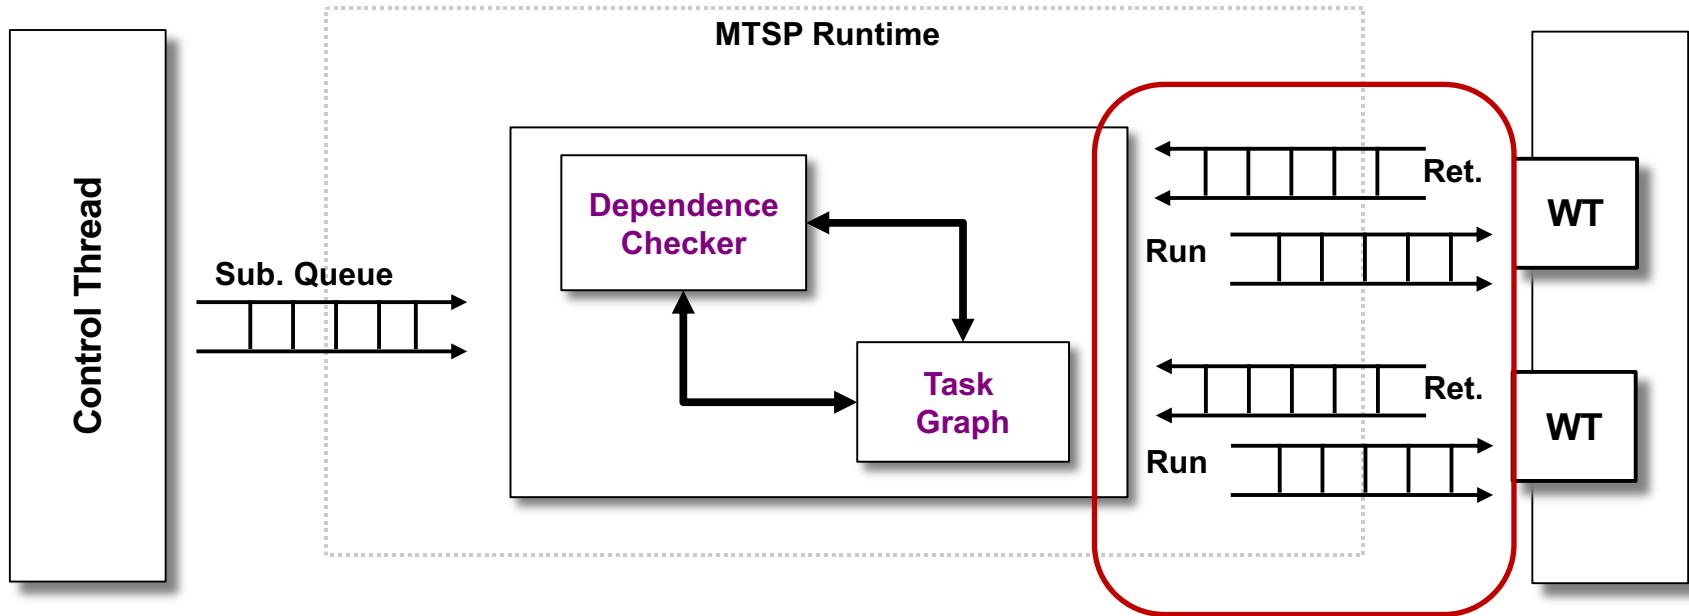
# Basic elements of a task scheduling system



# Basic elements of a task scheduling system

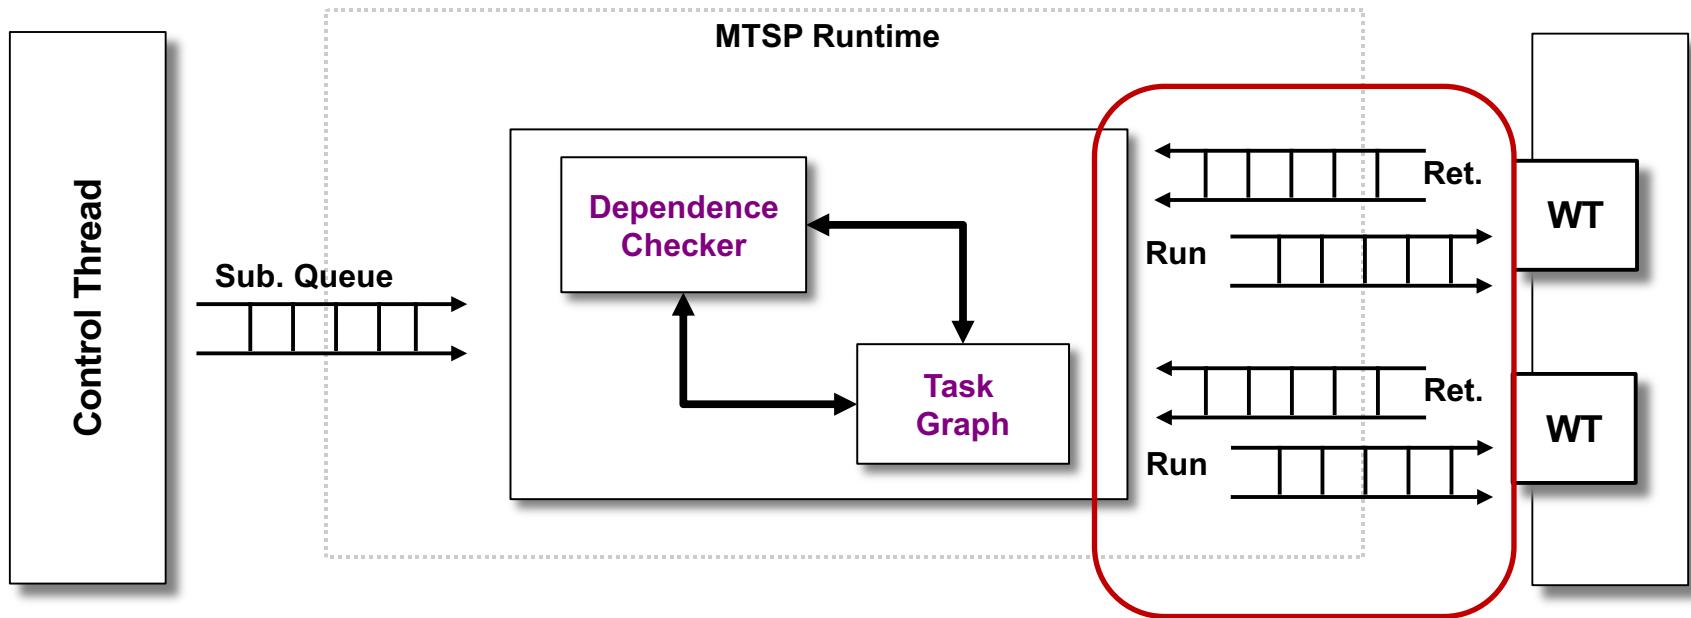


# Basic elements of a task scheduling system



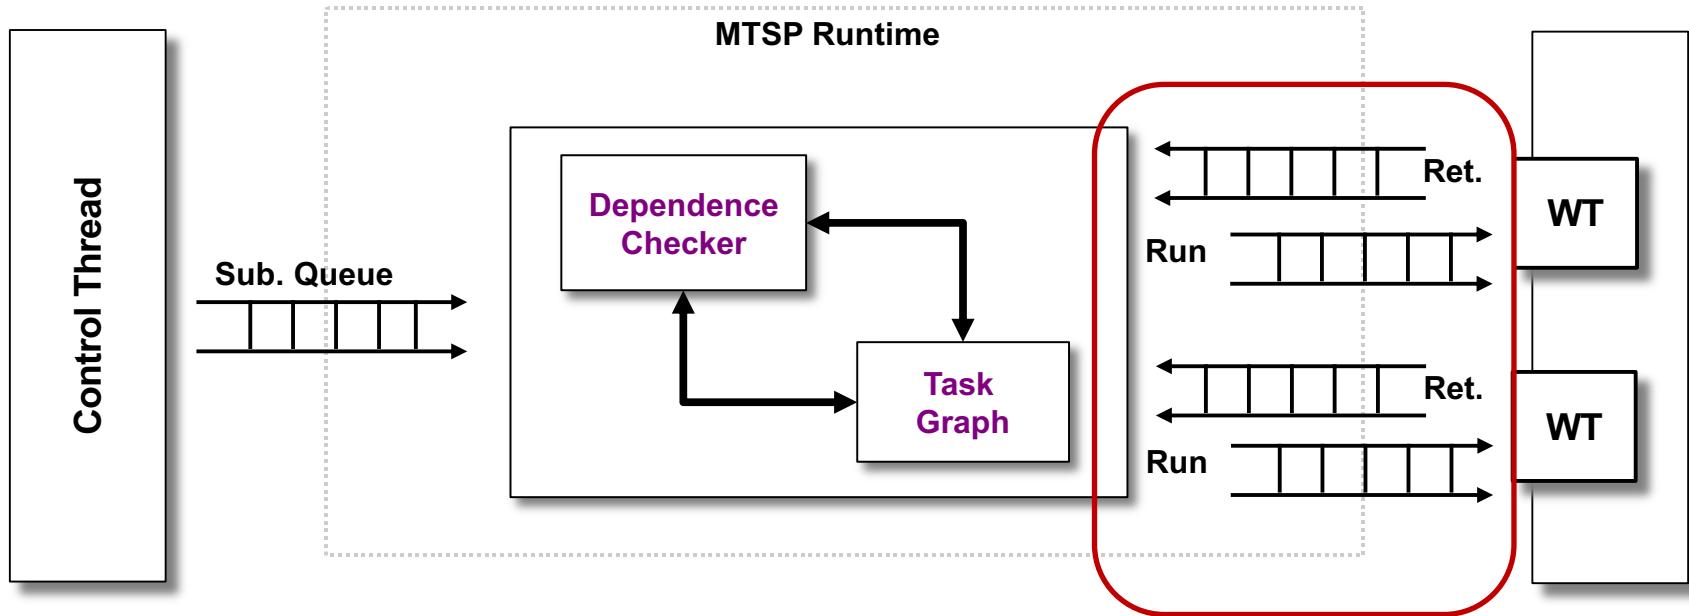
- MTSP actually has a runq/retq pair for each core

# Basic elements of a task scheduling system



- MTSP actually has a runq/retq pair for each core
- This avoids lock usage

# Basic elements of a task scheduling system

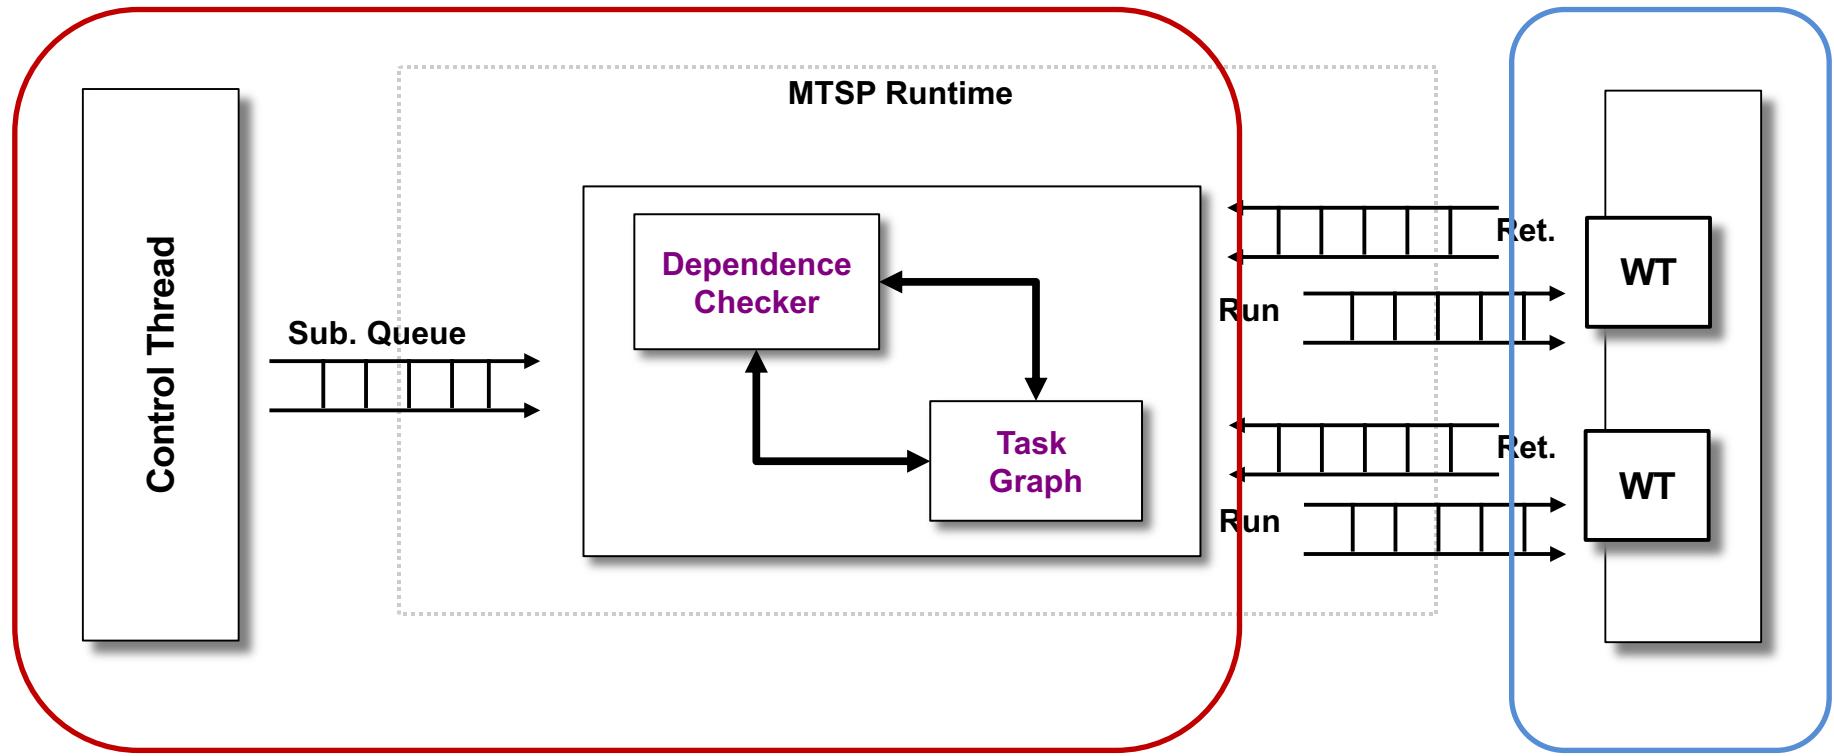


- MTSP actually has a runq/retq pair for each core
- This avoids lock usage
- But we must take care of load balancing

# Agenda

- Loop Parallelism
- Task parallelism: programming model
- MTSP: A lightweight runtime
- Where are the limitations?
- Task Graph Accelerator (TGA)
- Collaboration

# A small two stage pipe



$$D_{latency}$$

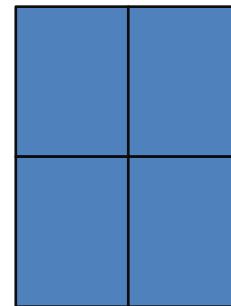
$$\frac{T_{et}}{N_{proc}}$$

# When is the pipe balanced?

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$



1 task / T<sub>et</sub>



4 tasks / T<sub>et</sub>

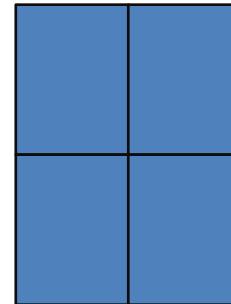
# When is the pipe balanced?

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$



1 task /  $T_{et}$

$$N_{proc} < \frac{T_{et}}{D_{latency}}$$



4 tasks /  $T_{et}$

Must keep pipeline balanced!!

# When is the pipe balanced?

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$

$$N_{proc} < \frac{T_{et}}{D_{latency}}$$

Upper bound for the number of serviced processors

# Software task scheduling limitations

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$

- Higher parallelism means higher  $N_{proc}$  ↑

$$\uparrow N_{proc} < \frac{T_{et}}{D_{latency}}$$

# Software task scheduling limitations

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$
$$\uparrow N_{proc} < \frac{\uparrow T_{et}}{D_{latency}}$$

- Higher parallelism means higher  $N_{proc}$  
- Higher  $N_{proc}$  means either
  - Higher  $T_{et}$    
(coarse grained tasks)

# Software task scheduling limitations

$$D_{latency} = \frac{T_{et}}{N_{proc}}$$
$$\uparrow N_{proc} < \frac{T_{et}}{D_{latency}} \downarrow$$

- Higher parallelism means higher  $N_{proc}$
- Higher  $N_{proc}$  means either
  - Higher  $T_{et}$   
(coarse grained tasks)
  - Lower  $D_{latency}$    
(faster scheduling)

# Software task scheduling limitations

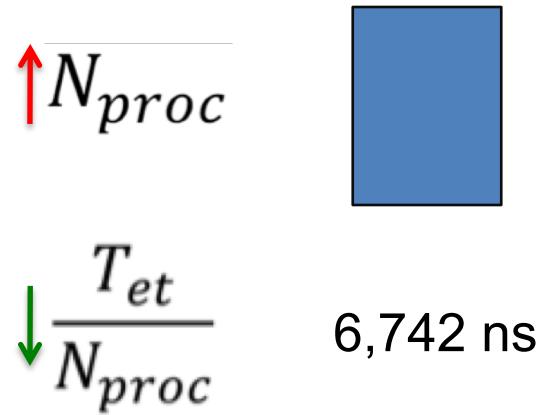
$$D_{latency} = \frac{T_{et}}{N_{proc}}$$

$$N_{proc} < \frac{T_{et}}{D_{latency}}$$

**Hardware accelerated task scheduling** could change what is possible to do with task parallelism

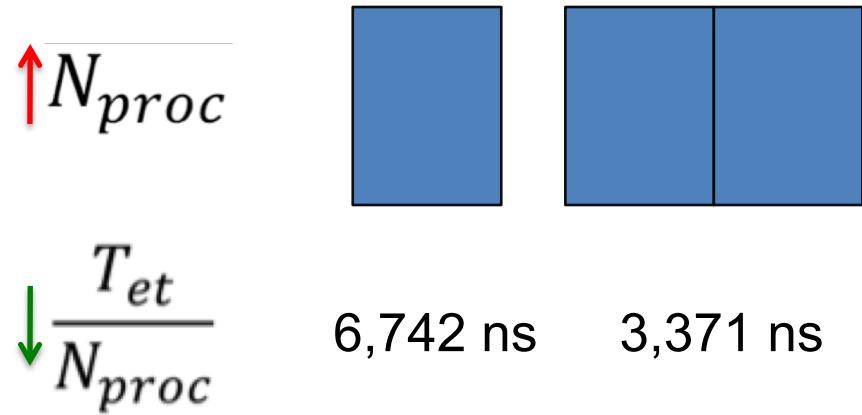
# Software task scheduling limitations

Type A – 6,742 ns



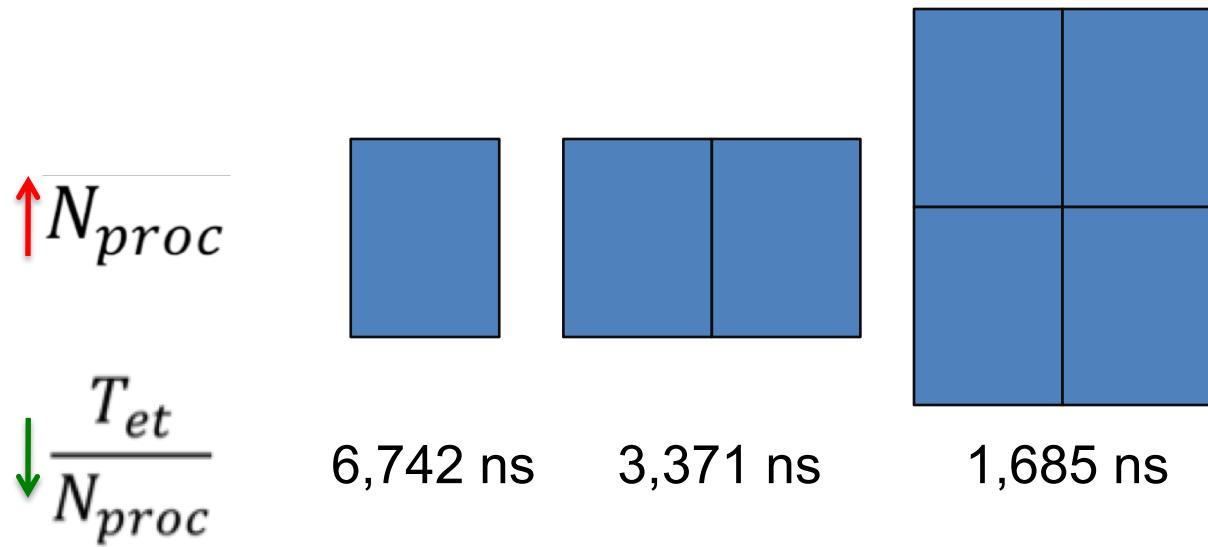
# Software task scheduling limitations

Type A – 6,742 ns



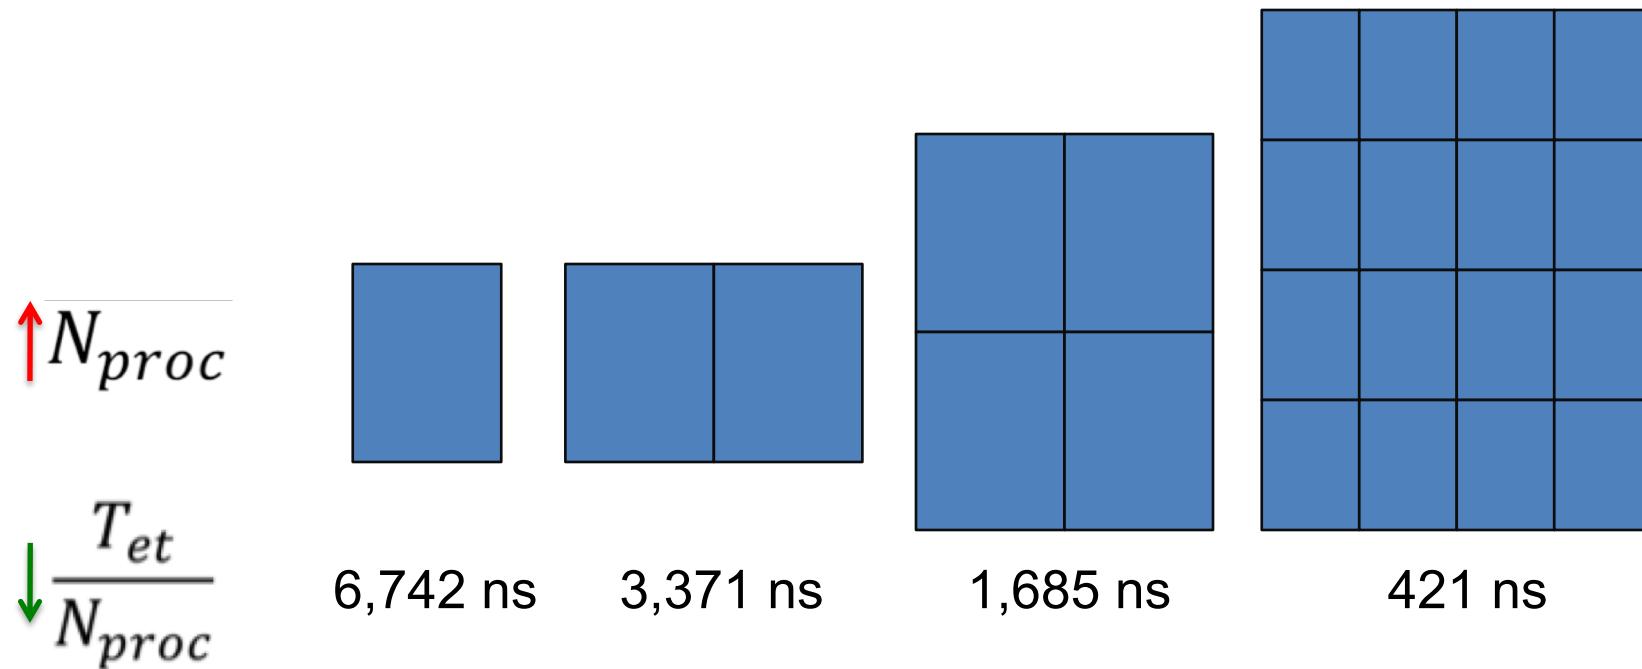
# Software task scheduling limitations

Type A – 6,742 ns



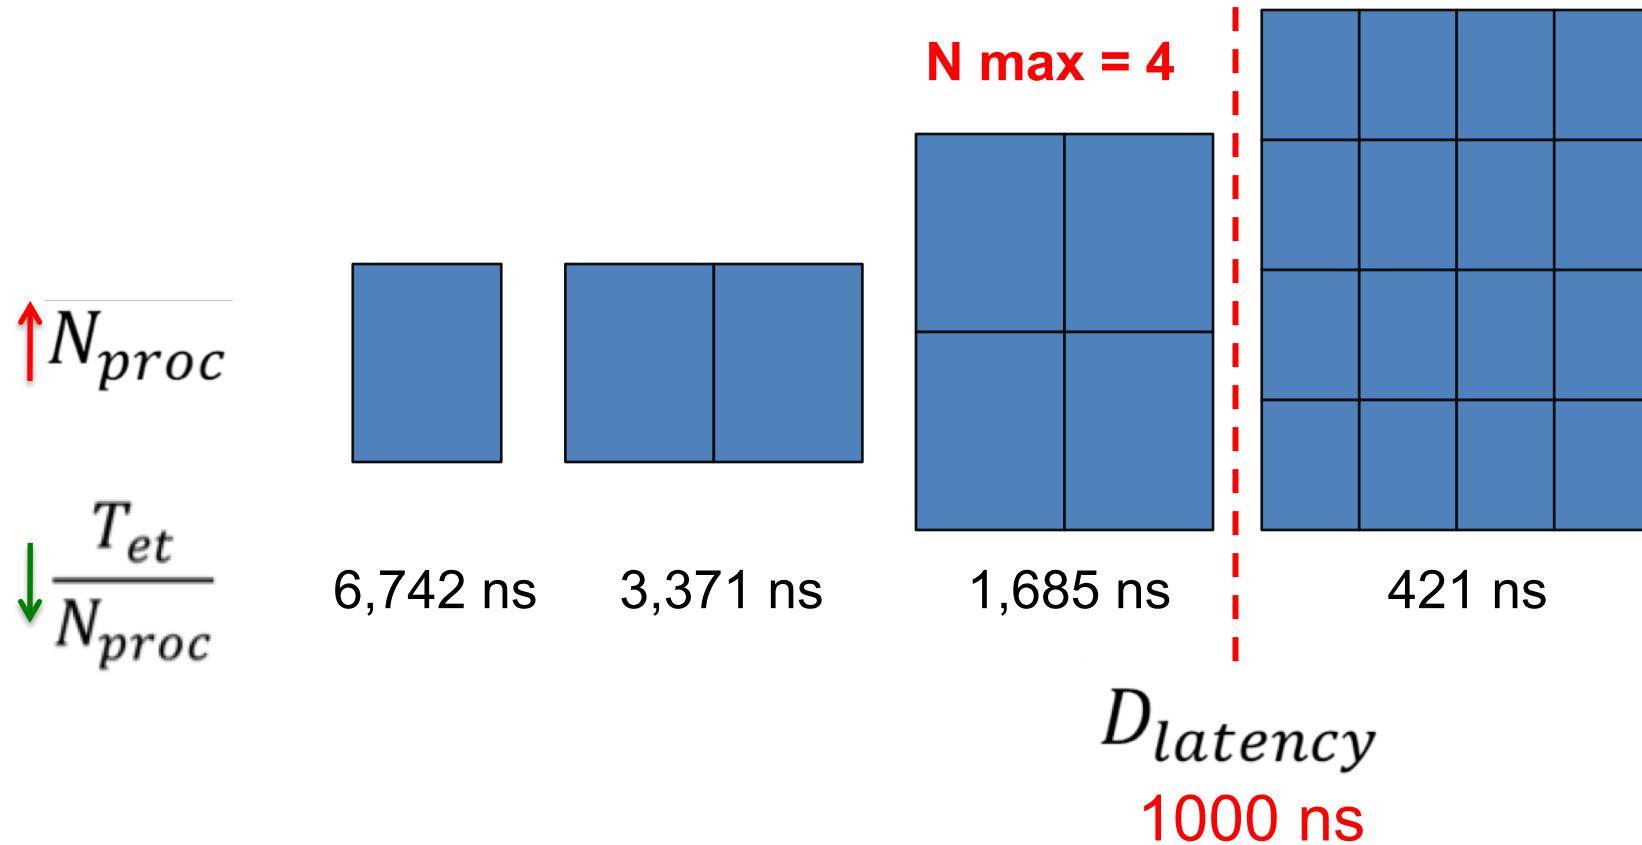
# Software task scheduling limitations

Type A – 6,742 ns



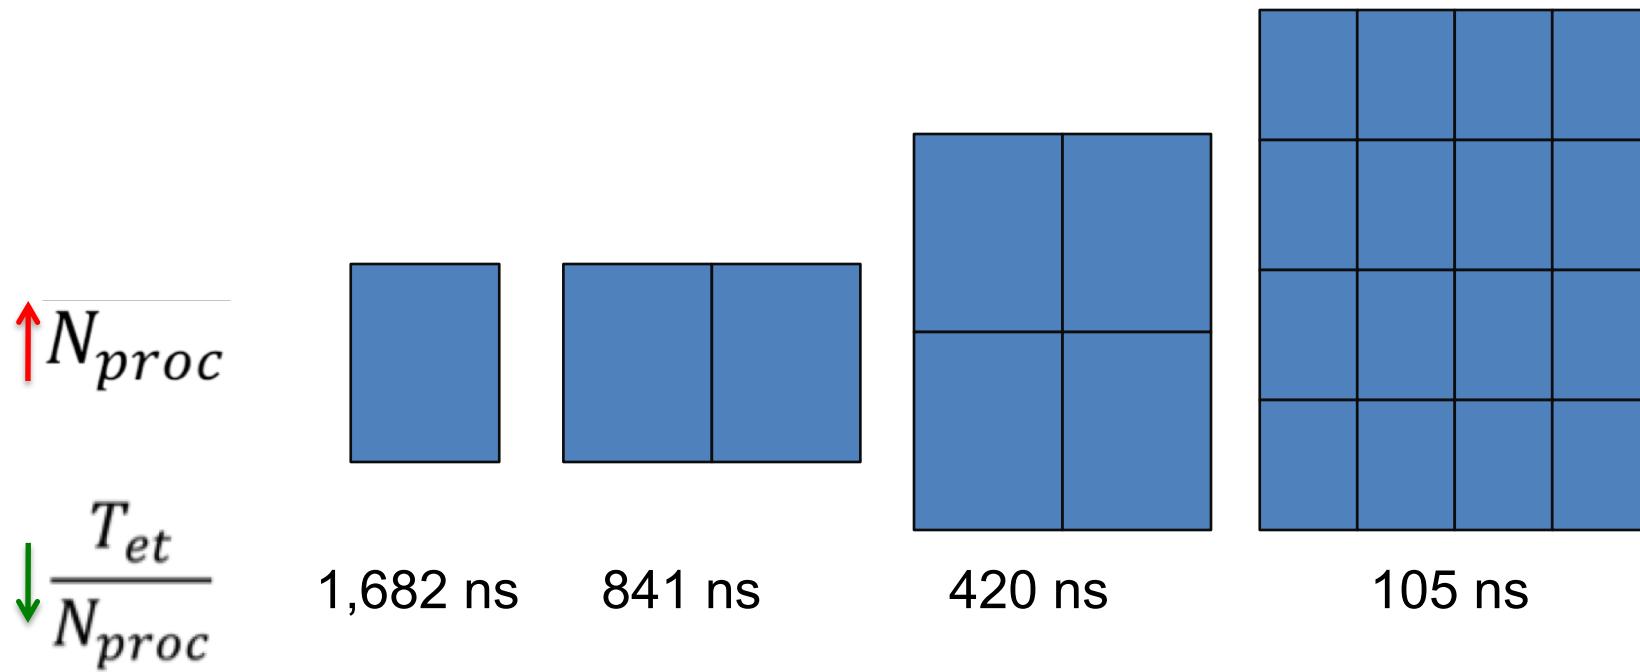
# Software task scheduling limitations

Type A – 6,742 ns



# Software task scheduling limitations

Type B – 1,682 ns



# Software task scheduling limitations

Type B – 1,682 ns

